

CS111: Operating Systems

Xu

Winter 2020

Contents

| | |
|---|----------|
| CS111: Operating Systems | 3 |
| Introduction to OS | 3 |
| Services | 6 |
| Delivery | 6 |
| Interfaces | 7 |
| Creating Programs and Linking | 9 |
| Linking Libraries | 10 |
| Process Virtualization | 11 |
| Subroutine Stack Frames | 12 |
| Process Overview | 13 |
| UNIX Process API | 15 |
| Process Mechanisms | 18 |
| Signals | 18 |
| Direct Execution | 18 |
| Restricted Operations | 19 |
| Process Switching | 21 |
| Process Scheduling | 23 |
| Feedback Priority Scheduling | 25 |
| Realtime Systems | 26 |
| Inter-Process Communication | 27 |
| Memory Virtualization | 28 |
| UNIX Memory API | 29 |
| Memory Mechanisms | 29 |
| General Partition Strategies | 29 |

| | |
|--|----|
| Address Translation | 31 |
| Segmentation | 32 |
| Free-Space Management | 33 |
| Paging | 36 |
| Translation Lookaside Buffer (TLB) | 38 |
| Swapping | 41 |
| Swapping Policies | 43 |
| Concurrency | 45 |
| Locks | 46 |
| Locks with Data Structures | 49 |
| Condition Variables | 50 |
| Semaphores | 52 |
| UNIX Synchronization API | 57 |
| Appendix | 61 |
| UNIX Syscalls | 61 |
| Sockets Example | 62 |

CS111: Operating Systems

Introduction to OS

- **Von Neumann** model of computing:
 - when a program is run, the processor repeatedly *fetches* an instruction from memory, *decodes* it, and *executes* it
- **OS Principles**
- *complexity* management principles:
 - layered structure and hierarchical decomposition
 - modularity and functional encapsulation
 - appropriately abstracted interfaces and information hiding
 - powerful abstractions
 - interface contracts
 - progressive refinement
- *architectural* paradigms:
 - mechanism/policy separation
 - indirection, federation, and deferred binding
 - dynamic equilibrium
 - criticality of data structures
- the **operating system (OS)** is in charge of making the system operates correctly and efficiently in an *abstracted*, easy-to-use manner:
 - acts as the software layer between hardware and higher level applications, abstracts and hides the low level details eg. hardware and ISAs
 - uses technique of **virtualization** to transform a *physical* resource into a generalized, easy-to-use *virtual* form
 - * OS thus also known as **virtual machine**
 - * eg. in order to virtualize memory, each running program seems to have its own private memory, instead of sharing the actual physical memory
 - provides services through interfaces and **system calls** in a **standard library** that users can use
 - acts as a **resource manager** to manage resources such as the CPU, memory, and disk
 - * eg. abstracts physical memory *disks* as *files*
 - virtualizes the CPU, ie. turning a small number of CPUs into infinite CPUs

that can run many programs at once

- * this **concurrency** can lead to different problems for the OS itself as well as **multi-threaded** programs that require certain mechanisms to solve
 - handles data **persistence** with the file system and I/O
 - deals with **drivers** and coordination with external devices
- basic OS goals include *abstraction*, minimizing *overhead* (ie. in time or space), providing *protection* and *isolation* between applications, and high *reliability*
- original role has changed over time from harnessing hardware, shielding applications from hardware, to providing an ABI platform, to acting as a “traffic cop”
 - over time, different OSs have converged, since they are so difficult to *maintain*
 - applications have to *choose* to support an OS
 - new OSs must have some clear advantages over alternatives
- **instruction set architectures (ISAs)** are a computer’s lowest-level supported instructions/primitives
 - many different, incompatible ISAs
 - * thus OS also responsible for running on *different* ISAs and abstracting them
 - only OS/kernel can work with the *privileged* ISA, but standard ISA is accessible by all
- OS abstracts ISAs into a set of management and abstraction *services* accessible through a **system call interface**
- system calls may be further abstracted into an **application binary interface (ABI)**
- the OS code is *unique* from application code:
 - eg. applications should not be able to read from anywhere on disk
 - thus, OS should distinguish between:
 - * **system calls** that require formal hardware instructions to use the OS (ie. jumps into the privileged *kernel* mode and raises hardware privilege level)
 - * **procedure calls** that are provided as a library and are accessible in the *user* mode
- **resources** have different types:
 - **serial** - used by multiple clients, one at a time, eg. printer
 - * serial multiplexing
 - * need *graceful* transitions when switching between clients, mechanisms for exclusive use, cleanup of incomplete operations, etc.
 - **partitonable** - divided into disjoint pieces for multiple clients, eg. memory
 - * spatial multiplexing

- * need access control for containment and privacy, transitions
 - **shareable** - used by multiple concurrent clients, eg. OS shared by multiple processes
 - * no need for transitions, no unique state for particular client
- the OS should handle *abstractions* in order to:
 - *encapsulate* implementation details
 - provide more *convenient* and powerful behavior
- core OS abstractions include:
 - processor, memory, and communications abstractions
- **process abstractions:**
 - the source code of a program specifies a program's behavior
 - when running as a process, the stack, heap, and register contents form its environment
 - * must be independent from other processes
 - but the CPU thus must be shared across many processes:
 - * CPU **schedulers** share the CPU among those processes
 - * **memory management** hardware and software gives the illusions of full exclusive memory use for each process
 - * **access control** mechanisms to keep processes independent
- **memory abstraction:**
 - at a low level, there are many different related data storage resources:
 - * variables, chunks, files, database records, messages
 - * all with unique, peculiar characteristics
 - OS must abstract these physical devices to create ones with consistent, more *desirable* properties:
 - * **persistence**
 - * user desired **size**
 - * **coherency** (reads reflect writes) and **atomicity** (full writes and reads)
 - * **latency**
 - OS will thus:
 - * have a *thorough* file system component
 - * *optimize* caching
 - * have sophisticated organizations to handle *failures*
- **communications abstractions:**
 - networks and interprocess communication mechanisms
 - different from memory:
 - * highly *variable* performance
 - * *asynchronous*
 - * *complications* from working with remote machines

Services

- a **service** is a provided functionality
 - in an OS, the client of its services are applications
- decomposed into:
 - **interface** - the *specification* of the service, ie. description of pre- and post-conditions
 - **implementation** of the interface
- main types of OS services include:
 - **CPU/memory** - processes, threads, virtual addresses, lowest latency memory
 - **persistent storage** - disks, files, and file systems, higher latency memory
 - **I/O** - terminals, windows, sockets, networks, signals (interrupts), highest latency memory
 - note each service family can be associated with a memory latency class
 - * when CPU is waiting for a process's slow memory access, may use **context switching** to switch to a different process
- *higher* level OS services:
 - used by clients
 - cooperating parallel processes
 - security, eg. authentication and encryption
 - providing a UI
- *lower* level OS services:
 - not as visible
 - hardware handling
 - software updates, config registry
 - resource allocation and scheduling
 - network and protocols

Delivery

- the OS *delivers* these various services at different layers:
 - **subroutines** (functions), eg. `malloc()` provided by `libc` library (implementation uses a system call)
 - * implemented at higher layers to provide richer operations
 - * simplest access, just call subroutines
 - at a lower level: push parameters, jump, return values in registers
 - * **pros:**
 - fastest, can be implemented to use the fewest system calls, eg.

- buffered read and writes
 - can bind implementations at runtime
- * **cons:**
 - services implemented in same virtual address space *associated* with the running program
 - limited to a language
 - can't use privileged instructions
- * provided in **libraries**
- * **pros:**
 - code reuse, single copy, encapsulates complexity
 - many bind-time options: *static* (included at link time), *shared* (mapped into address space at exec time), *dynamic* (choose at load time)
- **system calls**
 - * forces an entry into the OS, implementation uses privilege kernel
 - * **pros:**
 - can use privileged resources and operations
 - can communicate with other processes
 - * **cons:**
 - very specific use cases, eg. viewing status of a page table
 - slower, the process may have to switch to a privileged kernel mode
 - requires hardware to **trap** into the OS
- send **messages** to software that performs services
 - * used in distributed systems, exchange messages with a server
 - * **pros:**
 - server can be anywhere
 - service is highly scalable and available
 - * **cons:**
 - slowest method

Interfaces

- standardized **interfaces** in software are inspired by the concept of *interchangeable* parts
 - ie. every part has specifications that allow any collection of parts to be assembled together
 - * *pros*: standards end up being extensively reviewed, platform-neutral, and clear and complete
 - * *cons*: standards constrain possible implementations and consumers,

- and can be hard to evolve, leading to obsolescence
- * **proprietary** interfaces are controlled by a single organization, which puts the burden on the org to develop it
- * **open standards** are controlled by a consortium of providers, which may lead to reduced freedom and competitive advantage
- using interfaces for the components of a complex system architecture allows for modularity and independent designs and implementations
 - but interfaces and implementations should be defined *independently*
- an interface's specifications is a **contract** between developers and the implementation providers
 - if this contract is broken, programs are no longer portable and solving issues becomes more complex
 - **backwards compatibility** can still be maintained with some strategies:
 - * **interface polymorphism** for different versions of a method with unique signatures
 - * **versioned interfaces** with micro, minor, or major releases
- an **application programming interface (API)**:
 - defines *subroutines*, what they do, and how to use them
 - ie. a source level interface, helps write programs for the OS
 - includes discussion of signatures, options, return values and errors
 - eg. in a simple “Hello World”, two system calls are made using their respective APIs:
 - * `write(fd, p, num)` writes `num` bytes from the address at `p` to the file descriptor `fd`
 - * `exit_group(code)` exits the program with exit code `code`
- an **application binary interface (ABI)**:
 - *binds* an API to an ISA
 - * applications work *above* the ABI, while the kernel and machine level operations lie *under* the ABI
 - ie. a binary interface specifying the DLLs, data formats, calling sequences, linkage conventions
 - * help install binaries on the OS
 - describes the *machine language* instructions and convention to call routines for a specific ISA
 - * eg. the binary representation of data types, stack-frame structure, register conventions
 - usually used by the compiler, linker, loader, and OS
 - eg. in the above “Hello World”, the system call ABI for Linux x86-64 consists of the assembly instruction `syscall`
 - * where the register `rax` holds the system call number, and registers `rdi`

-r9 hold the 6 possible arguments

Creating Programs and Linking

- general software file classes:
 - **source** files are editable text files in a programming language
 - **object modules** are relocatable sets of compiled or assembled instructions from source files
 - **libraries** are collections of object modules, source files can fetch functions from them
 - * order in which libraries are searched can matter
 - **load modules** are complete programs that can be loaded into memory and executed into CPU
- software generation tool chain:
 - **compiler** produces lower-level assembly language code from source modules
 - **assembler** creates an object module in mostly machine language code from assembly language files
 - * handles lower-level operations including CPU initialization, traps/interrupts, synchronization
 - * however, some functions and data may not yet be present and not all memory addresses are finalized
 - ie. references and addresses can only be relative to the start of the module addresses
 - **linkage editor** reads a set of object modules, places them into a virtual address space (VAS), *resolves* external references in the VAS, and finalizes all symbol addresses
 - * creates an executable load module
 - * **resolution** searches through specified libraries to find object modules that satisfy unresolved references
 - * **loading** lays out text and data segments from modules into one VAS
 - * **relocation** fixes relocation entries and updates addresses
 - **program loader** is a part of the OS that creates a virtual address space, loads in instructions and data from executable, resolves references to additional *shared* libraries
 - * reads segments into memory, creates a stack, initializes stack pointer
 - * program can then be executed by the CPU
 - * *symbol tables* are used primarily for debugging

- **executable and linkable format (ELF)** is an object module format shared across different ISAs. Includes:
 - **header** with types, sizes, locations
 - **code and data**
 - **symbol table** for external symbols and references
 - **relocation** entries

Linking Libraries

- **static** libraries (linktime binding, mapped into memory at linktime):
 - library modules are directly and *permanently* embedded into the load module
 - *cons*:
 - * can lead to identical **copies** of the same library code in different programs
 - * difficulty keeping static libraries updated (version is *frozen*)
- **shared** libraries (linktime binding, mapped into memory at runtime):
 - reserve an address
 - linkage edit libraries into code segments
 - includes redirection table (stub library) with addresses for routines
 - at load time, libraries are *mapped* into memory
 - *pros*:
 - * only single library copy required (reduced memory consumption, cached libraries)
 - * version can be specified at load time
 - * library changes (eg. size, new routines) easy to update
 - * from client's perspective, *indistinguishable* from static libraries
 - *cons*:
 - * cannot use global data storage, since other programs will use this same library copy
 - * large, expensive libraries always loaded at startup
 - * unlike for a static library, executable will not work on clients without the used library
- **dynamic** libraries (DLLs, runtime binding, mapped into memory during runtime):
 - libraries that are not loaded until they are actually needed
 - application asks OS to load a library into its VAS
 - application receives standard *entry points* to make calls to the DLL through
 - maintains a *table* of entry points for different DLLs
 - on DLL shutdown, application asks OS to unload module

- loading DLLs is done through an API, but the actual loading mechanism is ABI-specific
- *pros*:
 - * runtime binding
 - * libraries can be unloaded when no longer required
- *cons*:
 - * more work for the client to load and manage DLLs

Process Virtualization

- the process of **virtualization** takes a *physical, limited* resource and creates the illusion of having *virtual, unlimited* copies of that resource
- the most fundamental abstraction provided by the OS to users is the **process**, or running instance of a **program**
 - a program is:
 - * *static*, an abstraction stored on disk as a **load module** with resolved references
 - * contains **headers**, code and data segments, **symbol table** for the linker
 - * but all addresses are relative, unloaded addresses
 - a process has different **segments** loaded into its address space:
 - * statically-sized **code segment** contains code read in from load module
 - read-only, executable-only, thus different processes can share the same code segment by mapping the addresses
 - * **data segment** containing heap, handles *initialized* global data as well as *dynamic* data
 - read-write, process private
 - can grow and shrink during process, grows upward
 - * **stack segment** handles procedure call stack frames (eg. local variables, invocation parameters, saved registers)
 - grows downward
 - * **stack overflow** occurs when stack and data segment meet, protects from data corruption
 - can also interpret a process as a virtual, private computer, or an *object*
 - * the **state** of a process should consistently, uniquely, characterize the process
 - * consists of the metadata, allocated memory, opened files, condition of an I/O operation, etc.

- in order to run many programs at once, the OS must *virtualize* the CPU
- OS uses a **time sharing** approach to virtualizing the CPU, as opposed to a **space sharing** approach (eg. for files)
- there are low-level **mechanisms** that help achieve this virtualization, eg. **context switching** that allows OS to switch between running programs on a CPU
- in addition, there are higher-level **policies** or decision algorithms used by the OS to choose which programs to run at a given time (*scheduling* policies)
- a process has an associated **machine state** or properties that it can read or write to at any given time. The machine state comprises of:
 - **memory** to store instructions and data, every process has an **address space**
 - * the address space is the *virtual memory addresses* reserved for a process (illusion of infinite memory)
 - **registers** that are used during execution
 - * some special registers include the **program counter** that indicates the next instruction, **stack pointer**, and **frame pointer**
 - **I/O information** for open persistent storage devices
 - other OS-related state information

Subroutine Stack Frames

- calling a subroutine:
 - *parameter passing* involves placing parameters into registers
 - *subroutine call* involves saving the *return* address on the stack, and transferring control to the entry point
 - *register saving* involves saving certain nonvolatile/callee-saved registers so that they can be restored
 - *space allocation* for local variables
- returning from a subroutine (symmetric steps):
 - place *return value*
 - pop *local* storage off stack
 - restoring *registers*
 - transfer control
- handling traps and interrupts:
 - **interrupts** inform software of an external event
 - **traps** are hardware instructions that inform software of an execution fault (type of interrupt)
 - similar to a procedure call, ie. have to transfer control, save state, restore

- state and resume process
- different from procedure call because *hardware*-initiated, so linkage conventions are defined by the hardware
 - * after event, computer state should be restored as if event never happened
- very expensive event to handle, since the CPU is moved to a privileged mode and new address space
- trap and interrupt *mechanism*:
 - a table associates a **program counter and processor status (PC/PS)** word pair with each possible interrupt/trap
 - when an event triggers an interrupt or trap:
 - * CPU uses exception number to index into table and load a new PC/PS onto the CPU stack
 - * exception continues at new PC address
 - *first level handler* saves registers, polls hardware for cause of exception, chooses and calls a *second level handler*
 - * on second handler termination:
 - first level handler restores registers, reloads PC/PS, resumes execution

Process Overview

- conceptual process *API*:
 - *create* - OS must provide method to create new processes
 - * may create a *blank* process with no initial state (Windows approach)
 - * or use the *calling* process as a template (UNIX approach)
 - this approach is useful when making processes with context from parent
 - * leads to *parent-child* process relationship
 - *destroy* - OS must provide method to destroy or kill processes, eg. with **signals**
 - * must clean up a terminating process:
 - reclaim resources
 - inform interprocess processes with signals
 - remove process descriptor
 - *wait* - wait for a process to stop running
 - *misc. control* - eg. suspending and resuming processes
 - *status* - retrieve status info for a process
- process *creation* consists of:
 - creating a new address space

- *loading* and *mapping* code and data into memory/address space of the program
 - * programs usually reside on **disk**, so the OS reads bytes from disk and places them in memory
 - * modern OSs use **lazy loading** to load data only when it is needed
- allocating and initializing the **stack** for the program (eg. with parameters, `argv`, `argc`)
- allocating the **heap** for dynamic memory
- initializing registers (PC, PS, SP)
- initializing I/O (eg. opening **file descriptors**)
- run program from its **entry point** (eg. `main`)
- in addition, processes may be loaded and *resumed* from a previous blocked state
 - in this case, registers must be loaded from the saved state
- **states** of a processes:
 - **running** - CPU is executing instructions for a process
 - **ready** - process is ready to run, but not currently executing
 - * when a process is *scheduled*, it moves from ready state to running
 - * when a process is *descheduled*, it moves from running state to ready
 - **blocked** - process has performed some operation that makes it unable to run until some other event takes place (eg. I/O request to disk)
 - * the OS recognizes when a running process becomes blocked, and will run a different process to maximize time sharing
 - **initial, final/zombie** (not yet cleaned up, allows **parent** process to check return code)
- the OS maintains key *data structures* or **process descriptors** to track the state of processes. These include:
 - **process/task table** for all ready or running processes, another list for blocked processes
 - * **process control block (PCB)** is a C structure maintained in Linux storing information for each process
 - used for saving the state of process and the registers of a process for **context switching**
 - eg. start and size of memory, process state (eg. scheduled, blocked) and ID, open files, CWD, context, parent, registers
 - * certain state of processes is additionally stored on a *per-process kernel stack*
 - retains the stack frames for in-progress OS system calls, and the state of interrupted processes so that the OS can return back to the process
 - must be separate from user stack for *security*, kernel stack used

for privileged operations

- must be per-process since different processes will experience *different* interrupts and system calls
- saves registers required for switching in and out of the kernel after interrupts, eg. PC, PS, SP, as well as user registers

UNIX Process API

Using `fork`:

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        /* child (new process) */
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {
        /* parent goes down this path (main) */
        printf("hello, I am parent of %d (pid:%d)\n",
               rc, (int) getpid());
    }
    return 0;
}
```

- `fork` creates an almost *exact* copy of the calling process
 - to OS, there are two programs running, both about to return from `fork`
 - thus, new **child** process starts running after call to `fork`, instead of from start of `main`
 - * child has a new address space
 - * child shares parent's *code segment*

- * a new *stack* is initialized to match the parent's
- * *data* initially points to the parent's original data
 - but when the child modifies the data segment, we need to set up a separate data segment copy for the child
 - copying large data segment can be expensive, so OS uses **copy-on-write** (lazy operation) when needed
 - copy-on-write occurs on a low granularity, eg. only copying and remapping specific page that is changed
- `fork` is non-deterministic, either child or parent will print first depending on the CPU **scheduler**
- new child has a copy of the address space, but the return of `fork` differs:
 - child receives return code of 0
 - parent receives new PID of child

Using `wait`:

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {
        int rc_wait = wait(NULL);
        printf("hello, I am parent of %d (rc_wait:%d) (pid:%d)\n",
            rc, rc_wait, (int) getpid());
    }
    return 0;
}
```


- `wait` waits for a child process to finish executing
 - returns PID of finished child process
 - this makes code snippet deterministic, child will always print before parent in this case

Using `exec`:

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/wait.h>

int main(int argc, char *argv[]){
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc");
        myargs[1] = strdup("p3.c");
        myargs[2] = NULL;
        execvp(myargs[0], myargs); /* counts words in p3.c */
        printf("this should not print");
    } else {
        int rc_wait = wait(NULL);
        printf("hello, I am parent of %d (rc_wait:%d) (pid:%d)\n",
            rc, rc_wait, (int) getpid());
    }
    return 0;
}
```

- `exec` allows us to fork a child of a *different* program

- does not *create* a new process, *transforms* currently running program (here, a forked child) into another running program
 - * OS loads new code and overwrites current code segment, reinitializes heap and stack, runs new program
- thus, a successful call to `exec` *never returns*
- separation of `fork` and `exec` have several advantages:
 - allows shell code to be run *after* `fork` and *before* `exec` and thus alter the environment of about-to-run program
 - allows for easy redirection
 - * eg. in order to redirect output to some file, after `fork`, close `STDOUT`, open the file, and then `exec`
 - piping can be achieved with the `pipe` system call
- **signals** and processes:
 - `kill` system call can send signals to a process or **process group**, eg. `SIGINT` for interrupt, `SIGSTP` for stop
 - processes can then use `signal` system call to catch signals

Process Mechanisms

Signals

- OS allows for processes to attach *event* callbacks
 - functions analogously to traps and interrupts, implemented and delivered to process by OS
 - eg. I/O devies and timers
- these *events* are defined by OS through many types of **signals**
 - processes can then choose to *ignore*, *handle*, or perform *default* action on certain signals

Direct Execution

- the **scheduler** is the component that actually determines which processes to run
- challenges associated with **virtualizing**, ie. **time sharing** the CPU:
 - maximizing **performance** with minimal *overhead*
 - * eg. minimal entering the OS, minimal use of the priveleged instruction set (no OS intervention)
 - handling processes while retaining *control*

An initial **direct execution protocol** without limits for maximum efficiency:

| OS | Program |
|--------------------------------|--|
| create entry for process list | |
| allocate memory for program | |
| load program into memory | |
| set up stack with argc/argv | |
| clear registers | |
| execute <code>call main</code> | run <code>main</code> |
| | execute <code>return</code> from <code>main</code> |
| free memory of process | |
| remove from process list | |

- this initial approach has some problems:
 - how do we ensure CPU doesn't do anything undesired or restricted, without reducing efficiency?
 - * occasional **traps** for syscalls
 - how do we efficiently switch processes in order to actually *virtualize* the CPU?
 - * occasional **timer interrupts** for time sharing

Restricted Operations

- some operations should be **restricted** to the OS, eg. I/O or accessing more system resources
 - eg. if any process could do I/O, all data protections would be lost
 - the solution is to introduce processing modes, a restricted *user mode* and an unrestricted *kernel mode* with elevated privileges
 - * kernel mode also has full access to hardware resources
 - OS provides an ABI to access these operations with **system calls** and **traps**
- some **exceptions** are routine that can be checked in programs:
 - end of file, arithmetic overflow, conversion errors
- however, sometimes will occur that aren't handled by the user
 - usually **asynchronous** exceptions such as segfaults, user abort, power failure
 - OS must handle these unhandled exceptions with a **trap** into the OS to perform restricted operations
- when a *user* program wants to perform a privileged operation, it can use a system call
 - system calls allow the kernel to expose important functionalities

- a **system call number** is associated with each syscall, this indirection is a form of **protection**
- user arguments/input are placed in ABI-specified registers, and must be validated by OS before performing the actual syscall in kernel mode
- to actually execute a system call:
 - process executes a syscall with the ABI's specifications
 - this causes a **trap** or exception that jumps into the kernel
 - hardware:
 - * raises the privilege level to kernel/supervisor mode so privileged operations can be performed
 - * uses exception number to index into a **trap vector table** to get the **program counter (PC)/program status (PS)** associated with the first handler of the exception
 - PS usually holds the return/error code from the syscall
 - * push PC/PS of process triggering trap to kernel stack
 - * load a new PC/PS onto the kernel stack for the trap handler
 - software continues at new PC address:
 - * **first level handler:**
 - saves registers
 - polls hardware for cause of exception
 - chooses and calls a **second level handler** from a **dispatch table**
 - * **second level handler:**
 - specifies the **trap gate**
 - actually handles the trap/syscall
 - on second handler termination:
 - * first level handler restores registers, reloads PC/PS, allows for resuming execution at next instruction
 - when finished, the OS calls a **return-from-trap** instruction that returns to user mode and reduces the privilege level
- a per-process **kernel stack**:
 - acts as a *call stack* for trap handlers and other privileged operation routines
 - * separate from user stack for *security* and *isolation*
 - allows execution of the user process to be *resumed*
 - must push PC/PS, flags, user registers, system call parameters onto kernel stack when entering OS
 - grows and shrinks alongside the syscall handler stack frames
- the kernel should carefully control which code executes on a trap
 - OS sets up a **trap table** at boot time that informs the hardware of the locations of **trap handlers** to call on a trap
 - note that the machine boots initially in kernel mode

An updated **limited execution protocol** to deal with system calls and traps:

| OS at boot | Hardware | Program |
|-----------------------|---|---------|
| initialize trap table | remember addresses of trap/syscall handlers | |

| OS at run (kernel mode) | Hardware | Program (user mode) |
|-------------------------------|--------------------------------|-------------------------------|
| create entry in process list | | |
| allocate memory for program | | |
| setup user stack with args | | |
| fill kernel stack with reg/PC | | |
| return-from-trap | restore regs from kernel stack | |
| | move to user mode | |
| | jump to <code>main</code> | run <code>main</code> |
| | | call syscall |
| | | trap into OS |
| | save regs to kernel stack | |
| | move to kernel mode | |
| | jump to trap handler | |
| handle trap/syscall | | |
| return-from-trap | restore regs from kernel stack | |
| | move to user mode | |
| | jump to PC after trap | |
| | | return from main |
| | | trap (via <code>exit</code>) |
| free memory of process | | |
| remove from process list | | |

Process Switching

- when a program is running on a CPU, the OS is *not* running
 - how can the OS *regain control* of the CPU so that it can switch between processes?
 - in a **cooperative** scheduling system, the OS simply *waits* for a program to make a syscall or `yield` in order to regain control
 - * this can lead to bugs with infinite loops or malicious programs

- in a **non-cooperative** scheduling system, a **timer interrupt** is used
 - * every so many milliseconds, an interrupt is raised automatically, and an **interrupt handler** in the OS runs
 - * this timer must be started on boot up
 - * the hardware must save the state of the program so that execution can resume on a return-from-trap
- once OS has control, the **scheduler** decides whether to continue running the current process (process A), or switch to a different one (process B) with a **context switch**
 - in a context switch:
 - * the *hardware* saves the *user* registers for A into kernel stack A so A can resume execution after interrupt
 - * the *OS* saves the *kernel* registers for A into memory in the process structure of A so A can resume execution after context switch
 - * the *OS* restores the *kernel* registers for B from memory in the process structure of B
 - * the *OS* switches from kernel stack A to kernel stack B by changing the stack pointer
 - * the *hardware* restores the *user* registers for B from kernel stack B
 - then, after return-from-trap, the system resumes execution of *another* process
 - context switching is *expensive*: have to enter the OS, switch OS context (stacks, address spaces), loss of caches
- to deal with the issue of **concurrency**, the OS may disable interrupts for a period of time, or use locking schemes

An updated **limited execution protocol** to deal with context switching:

| OS at boot | Hardware | Program |
|-----------------------|---|---------|
| initialize trap table | remember addresses of trap/syscall handlers | |
| start interrupt timer | start timer | |
| | interrupt CPU in X ms | |

| OS at run (kernel mode) | Hardware | Program (user mode) |
|-------------------------|---------------------------------------|---------------------|
| | | Process A |
| | timer interrupt | |
| | save regs(A) \rightarrow k-stack(A) | |
| | move to kernel mode | |
| | jump to trap handler | |

| OS at run (kernel mode) | Hardware | Program (user mode) |
|--|---|---------------------|
| handle trap call switch save regs(A) into proc. struct A restore regs(B) from proc. struct B switch to k-stack(B) return-from-trap | restore regs(B) \leftarrow k-stack(b) move to user mode jump to PC of Process B | Process B |

Process Scheduling

- because CPU is limited as a resource, the OS has to use a **scheduler** in order to schedule processes such that they have the illusion of having full access to the CPU's resources
 - scheduler has to choose which *ready* processes to run when:
 - * current process *yields* or traps to the OS
 - * *timer* interrupt occurs
 - * current process becomes *blocked* (eg. I/O)
 - *metrics* for performance are the **turnaround time**, **throughput**, **wait time**, **response time**, degree of **fairness**, achieving explicit **priority** goals, **real-time** scheduling
 - different scheduling *goals*:
 - * **time sharing** - fast response time for interactive programs, every user gets equal CPU share
 - * **batch** - maximize throughput, individual delays are unimportant
 - * **real-time** - critical operations must happen on time, non-critical operations may be deferred
- *ideal* throughput is impossible:
 - some overhead per dispatch
 - in general, want to reduce the overhead per dispatch (mechanism), as well as the number of dispatches (policy) so that performance *approaches* the ideal throughput
- response time *explodes* at a certain load:
 - finite queue sizes, requests or parts of requests may be *dropped* (infinite

- response time)
- dealing with **overloaded** systems:
 - continue service with *degraded* performance
 - maintain performance by rejecting work
 - resume normal service once load drops
 - *avoid* throughput dropping to zero or infinite response time
- in addition to lower level mechanisms associated with the process abstraction, OS also deals with high-level scheduling **policies** for processes
 - the OS policies should be implemented and chosen *independently* than the mechanisms (eg. dispatching and context switching)
 - scheduler can either be **preemptive** (interruptive) or **non-preemptive**
 - * *non-preemptive pros*: low overhead, high throughput, simple (fewer context switches)
 - * *non-preemptive cons*: poor response time, freeze at infinite loop bugs, not fair, difficult for real-time and priority scheduling
 - * *preemptive pros*: good response time, fair, good for real-time and priority scheduling
 - * *preemptive cons*: more complex, requires context switch mechanism, not as good throughputs, higher overhead
 - we will explore different scheduling algorithm approaches and how they fail when certain *assumptions* are *relaxed*:
 - * every job runs for the same amount of time
 - * all jobs arrive at the same time
 - * each job runs to completion once started
 - * all jobs only use the CPU
 - * run-time of each job is known
 - initially, aim for optimizing turnaround time (similar to throughput)
- **first in, first out (FIFO)** scheduling:
 - schedules jobs in the order that they arrive
 - highly variable delays
 - useful when response time is not important (*batch* programming), or *embedded* systems (brief processes, simple implementation)
 - effective until first assumption is relaxed and jobs no longer run for the same amount of time...
 - * issues with the **convoy effect**, where many low potential consumers may become queued *behind* a heavyweight resource consumer
- **shortest job first (SJF)** scheduling:
 - runs the shortest job first, next shortest, etc.
 - optimal until second assumption is relaxed and jobs no longer arrive at the same time...

- * shorter job could arrive while a job is still running
 - * ie. SJF is a **non-preemptive** scheduler that cannot interrupt jobs
- **shortest time-to-completion first (STCF)** scheduling:
 - also relax assumption three that jobs will run to completion
 - allow scheduler to use context switching and **preempt** jobs to run another job
 - AKA **preemptive shortest job first (PSJF)**
 - effective until we consider a new metric **response time**, the time for a job to be scheduled for the first time
 - * response time deals with interactivity for users
- **round robin (RR)** scheduling:
 - rather than running jobs to completion, run a job for a **time slice** before switching to the next job in the queue
 - * more fair CPU sharing and delays, good for interactive processes
 - length of time slice thus must be a multiple of the timer-interrupt period
 - tradeoff between smaller, faster time-slices and the overhead of more context switching, find good **amortization**
 - *however*, one of the worst policies in terms of turnaround time and number of context switches
- when considering other systems, exploit the **overlap** of operations:
 - eg. when a process becomes blocked waiting for the completion of an I/O request, schedule another process
 - involves treating each CPU *burst* as an individual job
 - *however*, for SJF and STCF, the run-time of each job is known

Feedback Priority Scheduling

- **multi-level feedback queue (MLFQ)** scheduling:
 - aims to optimize turnaround time *as well as* response time when assumption five is relaxed and the run-time of jobs aren't known
 - an example of a system that uses the *past* to predict the *future*
 - has a number of distinct **priority queues** with different priority levels
 - * (1) If the priority of A > priority of B, only A runs
 - * (2) If the priority of A = priority of B, A and B run in RR
 - * queues may have different time slices depending on priority:
 - shorter time slices for *foreground* high priority tasks, and long time slices for *background* low priority tasks
 - * can also have a queue dedicated to real-time tasks that run until completion
 - FIFO for low priority queue or real-time

- the scheduler will *vary* the priorities of a job based on its *observed behavior*
 - * eg. when a job repeatedly makes I/O requests to the keyboard, it will have its high priority *maintained* (interactive process)
 - * eg. when a job uses the CPU intensively for long periods of time, it will have its priority *reduced* (response time isn't important)
 - * (3) When a job enters the system, it is placed at the *highest* priority
 - * (4a) If a job uses an entire time slice, its priority is *reduced*
 - * (4b) If a job gives up CPU before time slice is up, it stays at the *same* priority
- when a new job comes along, the scheduler *assumes* it may be a short job with a high priority:
 - * if it is short, the job will run quickly and complete
 - * otherwise, the job will move down the queues and run in a batch-like process
 - * can also profile processes to estimate which queue to place them into
- issues with this initial implementation:
 - * with too many interactive jobs, they will consume *all* CPU time and long running jobs will be **starved**
 - * a program could *maliciously* issue an I/O operation right before the end of its time slice to *always* run at a high priority
 - * no mechanism for a CPU-bound job to transition to interactivity
- using **accounting**:
 - (4) Once a job uses up its time allotment at a given level, its priority is *reduced*
- using a **priority boost**:
 - in order to guarantee CPU-bound jobs will make progress against *starvation*, boost *all* jobs periodically
 - (5) After some time period S , move all jobs to the topmost queue
 - * S is a **voodoo constant**, if too high, starvation occurs, if too low, interact jobs would not get a proper share of the CPU
- involves many **parameters** for time slice length, number of queues, etc.
 - many implementations provide configuration files that can adjust these parameters
 - users can also give **advice** to the OS to modify scheduler behavior

Realtime Systems

- priority scheduling is a *best effort* approach
 - there are other systems whose correctness depend on certain *timing* requirements as well as *functionality*

- eg. space shuttle during reentry, reading sensor data at high speeds, playing media
- realtime systems are characterized by different metrics:
 - **timeliness** - how closely timing requirements are met
 - **predictability** - deviation in timeliness
- new realtime concepts:
 - **feasibility** - whether or not requirements for a task set can be met
 - **hard real-time** - strong requirements that tasks be run at specific intervals, not recoverable on failures
 - * dynamic behavior, unbounded loops should be avoided
 - * may have to disable interrupts and preemptive scheduling
 - **soft real-time** - good response time required, but recoverable on failures
- realtime characteristics that make scheduling *easier*:
 - task length will be known
 - **starvation** of low priority tasks is acceptable
 - work-load may be fixed
- *differences* between ordinary time-sharing:
 - **preemption** is no longer an optimal strategy:
 - * preempting running tasks will cause it to miss its deadline
 - * execution time is known, so there is little need for preemption
 - * real-time systems run fewer and simpler tasks, so code is not malicious or buggy (no infinite loops)

Inter-Process Communication

- data can be exchanged between processes *uni-directionally* or *bi-directionally*
- uni-directional byte streams are processing pipelines, where processes read from stdin and write to stdout
 - each program accepts a byte-stream input and produces a well defined byte-stream output
 - each program operates independently
- **pipes** are temporary files from **pipe** that are different from static files in the following ways:
 - reader does not get EoF until the write side of the pipe is closed
 - SIGPIPE from writing to a pipe without an open read end
 - file automatically closed when both ends are closed
- **named-pipes (FIFOs)** are a *persistent* pipe that can connect unrelated processes
 - not destroyed when I/O is finished

- writes from multiple writers can be interspersed
- no clean fail-overs on failed reads
- **mailboxes** are another inter-process mechanism
 - * not a byte-stream, but distinct, delivered messages
 - * every write has id from sender
 - * unprocessed messages saved on death of reader
- **general network connections** provide network communications
 - can use byte streams (TCP) or datagrams (UDP)
 - much more complexity due to online network
 - high latencies, limited throughput
- **shared memory** is the highest performance communication:
 - much faster than other models
 - processes must run on the same memory bus
 - race conditions, synchronization issues
 - no authentication
- **out-of-band** signals should supersede queued or buffered data
 - locally, could set up a handler that flushes all buffered data upon receiving an out-of-band signal on a different channel
 - with network services, open multiple communication channels
 - * server must periodically poll the out-of-band channel

Memory Virtualization

- addresses are *abstracted* to programs as **virtual addresses**
 - allows for *ease of use* for programmers, and **isolation** and **protection**
- an **address space** is a *virtual* abstraction of physical memory
 - virtual address independent from physical address
 - contains all of the memory state of a running program (code, stack, and heap segments)
 - by convention, stack and heap at opposite ends, grow in opposite directions
 - the program is not in *contiguous* physical memory like the address space, but loaded at *arbitrary* physical addresses
 - * eg. printing pointers in C prints virtual addresses
 - every process can have an address space of immense size
 - * supported using **dynamic paging** and **swapping**
- memory virtualization goals include:
 - *transparency*, ie. an invisible implementation by the OS
 - *efficient* virtualization through hardware support

- *protecting* and processes from one another through isolation

UNIX Memory API

- **heap** memory, as opposed to **stack** or *automatic* memory, is explicitly handled by the programmer
- `malloc` dynamically allocates space on the heap
 - is a library call that uses system calls such as `brk` or `sbrk`
 - `sizeof` is a *compile-time* operator
 - `free` frees heap memory
- `calloc`, `realloc`

```
#include <stdlib.h>

double *d = (double *) malloc(sizeof(double));
free(d);
```

- common errors:
 - failure to allocate memory (eg. `strcpy` into a unallocated pointer) often leads to a **segmentation fault**
 - **buffer overflow**, or not allocating enough memory can have nondeterministic behavior
 - **uninitialized read**, or not initializing allocated memory
 - **memory leak**
 - **dangling pointer**
 - **double free**
 - **invalid free**

Memory Mechanisms

General Partition Strategies

- **fixed partition**:
 - preallocate for a certain number of processes
 - useful when exact memory needs are known
 - partition sizes are fixed
 - using only *contiguous* physical addresses

- *pros*:
 - * simple implementation
 - * allocation/deallocation cheap and easy
- *cons*:
 - * inflexible, limits number of processes and their memory usage
 - * can't share memory
 - * **internal fragmentation** - wasted space inside *fixed* blocks
- **dynamic partition**:
 - similar to fixed, except *variable* sized blocks
 - each partition is contiguous
 - still using physical addresses
 - *pros*:
 - * sharable partitions
 - * process can use multiple partitions, with different sizes
 - *cons*:
 - * still not *expandable*, may not be space nearby
 - * still not *relocatable*, pointers will be incorrect, partitions tied to address range
 - * still subject to fragmentation
 - * not as large as virtual address space
- can use a **free list** to track unallocated memory
 - allow processes to use incontiguous, *variable* sized partitions
 - each element in the list has a metadata **descriptor** with data such as length, free or not, and a pointer to the next chunk
 - to *carve* a chunk:
 - * reduce the length, create a new header for leftover chunk, connect new chunk to list, and mark chunk as used
 - eliminates internal fragmentation, since process can use smaller sized chunks as needed
 - however, over time, leads to **external segmentation** - useless, small left-over chunks
 - different free space management strategies help counteract external segmentation:
 - * different free space algorithms, eg. first-fit, next-fit
 - * **coalesce** adjacent free memory chunks together
 - * is it possible to relocate free memory and **compact** it together?
 - compaction requires relocation
 - relocation requires **address translations** and **virtual memory**

Address Translation

- **hardware-based address translation** is a generic, hardware technique that extends the **limited direct execution** model
 - hardware performs an address translation on every memory reference, ie. *interposes* each memory access
 - * the **memory management unit (MMU)** is the CPU component dealing with memory virtualization
 - OS takes care of **managing memory**
 - address translation allows for easier relocating of memory
 - * without virtual memory, whenever memory is moved, would have to update all of a process's pointers and memory references
- **dynamic relocation** or *base-and-bounds* technique:
 - uses a **base** and **bounds** register in the CPU (different per process)
 - when a program starts running, the OS decides where in physical memory to load it and sets the base register to that value
 - every memory reference (virtual address) from the process gets *translated* by the CPU by adding the base register to produce a physical address
 - * occurs at runtime, *interposed*, little hardware logic required
 - if a virtual address is greater than the bounds register, an exception will be raised
 - * *limits* and *protects* address spaces
 - *hardware*:
 - * provides extra registers in the MMU and privileged instructions to modify these registers
 - * provides mechanisms for raising exceptions and registering handlers
 - OS:
 - * allocate memory for new processes using a **free list** (data structure holding free slots in physical memory)
 - * cleans up after a process ends by updating the free list and deallocating memory
 - * *save-and-restore* base-bounds pair registers using the PCB when context switching
 - * install exception handlers at boot time (along with other handlers eg. syscall handlers)
 - acts as an *extension* of LDE, where address translating is interposed during direct execution, and OS only intervenes when process misbehaves
 - allows for address space to be *relocated* when a process is stopped by copying between locations and updating the saved base register
 - **software relocation** is an alternative where the loader rewrites all instruc-

- tions by adjusting the addresses
 - * provides no protection, and difficult to relocate address spaces
- dynamic relocation can lead to **internal fragmentation**, where the space inside the allocated unit of a process is wasted, since its stack and heap are small
 - restricts address spaces in fixed-size slots
 - cannot run programs where the entire address space doesn't fit into memory
 - issue compounded with larger, **sparse address spaces**, eg. 32-bit, 4 GB address spaces

Segmentation

- instead of having a single base-bounds pair in the MMU, instead have a base-bounds pair for every **segment** in the address space
 - use segments as the *unit* of relocation
 - each segment is already a *continugous* portion of address space
 - * this is a **course-grained** segmentation, **fine-grained** segmentation involves more smaller segments, usually with a **segment table**
 - allows OS to place segments in different parts of physical memory *independently* and fill *unused fragments*
 - * more flexibility when allocating for processes with large, sparse address spaces
 - * allows specific segments, eg. code segment, to be shared between processes
 - now, during address translation, hardware must consider the **offset in the specific segment** the address or instruction belongs to, and add the offset to the base register of the segment
 - * eg. to read an address in the heap at virtual address 4200, offset = 4200 - virtual address of start of heap, physical address = offset + bounds(heap)
 - to read into the stack that *grows backwards*, the MMU uses a register bit for all segments to keep track of which way the segment grows
 - to allow for sharing of memory segments, the MMU uses register **protection bits**
 - * eg. code is read-execute, heap and stack are read-write
 - considering an illegal address beyond the end of a segment leads to a **segmentation fault**
- matching an address to its segment and segment base-bounds register:
 - the **explicit** approach is to *slice* the address space into segments based on the *top few bits* of its virtual address

- * eg. if the first two bits of a virtual address is 00, the hardware will use the code base and bounds pair and the remaining bits as the offset from the segment start
- the **implicit** approach is to use hardware that examines how the address was formed (eg. program counter, stack pointer, otherwise)

Example address translation process:

```
Segment = (VirtualAddr & SEG_MASK) >> SEG_SHIFT
Offset = VirtualAddr & OFFSET_MASK
if (Offset >= Bounds[Segment])
    RaiseException(SegmentationFault)
else
    PhysAddr = Base[Segment] + Offset
    Register = AccessMemory(PhysAddr)
```

- segmentation raises some issues that OS must deal with:
 - on a context switch, segment registers must be saved and restored (each process has a virtual address space)
 - OS must find space in physical memory for new address spaces:
 - * every segment can now be a different size
 - issue of **external fragmentation**, when physical memory becomes full of small holes of free space as segments change in size
 - * periodic **defragmentation**: OS can **compact** physical memory by rearranging segments contiguously and updating base registers
 - copying can be very expensive (especially for some types of disks)
 - * OS can coalesce as much as possible when free segments are contiguous
 - frequent allocation/deallocations the opportunity to coalesce
 - * another approach is to use a free-list **management algorithm** (many algorithms have been used)
 - avoid creating small fragments
 - recombine smaller fragments
 - still not segmented enough, eg. segments *themselves* may be sparse and largely empty (entire heap or stack mostly empty)

Free-Space Management

- managing free space can be simple with fixed sized chunks

- more difficult with variable-sized units, eg. when using segmentation or allocation libraries
 - * supporting variable-sized blocks counteracts **internal fragmentation**
 - * but leads to **external fragmentation** as the number of small, useless left-over chunks increases
- note that the **allocator** in a allocation library cannot utilize **compaction** to combat external fragmentation
 - * compactations are expensive, and run periodically
- allocator mechanisms:
 - uses a **free list** to reference free chunks of space on heap
 - * eg. would contain the starting address and length for each chunk in a linked list
 - * difficult to **scale** the performance of a linked list, special types of trees may be a better data structure
 - on a small request, **split** a chunk into two and update the chunk's length accordingly in the free list
 - on a memory free, **coalesce** multiple, contiguous chunks together into a single new chunk in the free list
- allocators store metadata for allocated memory in a **header** block immediately before allocated chunk:
 - could store chunk size, additional pointers, magic number for *integrity checking*
 - every allocation of N bytes will require enough space for N + K bytes for the header
- the free list must be *embedded* in the free space itself:
 - free list node minimally holds free chunk size, and a pointer to the next chunk
 - on an allocation, the free chunk is split in two:
 - * one chunk large enough for the request and header
 - * remaining free chunk with an updated size in the free list node
 - on a memory free,
 - * the allocator uses the size in the chunk header to add the free chunk back into the free list,
 - * adds a pointer to the next free chunk's node, and redirects the head pointer (requires coalescing and merging to clean up)
- usually, allocator starts with a smaller heap, and uses **sbrk** to ask OS to grow the

heap

- free space *strategies*:
 - want to minimize external fragmentation by avoiding smaller fragments
 - **best fit**, ie. smallest fit
 - * find the smallest possible block, waste minimal space
 - * quickly creates small fragments
 - * may involve expensive exhaustive search
 - **worst fit**
 - * find the largest possible block, leaving a large free chunk remaining
 - * creates larger fragments, for longer
 - * may involve expensive exhaustive search
 - **first fit**
 - * find first block that fits the request
 - * fast, but pollutes free list with many small objects
 - searches get longer over time
 - * could use **address-based ordering** in the free list to help coalesce
 - **next fit**
 - * maintains a pointer to where allocator was last checking for free space
 - guess pointer acts as a *lazy* cache
 - * spreads out searches more uniformly
 - * shorter searches
- these strategies combat external fragmentation, but *carving* and *coalescing* is expensive when allocating memory
 - can we minimize these actions?
 - **segregated lists** AKA **buffer pool**
 - * maintain several lists specifically dedicated to a few popular-sized, special requests
 - * with a uniform size of requests, allocation is more efficient, and fragmentation is eliminated
 - * need to balance how much to reserve in the pools
 - * the **slab allocator** uses **object caches** for kernel objects that are likely to be requested frequently (inodes, locks, etc.)
 - requests slabs of memory when the cache is running out of space
 - also keeps objects pre-initialized for even faster performance
 - **buddy allocation**
 - * on a request, recursively divides free space in *half* until a small enough block is created
 - * allows for extremely *fast* recursive coalescing, can just check if imme-

- diate “buddies” are both free and coalesce them
 - because of the recursive division, the address of buddies differ only by a bit (easy arithmetic)
 - * suffers from internal fragmentation (fixed powers of 2 sized chunks)
- allocator decides when an allocated resources should be returned to the pool
 - eg. after `close`, `free`, `delete`, `exit`, or returning from a subroutine
 - if a resource is *sharable* (eg. open file or shared memory segment), resource manager must maintain a count for each resource and free the resource only when the count drops to 0
 - however, keeping track of references to a resource is not always *practical*:
 - * some languages support copying references without using OS syscalls
 - * some languages don’t require programmers to explicitly free memory (Java, Python)
 - * some resources may be allocated and freed so often that keeping track of them becomes a significant overhead
- an alternative to count based freeing is **garbage collection**:
 - resources are allocated and *never* explicitly freed
 - only when pool of available resources becomes small does garbage collection occur:
 - * start with a list of original resources
 - * scan to find reachable resources by *chasing* pointers
 - * remove from original list if reachable
 - * free anything still in the original list (unreachable memory)
 - however, must be able to *identify* all *active* resource references
 - * language must *mark* resource references so they can easily be identified on the heap
 - * thus leads to an overhead when program must *pause* for garbage collection
 - * can be mitigated with progressive background garbage collectors

Paging

- rather than separate memory into *variable* sized pieces in the **segmentation** approach, separate memory into smaller, *fixed* sized chunks called **pages** or **page frames**
 - fixes **external segmentation** caused with segmentation:
 - * paging eliminates the requirement of contiguity
 - * pages themselves are never *carved* up, granularity is fixed

- no small, unused memory fragments
- fixes **internal fragmentation** to a degree:
 - * if the page frame is relatively small, internal fragmentation averages only half a page
- physical memory becomes an array of fixed-sized slots called **page frames**
- virtual memory (address spaces) is also virtualized with virtual pages
- pages can still be shared between virtual addresses
- allows for *flexibility* in abstracting the address space, no more need to keep track of which direction a segment grows
- provides *simplicity* when allocating space for processes from the free list (fixed sized pages)
- a *per-process* **page table** records where each virtual page of address space is placed in physical memory in a **page table entry (PTE)**
 - ie. stores address translations for each virtual page (replaces base-bounds registers)
 - simplest implementation is a **linear page table** or array
 - every virtual address can be *translated* by splitting it into components:
 - * the **virtual page number (VPN)** indicates which virtual page the address resides only
 - number of bits depends on how many pages in the address space
 - can replace VPN with the **physical frame number (PFN)** to generate the actual physical address by indexing into page table
 - * the **offset** indicates the offset within the page
 - stays consistent through address translation
 - also stores meta data such as:
 - * **valid bit** that is important for marking unused pages as invalid (no physical frame allocation required)
 - * **protection bit** indicating protection
 - * **present bit** indicates whether page is in memory or disk (required for page swapping)
 - * **dirty bit** if page is modified
 - * **reference bit** if page has been accessed
- page tables can become very large
 - aren't stored on-chip, but in physical memory
- since page tables are process specific:
 - need to load pointer to new page table on context switch
 - need to flush previously cached entries
- however for every memory reference, paging requires an *additional* memory reference in order to first fetch the translation from the page table
 - this can slow down the process by half or more

- thus the current iteration of paging can cause significant *slowdown* and *memory usage*
 - * note that every *instruction fetch* also generates two memory references, one to the page table the instruction is in and then the instruction itself

Example memory access with paging:

```
VPN = (VirtualAddr & VPN_MASK) >> SHIFT // shift by size of offset
PTEAddr = PTBaseReg + (VPN * sizeof(PTE))
PTE = AccessMemory(PTEAddr)

if (!PTE.Valid)
    RaiseException(SEGMENTATION_FAULT)
else if (!CanAccess(PTE.ProtectBits))
    RaiseException(PROTECTION_FAULT)
else
    offset = VirtualAddr & OFFSET_MASK
    PhysAddr = (PTE.PFN << SHIFT) | offset
    Register = AccessMemory(PhysAddr)
```

Translation Lookaside Buffer (TLB)

- paging makes address translation slower with an extra required memory reference
 - a **Translation Lookaside Buffer (TLB)** is part of the MMU, and is a hardware **cache** of popular address translations
 - on every virtual memory reference, hardware first checks if the TLB contains the translation
 - * if so, translation can be quickly performed without referencing the page table
 - in the common case, translations will be in the cache, and little overhead will be added
 - * want to *avoid* TLB misses as much as possible
 - * performance of program is thus as if memory isn't virtualized at all
- **caching** in general depends on two principles:
 - **temporal locality** wherean instruction or data that has been recently accessed will be referenced again soon in the future (loop variables or loop

- instructions)
 - **spatial locality** where programs access memories near each other repeatedly (traversing an array)
 - caches are generally small but fast
 - * want to minimize the **miss rate** and maximize **hit rate**
 - types of cache *misses*:
 - * a **compulsory miss** occurs because cache is empty to start upon first reference
 - * a **capacity miss** occurs because the cache ran out of space and had to evict
 - * a **conflict miss** occurs in hardware due to limits on items in a hardware cache
- TLB is usually **fully associative**, ie. one to one mapping between VPN and TLB entries
 - entry contains VPN, PFN, other bits such as a **valid bit**, **protection bits**, **ASID**, **dirty bit**, **global bit**
 - valid bit for entry indicates if entry contains a valid translation
 - * note that the valid bit for the page table indicates the page has not been allocated for the process
- address translation with TLB:
 - use the VPN to check if TLB holds translation
 - if so, **TLB hit**:
 - * PFN can be found from relevant TLB entry
 - otherwise, **TLB miss**:
 - * must go through page table for PFN, and update TLB with the PFN
 - * once TLB is updated, hardware retries the translation

Example memory access with TLB:

```

VPN = (VirtualAddr & VPN_MASK) >> SHIFT
(Success, TlbEntry) = TLB_Lookup(VPN)
if (Success) // TLB Hit
    if (CanAccess(TlbEntry.ProtectBits))
        Offset = VirtualAddr & OFFSET_MASK
        PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
        Register = AccessMemory(PhysAddr)
    else
        RaiseException(PROTECTION_FAULT)
else // TLB Miss

```

```
PTEAddr = PTBaseReg + (VPN * sizeof(PTE))
PTE = AccessMemory(PTEAddr)
if (!PTE.Valid)
    RaiseException(SEGMENTATION_FAULT)
else if (!CanAccess(PTE.ProtectBits))
    RaiseException(PROTECTION_FAULT)
else
    TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
    RetryInstruction()
```

- handling the TLB miss can either be done by hardware or software:
 - with **CISC**, hardware handles TLB miss entirely
 - * hardware needs a **page table base register**
 - with **RISC**, software handles the TLB miss
 - * hardware raises an exception, and a OS trap handler updates the TLB and returns from the trap
 - note that this return-from-trap must fall back to the original instruction so that it can be *retried* by the hardware
 - must ensure that no infinite chain of TLB misses occurs, so TLB handlers usually stored in permanent physical memory or permanent translation (**wired** translation) slots
- when context switching, have to somehow clear the TLB since every process has a unique virtual to physical set of translations
 - can **flush** TLB on context switches, eg. specifically when PTBR is changed
 - * however, every context switch will start with many TLB misses
 - can use an **address space identifier (ASID)** in the TLB entries that is process specific
 - when process share physical pages, two TLB entries simply map to the same PFN
 - * reduces physical pages in use and lowers overhead
- an issue is **cache replacement**:
 - a common approach is to evict the **least recently used (LRU)** entry in TLB
 - another is to use **random** eviction
- TLB is not a perfect solution:
 - if the number of pages a program frequently accesses exceeds the number of pages in the TLB, there will be many TLB misses
 - * known as exceeding the **TLB coverage**
 - * one solution is to support larger page sizes
 - TLB can become bottlenecked when using **physically indexed caches**

- * translations must take place *before* cache access
- * one solution is to use a **virtually indexed cache**

Swapping

- another level in the **memory hierarchy** is the **disk**
 - pages no longer all reside in physical memory, for very large address spaces, the OS needs to stash away unused parts of these spaces
 - * thus parts of the disk is reserved for swapping and known as **swap space**
 - * OS can swap pages in and out of disk in a page-sized granularity
 - * OS must also remember the **disk address** of a page
 - disk is *larger* and *slower* than physical memory
 - can also use *demand* paging, which only swaps in pages when they are used
 - * because of locality, demand paging is more efficient than swapping in all pages for a process at once
- disk allows for an even larger abstraction of memory, but requires more *machinery* for address translations:
 - when hardware checks the PTE, the page may *not be present* in physical memory
 - * stored in a **present bit**
 - if the page is present, the translation can proceed as usual
 - otherwise, this is a **page fault** or **page miss**, the page is in disk
 - * page faults never crash, only slow a program down
- on a page fault the OS uses the **page fault handler** software (even for hardware-managed TLBs):
 - needs to know the disk address, which is additionally stored in the page table
 - look in PTE for disk address, fetches page into memory from disk
 - * when I/O request to disk, process becomes blocked
 - update page table to mark page as present
 - update PFN for newly-fetched page in memory
 - backup PC to retry the instruction (could still lead to TLB miss, etc.)
- if memory is full, OS may have to **page out** pages to make room:
 - paging and swapping is handled by the **page-replacement policy**
 - OS may also proactively replace pages to maintain between a **low watermark** and **high watermark** number of pages
 - * this background replacement thread is called a **swap daemon** or **page daemon**

- different systems also **cluster** pages together when writing to the swap space, increasing disk efficiency

Example page-fault exception (hardware):

```
VPN = (VirtualAddr & VPN_MASK) >> SHIFT
(Success, TlbEntry) = TLB_Lookup(VPN)
if (Success) // TLB Hit
    if (CanAccess(TlbEntry.ProtectBits))
        Offset = VirtualAddr & OFFSET_MASK
        PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
        Register = AccessMemory(PhysAddr)
    else
        RaiseException(PROTECTION_FAULT)
else // TLB Miss
    PTEAddr = PTBaseReg + (VPN * sizeof(PTE))
    PTE = AccessMemory(PTEAddr)
    if (!PTE.Valid)
        RaiseException(SEGMENTATION_FAULT)
    else
        if (!CanAccess(PTE.ProtectBits))
            RaiseException(PROTECTION_FAULT)
        else if (PTE.Present)
            TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
            RetryInstruction()
        else if (!PTE.Present)
            RaiseException(PAGE_FAULT)
```

Example page-fault handling (software):

```
PFN = FindFreePhysPage()
if (PFN == -1)
    PFN = EvictPage()
DiskRead(PTE.DiskAddr, PFN) // sleep, waiting for I/O
PTE.Present = True
PTE.PFN = PFN
```

```
RetryInstruction()
```

Swapping Policies

- under **memory pressure**, OS must use a **replacement policy** to *evict* pages from main memory out to disk
 - can't control which pages are read in (demand paging), but we can choose which to kick out
 - the *optimal* policy is to replace the page that will be accessed *furthest in the future*
 - * impossible to actually implement, serves as a comparison policy
- **first-in-first-out (FIFO)** policy:
 - simple to implement
 - doesn't understand the importance of pages
 - * even if a page has been accessed more often, it may still be kicked out if it was the first one brought in
 - can lead to **Belady's Anomaly** (increasing cache size increases miss rate) because it does not obey the **stack property**, where a cache of size $N+1$ naturally includes the contents of a cache of size N
- **random** policy:
 - simple to implement
 - not intelligent in evicting pages
 - literally random performance
- **least-recently-used (LRU)** policy:
 - use *history* to guide decisions, *approximate* future behavior
 - more intelligent evicting, closer to optimal
 - **least-frequently-used (LFU)** also used
 - to support context switching, note that *per-process* LRU should be used instead of *global* LRU
 - implementing LRU requires updating some data structure on *every* page access or memory reference
 - * could have hardware support to update the time field in memory on each access (lower overhead)
 - * but scanning all time fields when evicting a page and holding so many time fields is extremely expensive
 - * could even lead to extra page faults in a purely software implementation (no saved time field)
 - * is there a way to *approximate* LRU?
- **approximating LRU**:
 - requires hardware support in the form of a **use bit** or **reference bit**

- * one bit per page of system, stored in the MMU
 - * whenever a page is referenced, use bit is set by MMU to 1
- using the **clock algorithm**:
 - * consider all pages are arranged circularly
 - * check if the currently pointed page has a use bit of 1 or 0
 - * if 1, clear use bit to 0, and increment pointer
 - * if 0, this page has not been recently used, so it can be evicted
 - * search continues at the pointer on the next eviction
 - * on worst case, will check all pages in the system
 - * the guess pointer acts as a recency approximation
 - if the rate of access on the page is faster than the clock *hand*, it has a lower chance of being evicted
- *modified* or *dirty* pages are expensive to evict, since they must be written back to disk
 - * thus some systems prefer to evict clean pages
 - * hardware should include a **dirty bit** to support this behavior
- to use the clock algorithm with per-process LRU:
 - * hardware also needs to maintain the owning process of a page and accumulated CPU time of each process
- almost as good performance as true LRU
- performance based on *workload*:
 - workload with *no-locality*:
 - * LRU, FIFO, and random all perform the same
 - * no locality to exploit
 - 80-20 workload (80% hot reference, 20% cold):
 - * LRU performs near optimal
 - *looping workload* ($N+1$ accesses, N cache size, in a loop):
 - * worst case for LRU and FIFO
 - consistently accessing older pages
 - * random performs near optimal
- OS also uses a **page selection** policy:
 - determines *when* to bring a page into memory
 - usually **demand paging**, or paged into memory when accessed
 - OS can also **prefetch**, but only on reasonable chance of success
- OS also has a policy to deal with **thrashing**, when memory is oversubscribed and constant paging occurs:
 - a **working set** size is an optimal number of pages for a process such that:
 - * increasing the number of allocated frames make little difference in performance
 - * decreasing the number of allocated frames decreases performance

greatly

- some OS use **admission control** to terminate a subset of running processes
 - * hopes that the reduced set of processes' **working sets** do not thrash the system
- other OS run an **out-of-memory killer** to kill the most intensive process

Concurrency

- in **multi-threaded** applications, each **thread** runs independently but access memory *shared* with other threads
 - ie. has *more than one* point of execution, or multiple PCs
 - * but *share* the same address space and data
 - issue if these shared resources aren't *coordinated* between threads
 - * can lead to **inndeterministic** results
 - multithreading benefits:
 - * allows for **parallelism** on multiple CPUs
 - * enables **overlap** of I/O with other operations *within* a single program
 - * easy to *share* data
 - OS must support primitives such as **locks** and **condition variables**
- **thread abstraction**:
 - each thread needs its own private set of registers
 - one *independent* stack per thread, ie. **thread-local storage**
 - switching threads thus requires a context switch (save and retore registers)
 - * state is stored in a **thread control block (TCB)**
 - however, address space remains the same
- issue when reading and writing to shared variables due to *uncontrolled* scheduling:
 - leads to **race condition** or a **data race**, where multiple executing threads enter a critical section at the same time
 - * race conditions lead to **indeterministic** programs, where results depend on the timing execution
 - * the piece of code where threads access a shared resource is a **critical section**
 - eg. incrementing a variable is not **atomic**, a read and write occurs in sequence
 - * one thread reads a variable, and a context switch occurs *immediately* before the subsequent write
 - * the next thread reads the *unincremented* variable, and writes the vari-

- able incremented by one
 - * the first thread writes the *original* value incremented by only one
 - * variable appears to be incremented just once, not twice
- race condition solutions:
 - have a hardware instruction that read and writes **atomically**
 - * for more general cases, such instructions don't exist
 - have hardware provided **synchronization primitives**
 - * ie. **mutual exclusion primitives**
 - * also need mechanisms to sleep and wake threads while awaiting I/O blocks

Locks

- **locks** are used around critical sections in order to ensure they are executed as an *atomic* instruction
 - AKA **mutex**, provides *mutual exclusion*
 - after being declared and initialized, locks start out **available**
 - exactly one thread can **acquire** a lock at a time
 - when a thread calls **lock** when another thread owns the lock in question, the function will not return until the owner calls **unlock** on the lock
 - note that there can be different locks for different critical sections
- when implementing locks, have to consider:
 - *mutual exclusion*: does the lock work?
 - *fairness*: does every thread waiting for a lock have the same chance to acquire it?
 - *performance*: is there significant overhead?
- interrupt **masking**:
 - an initial implementation involved simply disabling interrupts during a lock
 - many issues:
 - * process may maliciously use locks to exploit CPU
 - * fails with multiple CPUs
 - * interrupts can be lost (eg. I/O completion)
 - * requires an expensive, privileged instruction
- simple load/store attempt:
 - have a simple *flag* variable that is set to 1 on a lock
 - when another thread tries to lock the flag with a value of 1, **spin-wait** until value becomes 0
 - to unlock the flag, set it to 0
 - issues:

- * does not guarantee mutual exclusion!
 - reading and setting the flag is *still* not atomic, an interrupt can occur
- * spin-waiting is expensive
- spin-locks with *hardware support*:
 - need some hardware support for a **test-and-set** operation:
 - * an **atomic** instruction that returns sets a value and returns its previous value
 - use the same load/store implementation with test-and-set
 - the hardware guaranteed atomicity allows this lock to function correctly
 - issues:
 - * no guarantee of fairness, eg. a thread may spin forever
 - * heavy performance overhead, especially with only one CPU, eg. scheduler only schedules blocked threads
- other useful atomic hardware primitives:
 - **compare-and-swap** only updates a value if it has an expected value
 - **load-linked** is a typically load instruction
 - **store-conditional** only updates a value if no intervening store has occurred since its address was load-linked
 - **fetch-and-add** increments and returns a value atomically
 - * used in **ticket locks** that guarantee all threads progress
- issues with spin locks:
 - thread may spin-wait until an interrupt goes off as it waits for a lock
 - **priority inversion** may occur where a higher priority, *scheduled* thread is stuck waiting for a lower priority, *unscheduled* thread to give up its lock
- how to minimize spinning?
 - simply **yield** to the OS
 - * this works well with fewer threads, but with more threads, spinning threads may just *continuously* yield to one another (round robin)
 - * does not address *starvation* and fairness
 - instead use **queues** and sleeping:
 - * a queue avoids starvation
 - * using spin-waiting only *around* the lock itself, so the time spent spinning is limited to few lock and unlock related instructions
- **two phase locks** are an example of a *hybrid* approach with both a spin and a sleep phase
 - since spinning can be useful if lock is about to be released

Lock example with queues and sleeping:

```
typedef struct __lock_t {
    int flag;
    int guard;
    queue_t *q;
} lock_t;

void lock_init(lock_t *m) {
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}

void lock(lock_t *m) {
    while (TestAndSet(&m->guard, 1) == 1)
        ; // spin to acquire guard lock
    if (m->flag == 0) {
        m->flag = 1; // acquire lock itself
        m->guard = 0;
    }
    else {
        queue_add(m->q, gettid());

        // precaution against wakeup/waiting race:
        // if interrupt occurs and other thread releases the lock,
        // we don't want this thread to sleep forever.

        // setpark indicates thread is about to sleep, and if an interrupt occurs
        // and unparks before parks occurs, park immediately returns.
        setpark();

        m->guard = 0;
        park(); // put calling thread to sleep
    }
}
```



```

void unlock(lock_t *m) {
    while (TestAndSet(&m->guard, 1) == 1)
        ; // spin to acquire guard lock
    if (queue_empty(m->q))
        m->flag = 0; // let go of lock
    else
        unpark(queue_remove(m->q))
        // lock is not set to 0,
        // since the next thread does not hold guard lock anymore
        // ie. passing on the lock to the next thread
    m->guard = 0;
}

```

Locks with Data Structures

- concurrent counter:
 - to make a counter **thread-safe**, simply wrap each increment and read between a lock and unlock
 - expensive performance cost, using multiple threads makes scaled operations much slower
 - want **perfect scaling**, where threads complete just as quickly as the single thread
 - an approach is **approximate counting**, where each CPU maintains a *local* counter, and once a certain threshold on a local counter is met, a *global* counter is incremented by the local counter
 - * all operations have locks, but local counters won't be in contention with one another
 - * scales well, but the global counter is *inaccurate* and approximate
- concurrent linked lists:
 - to make linked lists thread-safe, make a *big* lock for the list, and surround critical sections of operations with locks
 - * make sure to surround the minimal, *actual* critical section, eg. `malloc` for a new node should be outside of the lock in case it fails
 - does not scale well
 - can use **hand-over-hand locking**, where each individual node has its own lock
 - * when traversing, code grabs next node's lock and releases the current node's lock

- * still much overhead for so many locks
- concurrent queues:
 - to make queues thread-safe, make a *big* lock for the queue
 - can also use two locks for head and tail of the queue
 - * allows for more concurrent operations
- concurrent hash table:
 - can treat hash table as an array of concurrent linked lists
 - thus, uses an individual lock for every bucket
 - * allows for more concurrent operations, scales well

Condition Variables

- how to allow a thread to check if a **condition** is true before continuing
 - eg. parent checking whether a child thread has completed
 - a simple implementation would have the parent spin-wait until a shared variable changes value
- threads can wait on a **condition variable** and be **signaled** to continue
 - use `wait` and `signal` in UNIX, used in conjunction with a state variable and lock
 - * the lock should be held when calling signal or wait
 - * sleeping, waking, and locking is built around the variable
 - without a state variable:
 - * if child runs first, parent will be stuck spin-waiting, ie. there is no state variable to record the child's completion
 - without using a lock:
 - * race conditions will occur when reading/writing to the state variable
- **covering conditions** are conditions where a thread should be woken up conservatively, regardless of the cost that too many threads are woken
 - eg. a memory allocation library that does not know which threads to signal when a certain amount of memory is freed
- the **producer/consumer** or **bounded buffer** problem:
 - multiple producer threads generate data in a buffer, consumers consume data from the buffer, eg. piping I/O
- issues in the initial example below:
 - **Mesa semantics**: after a producer wakes a consumer, but before the consumer runs, the bounded buffer is changed by another consumer
 - * possible because signaling a thread simply *wakes* it up, but this is only a *hint* that the shared state may have changed
 - in reality, when the woken thread runs, the state may not be as desired

- * can fix by replacing **if** with a **while**
 - when the woken thread runs, it rechecks the state
- all threads may end up asleep if a producer is signaled instead of a consumer and vice versa
 - * need to use two conditions, so consumers don't signal consumers and producers don't signal producers

Bounded buffer example:

```
int buffer;
int count = 0;

void put(int val) {
    assert(count == 0);
    count = 1;
    buffer = val;
}

int get() {
    assert(count == 1);
    count = 0;
    return buffer;
}

int loops;
// cond_t cond;
cond_t empty, full;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        pthread_mutex_lock(&mutex);
        // if (count == 1)
        while (count == 1)
            // pthread_cond_wait(&cond, &mutex);
        pthread_cond_wait(&empty, &mutex);
```

```
    put(i);
    // pthread_cond_signal(&cond);
    pthread_cond_signal(&fill);
    pthread_mutex_unlock(&mutex);
}
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        pthread_mutex_lock(&mutex);
        // if (count == 0)
        while (count == 0)
            // pthread_cond_wait(&cond, &mutex);
            pthread_cond_wait(&fill, &mutex);
        int tmp = get();
        // pthread_cond_signal(&cond);
        pthread_cond_signal(&empty);
        pthread_mutex_unlock(&mutex);
        // process tmp here
    }
}
```

Semaphores

- a **semaphore** is an object with an integer value that can be manipulated with two routines after initialization:
 - `sem_wait`: decrement value by one, and wait if value is negative
 - `sem_post`: increment value by one, if there are one or more threads waiting, wake one
 - when negative, the value of the semaphore is equal to the number of waiting threads
 - the semaphore should be initialized to the number of threads that should enter the critical section at once
- a **binary** semaphore is simply another way to use a lock:
 - the value is initialized to 1
 - to lock, thread calls `sem_wait`, which decrements to 0 and immediately re-

- turns
 - * if another thread tries to lock here, `sem_wait` would decrement to negative and thread would sleep
- critical section then executes
- to unlock, thread calls `sem_post`, which increments back to 0, and wakes any other waiting threads
- want the waiting thread to execute critical section as soon as possible
 - * ie. give away lock immediately after initialization
- using semaphores for *ordering* (similar to condition variables):
 - eg. parent waiting for completion of child thread
 - here, the value should be initialized to 0
 - if parent runs first:
 - * parent calls `sem_wait`, decrements to negative and sleeps
 - * child calls `sem_post`, increments back to 0, and wakes the parent
 - if child runs first:
 - * child calls `sem_post`, increments to 1
 - * parent calls `sem_wait`, decrements to 0 and immediately continues execution
 - want the waiting thread to execute only once a condition has been satisfied
 - * ie. nothing to give away at the start, waiting for child's completion
- using semaphores for **bounded buffer** problem in below example:
 - initially, when `MAX = 1`, example works
 - when `MAX` is increased, need to add mutex locks to make `put` and `get` atomic
 - * need to ensure scope of mutex lock is correct
 - * **deadlock** can occur if mutex lock is outside the conditional variable semaphores

Bounded buffer with semaphores example:

```
sem_t empty, full;
sem_init(&empty, 0, MAX); // 0 indicates semaphores are shared
sem_init(&full, 0, 0);

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        // sem_wait(&mutex); // leads to deadlock
        sem_wait(&empty);
        sem_wait(&mutex);
```

```

    put(i);
    sem_post(&mutex);
    sem_post(&full);
    // sem_post(&mutex); // leads to deadlock
}
}

void *consumer(void *arg) {
    int i, tmp = 0;
    while (tmp != -1)
        sem_wait(&full);
        sem_wait(&mutex);
        tmp = get();
        sem_post(&mutex);
        sem_post(&empty);
        // process tmp
    }
}

```

- using semaphores for *reader-writer locks*:
 - split up locks between reading and writing operations
 - * ie. many lookups can proceed concurrently *as long as* no insert is on going
 - the write lock functions as an ordinary binary lock
 - for readers:
 - * the first reader acquires the write lock
 - * the last reader to release read lock releases the write lock as well
 - not always useful, can introduce excessive overhead
 - * readers may *starve* writers

Reader-writer locks example:

```

typedef struct _rwlock_t {
    sem_t lock;        // basic binary semaphore lock
    sem_t writelock;    // allow ONE writer but MANY readers
    int readers;        // # readers
} rwlock_t;

```

```
void rwlock_init(rwlock_t *rw) {
    rw->readers = 0;
    sem_init(&rw->lock, 0, 1);
    sem_init(&rw->writelock, 0, 1);
}

void rwlock_acquire_readlock(rwlock_t *rw) {
    sem_wait(&rw->lock);
    rw->readers++;
    if (rw->readers == 1) // first reader gets writelock
        sem_wait(&rw->writelock);
    sem_post(&rw->lock);
}

void rwlock_release_readlock(rwlock_t *rw) {
    sem_wait(&rw->lock);
    rw->readers--;
    if (rw->readers == 0) // last writer lets writelock go
        sem_post(&rw->writelock);
    sem_post(&rw->lock);
}

void rwlock_acquire_writelock(rwlock_t *rw) {
    sem_wait(&rw->writelock);
}

void rwlock_release_writelock(rwlock_t *rw) {
    sem_post(&rw->writelock);
}
```

- the dining philosopher's problem:
 - philosophers around a table with a fork on either side
 - * each needs a pair of forks to eat
 - broken solution:
 - * every philosopher grabs left fork and then right fork

- * leads to deadlock
- solution:
 - * one philosopher has to grab forks in a *different* order

Semaphore Implementation

```
typedef struct __sem_t {
    int value;
    pthread_cond_t cond;
    pthread_mutex_t lock;
} sem_t;

void sem_init(sem_t *s, int value) {
    s->value = value;
    cond_init(&s->cond);
    mutex_init(&s->lock);
}

void sem_wait(sem_t *s) {
    mutex_lock(&s->lock);
    while (s->value <= 0)
        cond_wait(&s->cond, &s->lock);
    s->value--;
    mutex_unlock(&s->lock);
}

void sem_post(sem_t *s) {
    mutex_lock(&s->lock);
    s->value++;
    cond_signal(&s->cond);
    mutex_unlock(&s->lock);
}
```


UNIX Synchronization API

- multithreading used to implemented in a *user mode* library, without help from the OS
 - *pros*:
 - * user *sleep/yield* model can be much more efficient than context switching if non-preemptive scheduling is an option
 - *cons*:
 - * user mode lock operations may be much slower than kernel implementations
 - * when a system call blocked, the OS would not know that it is able to overlap other threads
 - all threads in the process would stop executing
 - * user mode library can not exploit multiple cores

Threads

- `pthread_create` handles thread creation
 - `thread` is a pointer to a structure to interact with thread
 - `attr` specifies any thread attributes, `NULL` for defaults
 - `start_routine` is a function pointer to a function that returns a void pointer
 - `arg` is the argument to be passed to start routine
- once a thread is created, has a new call stack but shared address space

`pthread_create` example:

```
#include <pthread.h>
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,
                  void *(*start_routine) (void*), void *arg)

#include <stdio.h>

typedef struct {
    int a; int b;
} myarg_t;

void *mythread(void *arg) {
    myarg_t *args = (myarg_t *) arg;
    printf("%d %d\n", args->a, args->b);
```

```
    return NULL;
}

int main() {
    pthread_t p;
    myarg_t args = {10, 20};
    int rc = pthread_create(&p, NULL, mythread, &args);
}
```

- `pthread_join` waits for a thread to complete
 - `thread` is the thread to wait for
 - `ret_ptr` is a pointer to the return value we are expecting (type `void **`)
 - * never return a pointer to a thread local variable

`pthread_join` example:

```
#include <pthread.h>
int pthread_join(pthread_t thread, void **ret_ptr);

typedef struct { int a; int b; } myarg_t;
typedef struct { int x; int y; } myret_t;

void *mythread(void *arg) {
    myret_t *rvals = malloc(sizeof(myret_t));
    rvals->x = 1;
    rvals->y = 2;
    return (void *) rvals;
}

int main() {
    pthread_t p;
    myret_t *rvals;
    pthread_create(&p, NULL, mythread, NULL);
    pthread_join(p, (void **) &rvals);
    printf("returned %d %d\n", rvals->x, rvals->y);
    free(rvals);
}
```

```
return 0;
}
```

Locks

- need `-pthread` option when compiling with gcc
- `pthread_mutex_lock` gives only control over the lock to the calling thread
 - takes a pointer `pthread_mutex_t` argument indicating the **mutex** lock
 - * locks must be initialized before use by calling `pthread_mutex_init`
 - takes pointer to the lock and an optional attribute structure
 - * locks must be destroyed after use with `pthread_mutex_destroy`
 - returns 0 on success
- `pthread_mutex_unlock` releases a locked lock
 - same args and return as lock
- **condition variables** allow for signaling between threads:
 - used to signal something in the program has changed while a thread was sleeping
 - * safer than using *ad-hoc* flags
 - `pthread_cond_wait` puts the calling thread to sleep until another thread signals it
 - * pass in pointer to condition and pointer to lock for the condition
 - * *releases* lock immediately before sleep, and *requires* it immediately before wake
 - `pthread_cond_signal` sends a signal to wake a thread
 - * pass in pointer to condition
 - * lock must be held when signalling

Example with locks:

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER; // alternative init
pthread_cond_t lock = PTHREAD_COND_INITIALIZER;

// thread 1
pthread_mutex_lock(&lock);
while (ready == 0) // while safer than if when checking condition, even with some
    overhead
    pthread_cond_wait(&cond, &lock);
pthread_mutex_unlock(&lock);
```

```
// thread 2
pthread_mutex_lock(&lock);
ready = 1;
pthread_cond_signal(&cond);
pthread_mutex_unlock(&lock);
```

Example with Linux **futex** locks:

```
void mutex_lock(int *mutex) {
    int v;
    // futex locks use the high bit to track if held
    if (atomic_bit_test_set(mutex, 31) == 0)
        return; // got the lock
    atomic_increment(mutex); // the rest of the bits track number of waiters
    while (1) {
        if (atomic_bit_test_set(mutex, 31) == 0) {
            atomic_decrement(mutex);
            return;
        }
        v = *mutex;
        if (v >= 0)
            continue;
        // if negative, futex is locked

        // if v is equal to mutex, sleep and go into queue
        // otherwise, return immediately
        futex_wait(mutex, v);
    }
}

void mutex_unlock (int *mutex) {
    // adding 0x80000000 results in 0 only if there are no other waiting threads
    if (atomic_add_zero(mutex, 0x80000000))
        return;
}
```

```
// wake another waiting thread up
futex_wake(mutex);
}
```

Appendix

UNIX Syscalls

- `sighandler_t signal(int signum, sighandler_t handler)` - handles signals, registers signal catchers
 - in `signal.h`
 - if `signum` is delivered to the process:
 - * if `handler` is set to `SIG_IGN`, the signal is ignored
 - * if `handler` is set to `SIG_DFL`, the default action occurs
 - * if `handler` is set to a function, the function is called with argument `signum`
 - note that the signals `SIGKILL` and `SIGSTOP` cannot be caught or ignored
 - returns the previous value of the signal handler, or `SIG_ERR`
 - * `errno` set on errors
- `int kill(pid_t pid, int sig)` - sends signals to a process
 - in `sys/types.h`, `signal.h`
 - if `pid` is positive, signal `sig` is sent to process with matching PID
 - if `pid` is 0, `sig` is sent to every process in the process group of the calling process
 - if `pid` is -1, `sig` is sent to every process possible
 - if `sig` is 0, no signal is sent, but existence and permission checks still occur
 - returns 0 on success, returns -1 and `errno` set on error
- `send(int sockfd, const void *buf, size_t len, int flags)` - sends message on a socket
 - in `sys/socket.h`
 - used with a connected socket
 - supports various flag options
 - returns number of bytes sent on success, return -1 and `errno` set on error
- `recv(int sockfd, const void *buf, size_t len, int flags)` - receive message from a socket
 - in `sys/socket.h`
 - supports various flag options

- if message is too long, excess bytes may be discarded
- returns number of bytes received, return -1 and `errno` set on error
- `mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset)` - map or unmap files or devices into memory
 - creates a new mapping in the virtual address space of the calling process, starting at `addr` for `length`
 - the contents of the file mapping are initialized by `length` bytes from `fd` starting at `offset`
 - `prot` specifies memory protections
 - returns a void pointer to the mapped area, return -1 and `errno` set on error
- `flock(int fd, int operation)` - apply or remove advisory lock on an open file
 - `operation` can be `LOCK_SH` to place a shared lock, `LOCK_EX` for exclusive lock, and `LOCK_UN` to remove an existing lock
 - file can only have one type of lock
 - duplicate file descriptors refer to the same lock
 - returns 0 on success, returns -1 and `errno` set on error
- `lockf(int fd, int cmd, off_t len)` - apply, test, or remove POSIX lock on an open file
 - applies to only a section of a file, starting at the current file position for `len` bytes
 - `cmd` can be:
 - * `F_LOCK` to set an exclusive lock, blocks until release if already locked
 - overlapped locks are *merged*
 - * `F_TLOCK` same but call never blocks
 - * `F_ULOCK` unlocks section
 - file locks released on file close
 - may split into two locked sections
 - * `F_TEST` tests the lock
 - 0 if unlocked or locked by process, -1 if other process holds lock
 - returns 0 on success, returns -1 and `errno` set on error

Sockets Example

- **stream socket vs. datagram socket:**
 - datagrams are more *unreliable*, ie. packets can be lost
 - * TCP protocol with streams will detect and *retransmit* lost messages
 - datagrams preserve message *boundaries*
 - * stream sockets may divide messages into chunks
 - much less *overhead* (no initialization/breakdown, no package acknowledgement), so used for short services

Server example:

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

int main(int argc, char *argv[])
{
    int sockfd, newsockfd, portno, clilen;
    char buffer[256];
    struct sockaddr_in serv_addr, cli_addr;
    int n;

    /* create socket:
     * AF_UNIX local, AF_INET network
     * SOCK_STREAM continuous stream, SOCK_DGRAM chunks */
    sockfd = socket(AF_INET, SOCK_STREAM, 0);

    bzero((char *) &serv_addr, sizeof(serv_addr));
    portno = atoi(argv[1]);
    serv_addr.sin_family = AF_INET;
    serv_addr.sin_addr.s_addr = INADDR_ANY;
    serv_addr.sin_port = htons(portno);

    /* bind socket to an address: */
    if (bind(sockfd, (struct sockaddr *) &serv_addr,
            sizeof(serv_addr)) < 0)
        error("ERROR on binding");

    /* listen for connections: */
    listen(sockfd, 5);

    clilen = sizeof(cli_addr);
    /* handle multiple connections */
```

```
while (1) {
    /* repeatedly accept a connection, return new fd: */
    newsockfd = accept(sockfd,
        (struct sockaddr *) &cli_addr, &clilen);
    if (newsockfd < 0)
        error("ERROR on accept");
    pid = fork();
    if (pid < 0)
        error("ERROR on fork");
    if (pid == 0) {
        close(sockfd);
        /* write and read from new fd */
        dostuff(newsockfd);
        exit(0);
    }
    else close(newsockfd);
}
```

Client example:

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>

int main(int argc, char *argv[])
{
    int sockfd, portno, n;
    struct sockaddr_in serv_addr;
    struct hostent *server;
    char buffer[256];

    portno = atoi(argv[2]);
```



```
/* create socket */
sockfd = socket(AF_INET, SOCK_STREAM, 0);
server = gethostbyname(argv[1]);
if (server == NULL) {
    fprintf(stderr, "ERROR, no such host\n");
    exit(0);
}

bzero((char *) &serv_addr, sizeof(serv_addr));
serv_addr.sin_family = AF_INET;
bcopy((char *)server->h_addr,
      (char *)&serv_addr.sin_addr.s_addr,
      server->h_length);
serv_addr.sin_port = htons(portno);

/* connect to server: */
if (connect(sockfd, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) < 0)
    error("ERROR connecting");

/* write and read from new fd */
bzero(buffer, 256);
printf("Please enter the message: ");
fgets(buffer, 255, stdin);
n = write(sockfd, buffer, strlen(buffer));
bzero(buffer, 256);
n = read(sockfd, buffer, 255);
printf("%s\n", buffer);
}
```