

CS118: Computer Network Fundamentals

Professor Lu

Spring 2020

Contents

CS118: Computer Network Fundamentals	2
Overview	2
Access Networks	2
Physical Media	4
Network Core	4
Packet Delay, Loss, and Throughput	6
The Internet	7
Appendix	8
Network Programming	8
Socket Programming API	9

CS118: Computer Network Fundamentals

Overview

- **computer networks** allow for interaction and communications between computers
 - requires certain hardware and software components
 - involves standardized **network protocols**
 - * protocols are complex, with different layers, eg. application, transport, or link layers
 - * used by developers for **network programming**
 - * eg. the **TCP** and **IP** protocol suite used in the today's internet
- the **internet** is a global network for computers
 - hierarchal, has global, regional, and local levels
 - * managed by different **internet service providers (ISP)**
 - *nuts and bolts* view:
 - * **hosts** are the end systems running various network apps
 - billions of connected computing devices
 - clients *and* servers
 - * **communication links**, eg. fiber, copper, radio
 - wired or wireless
 - each has an associated transmission rate and bandwidth
 - different types of connections, eg. phone-wireless, phone-base, router-router, router-server
 - * **routers and switches**
 - deals with transferring **packets** ie. chunks of data
 - act as the in-between between hosts and do not run network apps
 - the **network edge** is made up of the hosts, access networks, and various physical media
 - the **network core** acts as a backbone that deals with actually transferring the data
 - * consists of interconnected routers and the packet/circuit switching method used

Access Networks

- **digital subscriber line (DSL):**
 - uses the *existing* dedicated *telephone* line to connect to a central **DSL access multiplexer (DLSAM)**
 - * **splitter** sends data on the DSL line through internet and voice on the DSL line to telephone net
 - * DLSAM is handled by an ISP
 - requires a dedicated hardware device called a **DSL modem**
 - downstream transmission rate is usually *much faster* than the upstream transmission rate
 - * based on user patterns, users typically download much more than they upload
- **cable network:**
 - alternatively, use the *television* line
 - *similarities* with DSL:
 - * data and TV is *split* and transmitted at different frequencies over a shared cable distribution network
 - * requires hardware device called **cable modem**
 - * connected to a central **cable modem termination system (CMTS)** or **cable headend**
 - * CMTS is handled by an ISP
 - * asymmetric transmission rate
 - unlike DSL, multiple homes are connected via the cable network to the ISP's cable headend
 - * access network is shared, instead of having dedicated access to the central office as with DSL
- **home network:**
 - a *lower* hierarchy of networks
 - within the home, a **wireless access point** is connected to the DSL or cable modem
 - * various devices can *wirelessly* connect to the access point
 - * speed of access point is slower than a direct *wired* connection
 - speed also dependent on the wifi card of the device connecting to the access point
- **enterprise access network or Ethernet:**
 - uses a special hardware device called an **Ethernet switch**
 - connected with ISP through some institutional link and router
 - allows for *much higher* possible transmission rates
 - end systems typically connect into Ethernet switch, eg. WiFi router and PC
- **wireless access networks:**

- shared access networks that connect end systems to routers *wirelessly*
- **wireless local area network (LAN)** can reach within a building (100 ft)
 - * supports up to 450 Mbps rate
 - * eg. 802.11 b/g/n
- **wide-area wireless access** coverage is almost universal (10's km)
 - * provided by a cellular operator
 - * much slower, between 1 and 10 Mbps
 - * eg. 4G, 5G, LTE

Physical Media

- data is *physically* transferred using **bits** that propagate between transmitter/receiver pairs
- a **physical link** lies between the transmitter and receiver
 - eg. common **twisted pair** with two insulated copper wires
- **guided media**:
 - signals propagate through *solid* media, eg. copper, fiber, coax
 - coax cable is made of concentric conductors, allows for bidirectionality
 - * supports multiple channels, **hybrid fiber coax (HFC)**
 - fiber optic cable is a glass fiber carrying light pulses to represent bits
 - * allows for extremely high-speed operation
 - * *immune* to electromagnetic noise
- **unguided media**:
 - signals carried freely through electromagnetic spectrum
 - * no physical wire
 - has issues of reflection, obstruction, interference
 - eg. LAN, wide-area, satellite

Network Core

- the **network core** is a mesh of interconnected routers
 - its role is to send **packets** or chunks of data between hosts
- two key *functions*:
 - **forwarding** relays packets from a router's input to the appropriate router output
 - **routing** determines the source-destination route taken by packets
 - * these routes are computed locally and proactively, and are stored

within the router

- key *technologies*:
 - **packet switching**:
 - * hosts *break* application-layer messages into packets
 - * packets are forwarded between routers, across links, from source to destination
 - packets *hop* through a certain number of intermediate nodes
 - * each packet is transmitted *back-to-back*, not simultaneously, allowing for **full link capacity** transference
 - sending packets takes time $(L \text{ bits}) / (\text{transmission rate } R \text{ bits/sec})$
 - * entire packet must arrive before it can be transmitted (**store and forward**)
 - thus, the *end-to-end* delay is therefore *scaled* to the number of hops the packet must make
 - alternatively, **circuit switching**:
 - * used in traditional telephone networks
 - * no packets, switching granularity is in terms of **circuits**
 - * resources/circuits are *dedicated* for a particular call
 - * reservation-based, no sharing of an in-use circuit
 - * circuits are *released* on call completion
 - *sharing between users* with circuit switching:
 - * with **frequency division multiplexing (FDM)**, split up frequency domain
 - * alternatiely, with **time division multiplexing (TDM)**, use time slices and time sharing
- why is packet switching used by the internet over circuit switching?
 - circuit switching is less **robust**, in that if a part of a circuit fails, it may break the entire network
 - * on the other hand, with packet switching, the network infrastructure is maintained even if some routers go down
 - packet switching also allows for *more users* to use the network
 - * many users will be *idle* for a percentage of their time on the network
 - * eg. with a 1 Mbs link, and each user using 100 Kbs and active 10% of the time:
 - this user pattern is an example of *bursty data*
 - for circuit switching, can only support up to $1 \text{ Mbs} / 100 \text{ Kbs} = 10$ users at a time (*dedicated* circuits)
 - for packet switching, can support 35 users with a probability that > 10 are active that is less than 0.0004
 - the probability that x users are active is: $P(N, x) = \binom{N}{x} p^x (1 - p)^{N-x}$

- * in order to afford a certain number of users, the probability that more than the threshold number of users are active at the same time should be extremely small
- however, excessive **congestion** is still possible with packet switching:
 - * packet delay and loss may occur when the network becomes overloaded with active users
 - packets may have to jump more links in order to alleviate network congestion
 - * thus, certain protocols are needed for reliable data transfer and congestion control
- ie. circuit switching uses *reserved* resources and allows for consistent service, while packet switching uses *on-demand* allocation and less guaranteed service

Packet Delay, Loss, and Throughput

- if the arrival rate to a link *exceeds* the transmission rate for a time:
 - packets will **queue**, and await transmission
 - * the **queuing delay** is the time waited in the buffer before transmitted
 - * *different* from **transmission delay**, the total amount of time to transmit all bits of a packet
 - packets can then be **lost** or dropped if the memory buffer for the queue fills up
- thus **packet delay** overall has multiple sources:
 - **processing delay** from checking bit errors and determining output link
 - **queuing delay** from awaiting transmission, depends on congestion
 - * as $(L \text{ bits} * \text{a average arrival rate}) / R \text{ rate}$ approaches 1, queuing delay becomes large
 - * above 1, the average delay becomes infinite
 - **transmission delay** is how long it takes to push out all bits of the packet, depends on packet size
 - * $L \text{ bits} / R \text{ rate}$
 - **propagation delay** is the time for a bit to actually travel to another router
 - * $d \text{ length} / s \text{ speed}$
- the **traceroute** program provides delay measurement from source to destination
 - send three probe packet that reaches each router along the path
 - measures time interval between transmission and reply
- handling **packet loss**:
 - when a packet is lost, the source must slow its transmission, and also re-

- transmit the lost packet
 - * different *response* for different *applications*:
 - eg. for video streaming, the media will buffer and prioritize lower delay and allow dropping of some packets
 - eg. for emails and communications, delay is not as important as data integrity
 - the exact response is dictated by different transmission protocols eg. TCP
- the **throughput** is the rate at which bits are transferred between sender and receiver
 - can be *instantaneous* or *average*
 - often constrained by the slowest **bottleneck link** in the network

The Internet

- the **internet** is built as a network of networks
- given *millions* of access ISPs, how should they be connected to one another?
 1. pairwise connections, ie. connect each ISP to every other
 - fully distributed and requires $O(n^2)$ connections
 - this solution doesn't scale
 2. connect each ISP to a *global* transit ISP
 - full centralized solution
 - this global ISP becomes a *bottleneck* as all traffic passes through it
 3. use *multiple* global ISPs
 - a natural byproduct of a single global ISP from competition
 - each only serves a subset of its local networks
 - requires **peering links** and **internet exchange points (IXP)** between the global ISPs
 - * IXP are managed by a third party
 - * note that these are less of a bottleneck since global ISPs want to minimize user interaction with another ISP
 4. *hierarchical* structure
 - this is the current structure of the internet
 - at a lower level, several access ISPs are connected to a global ISP through a **regional net**
 - creates a **hierarchy** from access ISPs, to regional nets, to global ISPs
 - another unique level is the **content provider network** eg. Google that brings services and content directly to end users, bypassing the hierarchy
 - this structure is motivated more by business concerns than technical concerns

Appendix

Network Programming

- *application layer*:
 - applications using the network, eg. client-server model
 - **client** initiates communication and awaits the server's response
 - **server** responds to requests
 - * discoverable by clients
 - always running, waiting for client connections
 - * processes requests and sends replies
 - requests can be processed *concurrently*, *sequentially*, or some *hybrid*
 - however, client and server are not always disjoint
 - * eg. server can be a client to another server
- *transport layer*:
 - responsible for actually providing communication services (in conjunction with lower layers)
 - outlined by **protocols** such as TCP and UDP
- **transmission control protocol (TCP)**:
 - full-duplex byte stream
 - has guarantees for *reliable* data transfer:
 - * deliveries are completed
 - * data is ordered, with no duplicates
 - allows for regulated flow for flow and congestion control
- **user data protocol (UDP)**:
 - variable length datagram transfer
 - very basic transmission service
 - no reliability, order, or delivery guarantee
 - no flow or congestion control
- **socket programming**:
 - a **socket** is an endpoint on the network
 - * tuple of **IP address** and **port number**
 - socket programming helps to build the communication tunnel between application and transport layer
 - multiple types of sockets, eg. stream sockets, datagram sockets, and raw sockets for different protocols

- basic steps for working with TCP sockets:
 - create service
 - establish TCP connection
 - send and receive data
 - close TCP connection
- TCP server:
 - create a `socket`
 - `bind` socket to an address
 - `listen` for clients
 - `accept`, blocked until a connection from client
 - `read` and `write` data
- TCP client:
 - create a `socket`
 - `connect` to a server address
 - `read` and `write` data

Socket Programming API

`int socket(int domain, int type, int protocol)` creates a socket:

- returns socket descriptor or -1 and sets `errno`
- `domain` is the protocol family, eg. `PF_INET`, `PF_INET6`
- `type` is communication style, eg. `SOCK_STREAM` for TCP and `SOCK_DGRAM` for UDP
- `protocol` is the specific protocol within family, usually 0

`int bind(int sockfd, struct sockaddr* myaddr, int addrlen)` binds a socket to a *local* address:

- returns 0 or -1 and sets `errno`
- `sockfd` is the socket file descriptor
- `myaddr` is a structure including the IP address and port number
- `sockaddr` and `sockaddr_in` structures are the same size
 - typically, use `sockaddr_in` and cast to `socketaddr`
- `addrlen` is `sizeof(struct sockaddr_in)`

`sockaddr` and `sockaddr_in`:

```
struct sockaddr {
    short sa_family;
    char sa_data[14];
};
```

```
};

struct sockaddr_in {
    short sin_family;          // protocol family, eg. AF_INET (same as PF_INET)
    ushort sin_port;           // port number
    struct in_addr sin_addr;    // IP address
    unsigned char sin_zero[8]; // buffer for having same size as sockaddrr
};

struct in_addr { // used for IPv4 only
    uint32_t sin_port;
};
```

`int listen(int sockfd, int backlog)` waits for connections:

- returns 0 or -1 and sets `errno`
- `sockfd` is the socket file descriptor
- `backlog` is number of connections program can serve simultaneously

`int accept(int sockfd, struct sockaddr* client_addr, int* addrlen)` accepts a new connection:

- return client's socket file descriptor or -1 and sets `errno`
- `sockfd` is the socket file descriptor for server
- `client_addr` to be filled in with IP address and port number of client
- `addrlen` to be filled in with size of the `client_addr`
- a new socket is being cloned from the client
 - if there are no incoming connections to accept, there are multiple blocking modes
 - * non-block: `accept` can return -1
 - * blocking: `accept` operation is added to the wait queue

`int connect(int sockfd, struct sockaddr* server_addr, int addrlen)` connects a socket to a *remote* address:

- return 0 or -1 and sets `errno`
- `sockfd` is the socket file descriptor to be connected
- `server_addr` is the IP address and port number of the server

- `addrlen` is `sizeof(struct sockaddr_in)`

`int write(int sockfd, char* buf, size_t nbytes)` and `int read(int sockfd, char* buf, size_t nbytes)` read and write data from a TCP stream:

- returns number of bytes processed or -1
 - 0 if socket is closed
- `sockfd` is the socket file descriptor
- `buf` is a data buffer
- `nbytes` is the number of bytes to process
 - max to read, or desired number to send

`int close(int sockfd)` closes a socket:

- returns 0 or -1
- `sockfd` is no longer valid

Utilities

- note that **byte ordering** matters when transferring data between host systems
 - little endian vs. big endian
 - hosts may use different orderings, but the **network byte order** is always big endian
 - `ntohl` performs net-to-host long translation
 - `ntohs` performs net-to-host short translation
 - `htonl` performs host-to-net long translation
 - `htons` performs host-to-net short translation
 - thus, the port number and IP address in the API address structures should always be converted
 - * eg. `servaddr.sin_port = htons(servport)`
- other utilities are provided for working with network addresses:
 - `struct hostent* gethostbyname(const char* hostname)` translates a host name to an IP address
 - `struct hostent* gethostbyaddr(const char* addr, size_t len, int family)` translates an IP address to host name
 - `char* inet_ntoa(struct in_addr inaddr)` translates IP address to a dotted-decimal string
 - `int gethostname(char* name, size_t namelen)` reads the local host's name
 - `in_addr_t inet_addr(const char* str)` translates dotted-decimal string to IP address in network byte order
 - * `int inet_aton(const char* str, struct in_addr* inaddr)` same translation, different format

hostent structure:

```
struct hostent {  
    char*   h_name;      // canonical host name  
    char**  h_aliases;   // list of aliases, last element is NULL  
    int     h_addrtype;  // address type, eg. AF_INET  
    int     h_length;    // length of addresses, eg. 4 for IPv4  
    char**  h_addr_list; // list of IP addresses  
    // h_addr is an alias for h_addr_list[0]  
};
```