

CS136: Computer Security

Professor Reiher

Thilan Tran

Winter 2022

Contents

| | |
|--|-----------|
| CS136: Computer Security | 3 |
| Principles, Policies, and Tools | 6 |
| Design Principles | 6 |
| Policies | 7 |
| Tools | 9 |
| Access Control | 10 |
| Cryptography | 13 |
| Cryptanalysis | 14 |
| Symmetric Cryptosystems | 14 |
| Asymmetric Cryptosystems | 17 |
| Checksums and Signatures | 19 |
| Ciphers In Practice | 20 |
| Key Management | 22 |
| Key Exchange | 23 |
| Symmetric Key Exchange | 23 |
| Public Key Exchange | 25 |
| Key Generation | 26 |
| Key Infrastructures | 26 |
| Authentication | 29 |
| Operating Systems | 33 |
| Network Security | 37 |

| | |
|--------------------------------------|----|
| Traffic Control Mechanisms | 38 |
|--------------------------------------|----|

CS136: Computer Security

- why is computer security necessary?
 - people may have malicious intents
 - computers handle a lot of money and a lot of important information
 - society is increasingly dependent on correct operation of computers
- there is now a big problem with computer security:
 - only a matter of time before a real disaster
 - companies go out of business due to DDoS attacks
 - identity theft and phishing
 - cyberattacks e.g. Stuxnet
- examples of large scale security problems:
 - malicious code attacks:
 - * new viruses, worms, Trojan horses, etc. to create ransomware attacks
 - * increasing attacks on infrastructure systems
 - **distributed denial of service (DDoS)** attacks:
 - * use large number of compromised machines to attack, one target
 - * exploiting vulnerabilities or generating lots of traffic
 - * in general form, an extremely hard problem to tackle
 - vulnerabilities in commonly used systems:
 - * systems e.g. Android, Windows, iOS, macOS, etc.
 - * middleware e.g. Windows Installer, Apache, Node.js
 - * even security systems themselves
 - * critical hardware flaws in hardware e.g. Intel and AMD processors
 - electronic commerce attacks
 - * e.g. identity theft, ransomware, extortion, mining on compromised machines
 - cyberwarfare
 - * e.g. Stuxnet, attacks on Ukrainian power grid, cyberspying, Russian election hacking
 - privacy concerns:
 - * data mining by the government
 - * Facebook, Google, Amazon, etc.
 - **passive threats** are forms of eavesdropping, mostly threats to secrecy
 - **active threats** are more active e.g. destruction or interruption / misuse of services
 - social engineering attacks are also a common effective threat
 - * especially phishing
- recent Log4j vulnerability:
 - a programming flaw in a popular package for Java program logging

- allows attacker to force a server to execute arbitrary remote code
 - * essentially end goal of any attacker i.e. “game-over” type of vulnerability
- not a new or sophisticated type of flaw
- why aren’t all computer systems secure?
 - difficult due to hard technical problems
 - as well as cost / benefit issues:
 - * security only pays off when there is trouble
 - * buyers want 100% effectiveness with 0% overhead, learning, inconvenience
 - ignorance also plays a role
 - also constrained by legacy and retrofitting issues:
 - * e.g. core Internet design, popular programming languages, commercial OSs
 - * retrofitting security works poorly, considering the history of patching:
 - when to patch, patches are small and near-sighted, not all software will be patched, patches themselves can have vulnerabilities, etc.
 - malware spreads faster than the patching
- why isn’t security easy?
 - different than most other problems in CS
 - * universe is much more hostile and adversarial, since humans seek to outwit us
 - fundamentally, we want to share secrets in a controlled way
 - * classically hard problem in human relations
 - you have to get everything right:
 - * any mistake is an opportunity for attackers
 - * do we really have to wait for completely bug-free software?
 - computer itself isn’t the only point of vulnerability
 - * users, programmers, system administrators, supply chain
- how common are software security flaws?
 - about 1500 found per year
 - * only considering popular software, real security implications, and publicized flaws
- important definitions:
 - **security** is a policy e.g. no unauthorized user may access this file
 - **protection** is a mechanism e.g. the system checks user identity against access permissions
 - * protection mechanisms implement security policies
 - a **vulnerability** is a weakness that can allow an attacker to cause problems
 - * most vulnerabilities are never exploited

- an **exploit** is an actual incident of taking advantage of a vulnerability
- **trust** or doing certain things for those you trust and not doing them if we do not trust them:
 - * how to express trust, how do we know who someone is, what if trust is situational, what if trust changes
 - * trust relationships such as transitive trust
 - e.g. peer applications, chained certificates, database used by a web server, code that calls code that calls code
 - * most vulnerabilities are based on trust problems
 - taking advantage of misplaced trust
- what are our security goals?
 - confidentiality, integrity, availability
 - involves prevention, detection, or recovery
- what are the categories of threats that security faces?
 - disclosure, deception, disruption, usurpation
- the principle of easiest penetration:
 - an intruder must be expected to use any available means of penetration
 - * not necessarily the most obvious or the one against which there is the most solid defense
 - opponents attack where we are weak
- the principle of adequate protection:
 - worthless things need little protection
 - things with timely value need only be protected for a value

Principles, Policies, and Tools

Design Principles

- each principle has its own tradeoffs
- **economy:**
 - security tool must be economical to develop, use, and verify
 - should add little or no overhead
 - keep it small and simple
- **complete mediation:**
 - apply security on every access to a protected object
 - * e.g. each read of a file, not just the open
 - check access on everything that could be attacked
- **open design:**
 - don't rely on security through obscurity
 - * specifically, secret of how it works vs. secret keys
 - assume all potential attackers know everything about the design
 - obscurity can provide some security, but it's brittle
- **separation of privileges:**
 - provide mechanisms that separate the privileges used for one purpose from those used for another
 - * e.g. separate access control on each file, different passwords for every website
 - allows flexibility in security systems
- **least privilege:**
 - give bare minimum access rights required to complete a task
 - * e.g. don't give write permissions if program asked for read
 - require another request to perform another type of access
 - extremely important when building complex systems
- **least common mechanism:**
 - avoid sharing parts of the system's mechanism
 - * among different users or different parts of the system
 - coupling leads to possible security breaches

- **acceptability:**
 - mechanism must be simple to use
 - * people use it without thinking about it
 - must rarely or never prevent permissible accesses
 - e.g. principle of least astonishment
- **fail-safe designs:**
 - default to lack of access
 - so if something goes wrong or is forgotten or isn't done, no security is lost
 - * if false negatives, we can change the default on an individual basis

Policies

- policies describe how a secure system should behave:
 - describes what should happen, not how you achieve that
 - if you don't have a clear policy, you don't have a secure system
 - * you don't know what you're trying to do
 - should address all relevant aspects of confidentiality, integrity, and availability
 - *difficulties*:
 - * hard to define policies properly
 - * hard to go from policy to the mechanisms
 - * hard to understand implications of policy
- informal policies:
 - e.g. “system executable should only be altered by system administrators”
 - e.g. “users should only be able to access their own files”
- formal policies:
 - typically expressed in a mathematical security policy language
 - * e.g. Bell-La Padula model
 - tending towards precision
 - hard to express in formal ways and reason about them
 - mathematically, a policy partitions the system states into a set of authorized and unauthorized states:
 - * secure system starts in an authorized state and cannot enter an unauthorized state
 - * can reason about the system as an FSM moving between the states
- the **Bell-La Padula model** is the best-known formal computer security model:
 - two parts of clearances and classifications
 - * real systems use classes of information with different classifications

- corresponds to military classifications
- combines mandatory and discretionary access control
- each object has a **classification**:
 - * describes how sensitive the object is
 - * using same categories as clearances
 - * informally, only people with the same or higher **clearance** should be able to access objects of a particular classification
 - a subject S can read object O iff. $l_O \leq l_S$
- also concerned with object contents, not just objects themselves
 - * what if someone with top secret clearance writes the information to a lower classification object?
 - * additional Bell-La Padula *-property:
 - S can write O iff. $l_S \leq l_O$
 - prevents **write-down**
- how do you use the system?
 - * due to write-down, cannot communicate with someone lower privilege
 - * needs mechanisms for reclassification, requiring explicit operation
- can prove a system meeting these properties is secure in terms of confidentiality:
 - * doesn't address integrity at all
 - * **confidentiality policies** place no trust in object, just whether an object can be disclosed
- on the other hand, **integrity security** policies are designed to ensure that information is not improperly changed:
 - key issue for commercial systems
 - secrecy is nice, but not losing track of inventory is crucial
 - integrity policies operate based on how much some object can be trusted
 - * policies then dictate what a subject can do with that object
- the **Biba** integrity policy addresses integrity:
 - subject set S , object set O , set of ordered integrity levels I :
 - * subjects at high integrity levels are less likely to screw up data
 - * data at a high integrity level is less likely to be screwed up
 - s can write to o iff. $i(o) \leq i(s)$
 - s_1 can execute s_2 iff. $i(s_2) \leq i(s_1)$
 - s can read o iff. $i(s) \leq i(o)$
- in hybrid models, sometimes the issue is keeping things carefully separated:
 - issues of *both* confidentiality and integrity
 - e.g. in the Chinese Wall model, all the resources, computers, people are separated and do not touch each other

Tools

- **physical security:**
 - lock up your computer
 - but what about networking, and mobility?
 - in any case, lack of physical security often makes other measures pointless
- **access controls:**
 - only let authorized parties access the system
 - difficult, particularly in a network environment
- **encryption:**
 - algorithms to hide the content of data or communications
 - only those knowing a secret can decrypt the protection
 - one of the most important tools
- **authentication:**
 - methods of ensuring that someone is who they say they are
 - vital for access control
 - often based on cryptography
- **encapsulation:**
 - methods of allowing outsiders limited access to resources
 - * preferably making inaccessible things invisible
 - challenging in practice
- **intrusion detection:**
 - need to notice failures and take steps
 - reactive, not preventative
 - should be automatic to be really useful
- **common sense** is also a tool
 - social engineering attacks

Access Control

- how do we give access to only the right people?
 - at the right time and circumstances
 - similarly, how do we ensure a given resource can only be accessed when it should be
 - *goals*:
 - * complete mediation
 - * least privilege
 - * useful in a network environment
 - * scalability
 - * acceptable cost and usability
- main types:
 1. access control lists
 2. capabilities
 3. access control matrix (both of the first two)
 4. role based access control
- definitions:
 - **subjects** are active entities that want to gain access to something e.g. users or programs
 - **objects** represent things that can be accessed e.g. files, devices, records
 - **access** is any form of interaction with an object
 - **mandatory** access control is dictated by the underlying system
 - * can't be overridden by individual users
 - **discretionary** access control is under the command of the user:
 - * system enforces what they choose
 - * most users never change the defaults
 - * not wise to rely on it to protect important information, for system designers
- **access control lists (ACL)** is the first mechanisms for implementing access control:
 - for each protected resource, maintain a single list
 - each list entry specifies a user who can access the resource, and allowable modes of access
 - when user requests access, check the ACL
 - * can also have lower granularity than per-user, e.g. dictate per-process access
 - *issues*:
 - * how do we know subject is who he says he is?
 - * how do we protect the ACL itself from modification?
 - * how do you determine what resources a user can access?

- would have to check every single ACL, inefficient
- *pros*:
 - * easy to find who can access a resource
 - * easy to change permissions
- *cons*:
 - * hard to find which resources a subject can access
 - * changing access rights requires getting to the object e.g. across the Internet, distributed systems
- used by most modern systems:
 - * e.g. Linux, Windows, Android
 - * to prevent additional lookups, only check ACL on *first* open:
 - maintain metadata in file descriptor, which begins to act as a capability
 - not the safest
- utilizing **capabilities** is the second main mechanism for access control:
 - each subject keeps a set of data items that specify his allowable accesses
 - * i.e. a set of tickets
 - possession of the capability for an object implies that access is allowed
 - capabilities *must* be unforgeable:
 - * in single machine, OS is in charge of capabilities
 - * what about networked systems?
 - in most systems, some capabilities allow creation of other capabilities
 - * allows process to pass a restricted set of capabilities to a subprocess
 - much more difficult with ACLs
 - *pros*:
 - * easy to determine what a subject can access
 - * potentially faster than ACLs
 - * easy model for transfer of privileges
 - *cons*:
 - * hard to determine who can access an object
 - * requires extra mechanism for revocation
 - * in a network, need cryptographic methods to prevent forgery
- how can we revoke a capability?
 - destroy the capability
 - * how can we find it?
 - revoke on use?
 - * requires checking on use
 - * essentially turning capability into ACL
 - generation numbers i.e. generations of capabilities?
 - * requires updating non-revoked capabilities
 - * needs another list of subjects
- in distributed access control:
 - ACLs still work OK:

- * provided we have a global namespace for subjects
 - * and no one can masquerade
- capabilities are more problematic:
 - * relies on unforgeability
 - * provided by cryptographic methods
 - prevents forging, not copying
- **role based access control** is an enhancement to ACLs or capabilities:
 - each user has certain roles he can take while using the system
 - * at any time, the user is performing a certain role
 - give the user access to only those things required to fulfill that role
 - available in some form in most modern OSes
 - only helps if changing roles isn't trivial
 - * typically requires secure authentication
 - *limitations*:
 - * number of roles per user
 - * disjoint role privileges
 - * system administration overheads
 - * usability and management problems
- whatever form it is, access control must be instantiated in actual code:
 - needs to check if a given attempt to reference an object should be allowed AKA a **reference monitor**
 - good reference monitors are critical for system security
 - *properties*:
 - * correctness
 - * proper placement
 - * efficiency
 - * simplicity
 - * flexibility

Cryptography

- the goal of **cryptography** is to keep enciphered information secret:
 - encryption is the process of hiding information in plain sight through transformation:
 - * transform the secret data into something else
 - * make the secret hard for others to read
 - while making it simple for authorized parties to read
 - counters disclosure
 - can be used to provide integrity of data and origin
 - * counters modification and masquerading
 - can be used to provide non-repudiation
 - * counters repudiation of origin
- the basic component of cryptography is a **cryptosystem**:
 - sender S , receiver R , attacker O :
 - * **encryption** makes the message unreadable or unalterable by O
 - * **decryption** makes the encrypted message readable by O
 - * rules for transformation called the **cipher**
 - 5-tuple (E, D, M, K, C)
 - $E : M \times K \rightarrow C$ is the set of **encryption functions**
 - $D : C \times K \rightarrow M$ is the set of **decryption functions**
 - M is the set of **plaintexts**
 - K is the set of **keys**:
 - * most algorithms use a key (usually secret) to perform encryption and decryption
 - * if you change only the key, a given plaintext encrypts to a different ciphertext
 - C is the set of **ciphertexts**
- desirable characteristics of ciphers:
 - amount of secrecy required should match labor to achieve it
 - freedom from complexity
 - simplicity of implementation
 - * probability of error is lower
 - errors should not propagate e.g. consider if bits get flipped
 - ciphertext size should be same as plaintext size
 - encryption should maximize **confusion** i.e. plaintext and ciphertext relationship should be complex
 - encryption should maximize **diffusion** i.e. plaintext information is distributed throughout ciphertext

Cryptanalysis

- **cryptanalysis** is the process of trying to break a cryptosystem:
 - finding the meaning of an encrypted message without knowing the key
 - successful when you don't get garbage when decrypting
 - * almost all messages will be garbage if the key is wrong, only $\frac{1}{2^N}$ are sensible
- forms of cryptanalysis:
 1. analyze an encrypted message and deduce its contents
 2. analyze one or more encrypted messages to find a common key
 3. analyze a cryptosystem to find a fundamental flaw
- types of attacks:
 1. ciphertext only:
 - no plaintext knowledge or details of algorithm
 - must work with probability distributions, patterns of common characters, etc.
 - hardest type of attack
 2. known plaintext
 - have matching sample of ciphertext and plaintext
 3. chosen plaintext e.g. differential cryptanalysis
 - clever choices of plaintext may reveal many details
 4. algorithm and ciphertext:
 - can use exhaustive runs of algorithm against guesses at plaintext
 - or try and brute force
 - or, in a **timing attack**:
 - * have ability to watch algorithm encrypting and decrypting
 - * some algorithms perform different operations based on key values
 - * watch timing or observe power use to try to deduce keys
 - * successful against some smart card crypto
 - in many cases, intent is to guess the key
- most cryptosystems are breakable:
 - some just cost more to break than others
 - the job of the cryptosystem designer is to make the cost infeasible or incommensurate with the benefit extracted

Symmetric Cryptosystems

- **symmetric** cryptosystems have the same key for encipherment and decipherment

- i.e. there is a $D_k \in D$ such that $D_k(E_k(m)) = m$ for message m
- *pros*:
 - * encryption and authentication performed in a single operation
 - * well-known and trusted ones perform faster than asymmetric key systems
 - * doesn't require any centralized authority
 - key servers can help
- *cons*:
 - * makes signature more difficult
 - * non-repudiation
 - * key distribution
 - * scaling
- **transposition** or **permutation** ciphers diffuse the data in the plaintext:
 - the letters are not changed, only rearranged
 - * e.g. columnar transpositions, double transpositions
 - detected by comparing character frequencies with a model of the language
 - can be attacked by anagramming i.e. rearranging the ciphertext:
 - cannot be attacked by examining individual letter frequencies
 - * could check frequencies of **digrams** i.e. pairs of letters
- **substitution** ciphers change characters in the plaintext:
 - decrypt by reversing the substitutions
 - e.g. in a Caesar cipher, we translate each letter a fixed number of positions in the alphabet:
 - * simple, but no good diffusion or confusion
 - * could attack using letter frequencies to figure out the offset
 - the more ciphertext we have, the easier the attack
 - a **monoalphabetic** cipher maps every character into another character in one alphabet
 - * preserves the statistics of the underlying message
 - a **polyalphabetic** cipher uses multiple alphabets, obscuring the statistics:
 - * if patterns aren't hidden well, we don't gain much
 - * can be attacked by examining repetitions:
 - **index of coincidence** predicts the number of alphabets used to perform the encryption
 - requires lots of ciphertext
- there is a “perfect” substitution cipher, the **one-time pad**:
 - one that is theoretically and practically unbreakable without the key
 - * and the key cannot be guessed, if we chose the key correctly
 - we use non-repeating keys, where we use a new substitution alphabet for *every* character:
 - * substitution alphabets chosen purely at *random*, and these consti-

- tute the key
 - e.g. flip a coin many times to create a key stream
- * *any* key was equally likely
- * *any* plaintext could have produced this message
- * no longer has the property that only one key gives a non-garbage answer
- usually done in practice with bits, not characters
- *pros*:
 - * if key is truly random, provable that it cannot be broken
- *cons*:
 - * need one bit of key per bit of message
 - * key distribution is painful
 - * key synchronization is vital
 - * good random number generator is hard to find
- typically not used, very difficult in practice:
 - * pads distributed with some other cryptographic mechanism
 - * pads generated non-randomly
 - * pads reused
- in **quantum cryptography**, we use quantum mechanics to perform cryptography:
 - mostly for key exchange
 - relies on quantum entanglement or indeterminacy
 - can also use quantum computers to break cryptography:
 - * famously can potentially break RSA
 - but has no use in cracking AES
 - * currently non-feasible in reality
- modern ciphers tend to use both transposition and substitution
 - hide text patterns and also hide underlying text characters
- the **data encryption standard (DES)** is a classic symmetric cryptosystem:
 - bit-oriented
 - uses both transposition and substitution i.e. is a product cipher
 - input, output, and key are each 64 bits AKA one **block** long
 - consists of 16 rounds:
 - * each round uses a separate key of 48 bits
 - generated from the key block by dropping parity bits, permuting, and extracting 48
 - * if the order in which the round keys is used is reversed, input is deciphered
 - * input of one round is output of the previous round
 - * right input half and round key are ran through a function f that produces 32 bits of out
 - output is XORed into left half, and halves are swapped
 - f takes the right half of the input, expands it, and XORs it with the

- round key
 - * the resulting 48 bits are split into eight sets of six bits each
 - * each set is put through a substitution table called the S-box that produces four bits of output
 - * results are concatenated into a single 32-bit quantity, which is then permuted
- used from 1976 to 2001 (until the release of AES) as an official cryptography standard
- *weaknesses*:
 - * key length of 56 bits is too short
 - * had weak and semiweak keys
 - * S-boxes were classified, suggesting that the classification hid ways to invert the cipher
 - * S-boxes exhibited non-randomness
- the **advanced encryption standard (AES)** succeeded the DES:
 - another bit-oriented product cipher
 - can use keys of 128, 192, or 256 bits
 - operates on 128 bits of input, producing 128 output bits
 - initial state array is transformed over the rounds into the output
 - consists of between 10 and 14 rounds:
 - * round key for each round generated by rotating and substituting the words in the original key
 - * round key added into state array, substitutions performed rows are shifted, and columns are mixed
 - basic operations such as XOR allows for high-performance implementations
 - *advantages over DES*:
 - * larger keys and better round key generation
 - * S-box values are nonlinear and algebraically complex
 - * inputs bits are rapidly diffused
 - * no weak or semiweak keys
 - *weaknesses*:
 - * attacks work on version of AES using fewer rounds
 - * attacks get keys quicker than brute force, but not practical time
 - * unusable flaws often suggest presence of usable ones

Asymmetric Cryptosystems

- a new type of cryptography proposed in 1976 had different keys for encoding and decoding:
 - keys created in pairs:

- * one key is public and its complementary key must remain secret
- * if you want to send an encrypted message, encrypt with his public key, and only he can decrypt
- thus, this public key system should meet the following conditions:
 1. computationally easy to encode or decode given the key
 2. computationally infeasible to derive the private key from the public key
 3. computationally infeasible to determine the private key from a chosen plaintext attack
- typically based on either NP-complete problems or hard mathematical problems e.g. finding factors
- *vs. symmetric cryptosystems:*
 - * easier authentication
 - no need to distribute a shared key
 - * nicer scaling properties
 - each user just needs a key pair
- new challenge is publishing public keys in a trustworthy manner:
 - security depends on using the right public key
 - need high assurance a given key belongs to a particular person
 - needs some sort of **key distribution infrastructure**
- quick authentication with public keys:
 - to sign a message, simply encrypt it with your own private key
 - * only you know the private key, so no one else could create that message
 - everyone knows the public key, so everyone can check the claim
 - solves some issues with shared key authentication
- ideally, we want to use both symmetric and asymmetric cryptography:
 - public key used to “bootstrap” symmetric communication
 - e.g. RSA to authenticate and establish a session key
 - * use AES with that session key for the rest of the transmission
- the notable **RSA** cryptosystem was introduced in 1978:
 - most popular public key algorithm, in wide use
 - * has withstood much cryptanalysis
 - an exponentiation cipher based on factoring large numbers
 - 1. given two large prime numbers p and q , the **totient** $\phi(n)$ of $n = pq$ is the number of numbers less than n with no factors in common with n
 - alternatively, $\phi(n) = (p - 1)(q - 1)$
 - 2. choose an integer $e < n$ that is relatively prime to $\phi(n)$
 - 3. find a second integer such that $ed \bmod \phi(n) = 1$
 - public key is (e, n) and the private key is d
 - * i.e. functions of a pair of 100-200 digit prime numbers
 - recovering plaintext without private key is supposedly equivalent to factoring product of the prime numbers

- vs. *AES*:
 - * *AES* is much more complex
 - but only arithmetic, logic, and table lookup
 - * *RSA* uses exponentiation to large powers
 - much more computationally expensive
 - * *RSA* key selection also more expensive
- elliptic curve cryptography:
 - another math problem
 - can give good security with much smaller keys
 - often used for small devices
- attacking public key systems:
 - nobody uses brute force attacks of checking 2^{2048} keys
 - instead, attack the mathematical relationship between public and private key

Checksums and Signatures

- in some cases, secrecy isn't necessary, but authentication is required:
 - data must be guaranteed to be unchanged
 - important for long-lived data
- desired signature properties:
 - unforgeable
 - verifiable
 - non-repudiable
 - cheap to compute and verify
 - non-reusable
 - no reliance on trusted authority
- signatures with shared key encryption require a trusted third party:
 - third party needed so receiver cannot forge the signature
 - instead, third party checks validity with secret keys shared with them
- with public keys:
 - signer can simply encrypt the document with his private key
 - receiver decrypts with signer's public key
 - no trusted third party needed, but receiver must be certain he has the right public key
 - to save on computation, or if we don't need encryption, just sign a checksum only
- a **checksum** or **message digest** is used to check against tampering:
 - e.g. parity bit is a simple checksum
 - should meet the following conditions:
 1. checksum is easy to compute

2. computationally infeasible to find the input from a checksum value
3. computationally infeasible to find another different input that gives the same checksum value:
 - * by the pigeonhole principle, several messages *must* produce the same checksum
 - * ideally, the hashes of all possible messages will be evenly distributed over the possible checksums
- HMAC is a generic term for an algorithm that uses a keyless hash function and a cryptographic key to produce a keyed hash function:
 - used in public key systems to validate data is unchanged in transit
 - without the key, anyone can change the data and recompute a digest

Ciphers In Practice

- some issues can arise when using cryptosystems in practice:
 - messages can be precomputed
 - * in a small set of possible plaintexts, an attacker can use a “forward search” to precompute and compare ciphertexts
 - blocks can be misordered:
 - * e.g. over a network, parts of a message can be deleted replayed or reordered
 - * can checksum the entire message or have a sequence number in each block
 - statistical regularities
 - * independence of parts of ciphertext can give information relating to the structure of the message, even if it is unintelligible
 - type flaw attacks
 - * exploiting the structure or components of messages
- ciphers will often divide a message into a sequence of blocks:
 - can encipher each block with the same key, or use a nonrepeating stream of key elements AKA **stream ciphers**
 - **block ciphers** work on a given sized chunk of data at a time
- stream ciphers:
 - how can we generate a random, infinitely long key?
 - * an algorithm is used to create the new key
 - e.g. RC4 cipher creates a changing, supposedly unpredictable, key stream
 - * can use shift registers, or even obtain the key from the plaintext or ciphertext
 - *pros*:
 - * speed of encryption and decryption

- each symbol encrypted as soon as its available
 - * low error propagation:
 - errors affect only the symbol where the error occurred
 - depends on cryptographic mode
- *cons*:
 - * low diffusion, each symbol separately encrypted
 - * susceptible to insertions and modifications in the middle of a stream cipher
 - * not good match for many common uses of cryptography
 - can mitigate some issues with proper cryptographic mode
- block ciphers:
 - most common Internet cryptography done with block ciphers
 - *pros*:
 - * good diffusion
 - * immunity to insertions
 - *cons*:
 - * slower
 - * worse error propagation
- we have a bunch of data to encrypt using the same cipher and key:
 - block ciphers have limited block size and stream ciphers just keep going
 - if we encrypt naively:
 - * two blocks with identical plaintext encrypt to the same ciphertext!
 - * each block of data was independently encrypted with the same key
 - * we used the wrong **cryptographic mode** i.e. way of applying a particular cipher!
- cryptographic modes:
 - a combination of cipher, key, and feedback
 - in **electronic codebook (ECB)** mode, simply perform block cipher encryption block by block
 - in **cipher block chaining (CBC)** mode, a group of related encrypted blocks are tied together:
 - * hides that two blocks are identical, foiling insertion attacks
 - * the encryption version of the previous block is used to encrypt this block by XORing them together
 - adding feedback into the encryption
 - * however, we have to fix the first block:
 - use **initialization vectors (IV)**
 - XOR a random string with the first block
 - ensures encryption results are always unique
 - cipher-feedback mode and output-feedback mode both convert block to stream cipher

Key Management

- it doesn't matter how strong the algorithm is if the keys are insecure:
 - proper use of keys is crucial
 - ciphers don't get cracked often, but keys get leaked all the time
- if algorithm is otherwise completely secure, strength depends on key length:
 - since the only attack is a brute force attempt
 - however, with longer keys, encryption costs more and is slower
 - some algorithms have defined key lengths only
- **perfect forward secrecy** means that the compromise of any one session key will not compromise any other
 - keys get divulged, so minimize the resulting damage
- key lifetime consideration:
 - long-lived keys are more likely to be compromised
 - more data is exposed
 - easier cryptanalysis
 - more resources attackers can devote to breaking it
 - even old keys can be found in multiple places after being destroyed:
 - * e.g. caches, virtual memory, freed file blocks, stack frames, etc.
 - * need to zero out the key value
- key lifetime examples:
 - symmetric session keys:
 - * e.g. keys for specific communications sessions should be changed often
 - * avoid storing them permanently
 - long term symmetric keys:
 - * e.g. disk encryption
 - * safe storage is critical
 - private asymmetric keys:
 - * long-term storage as well
 - * safe storage is critical
- storing a user's keys:
 - permanently on machine
 - * machine can be cracked
 - difficult to remember keys
 - * hash keys from passwords or passphrases
 - smart cards
 - key servers
- key secrecy breaches:
 - private keys are often shared:
 - * for convenience

- * to share expensive certificates
- * don't know any better
- entire security of public key system depends on the secrecy of the private key

Key Exchange

- an **interchange key** is associated with a principal i.e. user
 - while a **session key** is associated with the communication session itself
- the first hurdle to overcome is transmitting the session key:
 - session key must be encrypted when it is exchanged
 - * in order to exchange, may need a trusted third party

Symmetric Key Exchange

- simple symmetric key exchange:
 1. A asks third party C to start a session with B
 2. C sends to A the session key encrypted with A's key, followed by the session key encrypted with B's key
 3. A sends to B the session key encrypted with B's key
 - note that A's key and B's key are keys shared with them and the trusted third party
 - * symmetric, not public keys
 - vulnerable to a man-in-the-middle attack, before the following minor changes:
 - * encrypt request with A's key
 - * include identity of other participant in response from C
 - however, still compromised using repeating messages
- types of security protocols:
 - **arbitrated protocols** involve a trusted third party
 - **adjudicated protocols** involve a trusted third party, after the fact
 - **self-enforcing protocols** do not involve a third party
- Needham-Shroeder protocol:
 - another symmetric key exchange and authentication protocol
 - uses **nonces** or randomly generated numbers to defend against replay attacks
 - 1. A sends to third party C: A's name, B's name, and a nonce r_1
 - 2. C sends to A: A's name, B's name, r_1 , the session key, A's name plus the session key encrypted with B's key, all encrypted with A's key:
 - i.e. $\{A||B||r_1||k_{session}||\{A||k_{session}\}_{k_B}\}_{k_A}$
 - A is now sure of who they are talking to, and the nonce assures against replay attacks

3. A sends to B: A's name and the session key, encrypted with B's key
 - B now knows who they are talking to
4. B sends to A: another nonce r_2 encrypted with the session key
5. A sends to B: $r_2 - 1$ encrypted with the session key
 - cannot be easily compromised with repeated messages:
 - * still possible for old session keys to be cracked by attackers, and B's challenge to A can be forged
 - * in this case, can add timestamps to further counter repeats, which requires synchronized clocks
 - e.g. Kerberos protocol with tickets
- global clocks and timestamps:
 - often hard to obtain a globally synchronized set of clocks
 - * attacker can attack clocks as well
 - in a suppress-replay attack, attacker can intercept and replay if the sender's clock is behind
 - clock solutions:
 1. rely on clocks that are fairly synchronized and hard to tamper with e.g. GPS signals
 2. make all comparisons against the same clock
- Otway-Rees protocol:
 - avoids timestamps
 - uses an integer n to associate all messages with a particular exchange
 - 1. A sends to B: n , A's name, B's name, and r_1 plus n plus A's name plus B's name encrypted with A's key
 - i.e. $\{n||A||B||\{r_1||n||A||B\}_{k_A}\}$
 - 2. B sends to third party C: n , A's name, B's name, r_1 plus n plus A's name plus B's name encrypted with A's key, and r_2 plus n plus A's name plus B's name encrypted with B's key
 - i.e. $\{n||A||B||\{r_1||n||A||B\}_{k_A}||\{r_2||n||A||B\}_{k_B}\}$
 - 3. C sends to B: n , r_2 and the session key encrypted with A's key, and r_2 and the session key encrypted with B's key
 - 4. B sends to A: n , and r_1 and the session key encrypted with A's key
 - goal is to prevent replay attacks
- Bellare-Rogaway protocol:
 - considers authentication and symmetric key exchange different problems
 - * protocol only provides the key exchange, trusted server sends to both parties
 - 1. A sends to B: A's name, B's name, r_1
 - 2. B sends to third party C: A's name, B's name, r_1 , r_2
 - 3. C sends to B: session key encrypted with B's key, keyed hash of A's name, B's name, r_1 , and the session key encrypted with B's key
 - keyed hash utilizes the user's interchange key

4. C sends to A: same as (3), but with A's keys

Public Key Exchange

- conceptually, public keys makes exchange keys very easy:
 - A sends to B: the session key encrypted with B's public key
 - * attacker can easily forge message
 - (revised) A sends to B: A's name and the session key encrypted with A's private key, all encrypted by B's public key
 - * after receiving the message, B can use A's public key to obtain the session key
- **man in the middle** attack:
 - occurs when A has to first obtain B's public key
 - 1. A asks C for B's public key
 - attacker intercepts and asks C themselves for B's public key
 - 2. C responds to attacker with B's public key
 - 3. attacker sends to A their own public key
 - 4. A sends to B the session key encrypted with the attackers public key
 - attacker intercepts again and sends to B themselves the session key encrypted with B's public key
 - no binding of identity to a public key
 - * to resolve, need to look at management of cryptographic keys
- Diffie-Hellman key exchange:
 - securely exchange a key:
 - * without previously sharing any secrets
 - * no public key available or symmetric key
 - * using an insecure channel
 - first two parties need to agree on a large prime n and a number g
 - * n, g don't need to be secrets, typically predefined in their software
 - 1. A chooses a large random integer x and sends B $X = g^x \mod n$
 - 2. B chooses a large random integer y and sends A $Y = g^y \mod n$
 - 3. A computes $k = Y^x \mod n$
 - 4. B computes $k' = X^y \mod n$
 - $k = k' = g^{xy} \mod n$
 - but nobody else can compute k, k' !
 - * others know n, g, X, Y , but not x, y
 - * knowing X, Y gets you nothing, unless you compute the discrete logarithm to obtain x or y
 - believed to be hard
 - * typically, x, y are just the users private keys
 - D-H guarantees that two parties share a secret:
 - * but it doesn't guarantee who those two parties are
 - * how does A know whether the Y she heard was sent by B?

- * D-H does not authenticate the parties
- authentication in any key distribution is a core Internet problem, TC/IP does no authentication!
 - * however, D-H is used all the time

Key Generation

- a sequence of **random** numbers is a sequence such that an observer cannot predict any x_k even if all the previous numbers are known:
 - requires physical source of randomness or noise, e.g. background radiation, electromagnetic phenomena, biometrics, disk drive delay
 - * done in the background AKA **gathering entropy**
 - on the other hand, a sequence of **pseudorandom** numbers is a sequence generated by an algorithm intended to simulate random numbers
 - * need statistical properties and non-reproducibility
- pseudorandom generators:
 - how good is that generator?
 - * don't use `rand`
 - linear congruential generator $x_k = (ax_{k-1} + b) \bmod n$ has been broken, as well as polynomial congruential generator
 - the outputs of a **strong mixing function** depend on some nonlinear function of all input bits e.g. SHA:
 - * best generator algorithms
 - * one approach is to continue to hash old ones to produce new keys
 - does not have perfect forward secrecy, and depends on strength of the hash algorithm

Key Infrastructures

- how can we guarantee the true owner that a public key belongs to?
 - need a trusted third party or authority to sign some sort of **certificate** binding an identity to a cryptographic key
 - * or some kind of central server
 - but now we need to distribute the third party's public key... which needs to be verified by an additional certificate?
 - * there is no universally trusted single authority
 - * does everyone need the public keys for all certificate authorities?
- **key servers** are machines whose job it is to distribute keys to other machines:
 - clients can authenticate themselves to the server
 - server can authenticate itself to the clients

- bootstrapping and transitive trust issue
- not the popular solution
- certificate is essentially a copy of a public key together with an identity signed by a trusted authority:
 - usually has an expiration date
 - presentation of the certificate alone serves as authentication of your public key
 - problems during certification process:
 - * what measures did CA take before issuing?
 - * how long is certificate valid for?
 - * is CA's own certification still valid?
 - * who is trustworthy enough to be at the top of the hierarchy?
 - * what do we do when keys are compromised?
 - * user may have different standards than the CA
 - revocation is a general problem for keys, certificates, etc.
 - * how does the system revoke something related to trust, in a network environment?
 - * related to revocation problem for capabilities
 - * one approach is OCSP, an online system that indicates if certificates have been revoked
 - used in different ways by different OSes and browsers
 - typically, most attackers do not break in using certificate validity:
 - * not the weakest link
 - * but now being exploited, mostly by sophisticated adversaries
- **Merkle's tree authentication:**
 - keeps public keys and their associated identities as data in a file
 - * uses checksums to detect data integrity breaches
 - keys and identities in the file are organized into a tree structure
 - * hash of the entire file is the **root**
 - during validation, can traverse an **authentication path** on the tree to verify the checksums
 - * if the root value does not match, an identity / key pair has been compromised
 - *pros:*
 - * creates certificates without using public key signatures
 - * suggests natural hierarchies
 - *cons:*
 - * requires entire file
 - * any changes requires wide redistribution
- a **certificate authority (CA)** is an entity that issues certificates:
 - there is no one CA for the entire Internet
 - CAs could be organized into a single hierarchy:
 - * single CA at the top supplies certificates for the next layer, etc.

- * in practice however, we rely on large numbers of independent certifying authorities, each of which may have its own internal hierarchy
- for new certificates by an unknown CA, the certificates also contain that authority's certificate
- in reality, most OSes or browsers come with a set of “pre-trusted” certificate authorities (sometimes around hundred certificates):
 - * system automatically trusts certificates they sign
 - usually no hierarchy
 - * if not signed by one of those, present it to the user

Authentication

- generally, **authentication** is the binding of an identity to a subject:
 - e.g. process, machine, human user
 - physically identify through credentials, recommendation, knowledge, location, etc.
 - * these all have cyber analogs
 - * but, authentication is done over a network, even if the party is human
 - everything is converted to digital signal
 - * in addition, identity might not be rechecked
 - more general than authentication in cryptography
 - * access control only works if you have good authentication
 - authorization is determining what someone can do
 - there is a certain set of specific information with which entities prove their identities
 - * can be passwords, biometrics, etc.
 - importantly, there is another set of information which the system stores that is used to validate the authentication information from the user
 - * i.e. the complementary information
- the simplest authentication mechanism is a **password**, often a sequence of characters:
 - i.e. authenticated by what you know
 - complement can simply be the password in plaintext:
 - * instead, should hash the password into a complement using a *one-way* function
 - * retrieving the password file does not allow you to log in to the system
 - password selection:
 - * random selection of passwords
 - strength of the pseudorandom generator
 - * computer-generated pronounceable passwords
 - less strong, but easier to remember
 - * user selected passwords
 - should avoid names, dictionary words, keyboard patterns, short passwords, etc.
 - * graphical passwords
 - typically, passwords are **salted** by adding random data before the password is hashed:
 - * random number need not be secret
 - * just different for different users

- * makes dictionary attacks much more difficult
 - * similar to nonces and initialization vectors
- passwords have an aging issue:
 - * can be cracked over time
 - * should change passwords periodically
 - * one-time passwords invalidate immediately
 - * many systems ask for password once, trading security for convenience
- proper use of passwords:
 - sufficiently long
 - contains non-alphabetic characters
 - unguessable
 - changed often
 - never written down
 - never shared
- attacks:
 - in an **offline dictionary attack**, the attacker knows the complementation functions and stored complementary information:
 - * e.g. has the encrypted password file
 - * repeatedly guesses different passwords and applies the functions
 - * real dictionary attacks use probability of words being used as passwords
 - in an **online dictionary attack**, the attacker guesses directly into the system, without other previous knowledge:
 - * with **backoff**, systems increase the time between interactions with more tries
 - * with **disconnection**, the connection is broken after a number of failures
 - * with **disabling**, the account is disabled
 - * with **jailing**, the user gets false access to a limited part of the system
 - can also **honeypot** the system with false data to trap attackers
 - modern machines are very fast, so even with salting, huge dictionaries can be checked against encrypted passwords:
 - * GPUs excel at password cracking
 - * even salted, hashed passwords are not safe
- password management:
 - limit login attempts:
 - * prevents dictionary attacks “over the wire”
 - * lock account, slow down, etc.
 - encrypt passwords:
 - * store unencrypted passwords as briefly as possible e.g. no temp files
 - * same with password attempts into a log file, etc.
 - * passwords should be sent over HTTPS

- protect the password file
 - * make dictionary attacks more difficult
- for forgotten passwords, should generate new passwords
 - * site should never be able to send back forgotten passwords, implies that there is a way to decrypt encrypted passwords
- transporting new passwords:
 - * generally sent encrypted via email or text message
 - * both are compromisable
 - * some banks require surface mail
- user passwords:
 - * using same vs. different passwords for sites
 - * security vs. usability
 - * password vaults, write down passwords
- another authentication mechanism is **challenge / response**:
 - authenticate based on questions you can answer correctly i.e. what you know
 - * e.g. security questions, or smart card
 - can ask for different information every time
 - * or challenge the hardware to perform something e.g. encrypt it with a unique key
 - security depends on encryption of the challenge
 - question is too hard to answer without special hardware, or too easy for intruders to spoof the answer
 - smart card details:
 - * cryptography should be performed only on smart card
 - * user should enter password into card
 - *cons*:
 - * if lost or stolen, can't authenticate, and maybe someone else can
 - * susceptible to sniffing attacks
 - * requires special hardware
- **biometrics** is another mechanism based on who you are:
 - fingerprints, voice patterns, retinal patterns, etc.
 - to authenticate, allow system to measure physical characteristics
 - * biometrics converted to digital
 - interplay vs. false positive and false negatives:
 - * more sensitivity means lower false positive rate, but also higher false negative rate
 - * the **crossover error rate (CER)** is the point where the rates meet
 - * for usability, false negatives are very undesirable
 - *good use cases*:
 - * use them for authentication with clean readings
 - * when biometric readers themselves are secure
 - * when attacks are rare or difficult

- * together with other authentication
- *poor use cases:*
 - * working off low-quality / noisy readings
 - * finding “needles in haystacks”
 - * when biometric reader is easy to bypass or spoof
 - anything across a network is suspect
- *cons:*
 - * requires very special hardware
 - * not as foolproof as you might think
 - * generally not helpful for authentication programs or roles
 - * many physical characteristics vary too much for practical use
- authentication by where you are
 - requires sufficient proof of physical location and ability to tie a device at that location to its messages
- **multifactor authentication:**
 - something you know + something you have
 - * at least one factor needs to be non-replayable
 - e.g. PIN + ATM card, password + phone
 - either can go wrong for a false negative
 - are the factors really orthogonal?
 - are both factors non-trivial?
 - is one factor likely to suffer a catastrophic break?

Operating Systems

- what does the OS protect?
 - authentication for operating systems
 - memory protection e.g. buffer overflows
 - IPC protection e.g. covert channels
 - stored data protection e.g. full disk encryption
- the OS provides the lowest layer of software visible to users:
 - close to hardware, often with complete hardware access
 - OS flaws compromise all security at higher levels
 - OS controls memory, scheduling, devices, other resources
 - systems may be single user, multiple user, embedded with no human user
 - * all still require OS security
 - almost all other security systems must assume a secure OS at the bottom
- security *depends* on running the right OS and version, not altered by an attacker:
 - i.e. **trusted computing**
 - need trusted hardware that makes sure the boot program behaves and runs the right OS:
 - * AKA **security enclaves**
 - * hardware implementation is challenge, often has known flaws
 - the **trusted platform module (TPM)** is special hardware designed to improve OS security:
 - * proves OS was booted with a particular bootstrap loader using tamperproof hardware and cryptographic techniques
 - * provides secure key storage and crypto support
 - * checks signatures of the OS etc.
 - * bootloader and users can request TPM to verify applications or OS
 - * not *guaranteed* security, but creates a chain of transitive trust
 - TPM hardware is widely installed, but not widely used:
 - * e.g. Microsoft Bitlocker, secure Linux boot loader
 - * Microsoft's SecureBoot is another build software alternative that only boots systems with pre-arranged digital signatures
- authentication and authorization in OS:
 - OS must *authenticate* all user requests
 - human users log in locally and remotely, and processes run on their behalf
 - once authenticated, requests must be *authorized*
 - remote user authentication timeline:
 1. user authenticates via password, public key crypto, sometimes a

- particular process, etc.
- 2. successful login creates a primal process under ID of logged in user
- 3. OS ties a process control block to the process with owner ID
- 4. process can fork off more processes
 - * invoking system calls checks owner IDs through **reference monitors**
 - * special system calls can change a process's ID
- how often should OS perform authorization?
 - * passing operations through reference monitors add overhead
 - * balance between overhead and necessary authorization
 - * e.g. only on first check, incrementally, periodically, etc.
- protecting memory:
 - memory contains executable code, copies of permanently stored data, and temporary process data
 - virtual memory provides a logical separation of processes:
 - * for error containment more so than security
 - * main memory divided into page frames, every processes has an address space divided into logical pages
 - * each process is given a table, and all addressing goes through the page table at the hardware level
 - * a process shouldn't be able to name other processes' pages
 - security issues of page frame reuse:
 - * OS switches ownership of page frames as necessary
 - * when a process acquires a new page frame, can the process read the old page frame data?
 - need to clean page:
 - * e.g. zero on deallocation, zero on reallocation, zero on use, clean pages in background
 - * Linux zeroes on reallocation, Windows cleans in background
 - **buffer overflow** is one of the common causes for compromises of operating systems:
 - * process messing with its own memory, running different code by changing the function return address:
 - i.e. choosing what gets written into the instruction pointer
 - programs often run on behalf of others, so this is dangerous
 - * can be interpreted as a flaw in OS input processing, programming languages, or even programmer training
 - * **stack overflow** is a kind of buffer overflow intended to alter the contents of the stack
 - * **heap overflow** does not offer the direct ability to jump to arbitrary code (heap is mostly non-executing), but potentially quite dangerous
 - fixing buffer overflows:

- * write better code
- * use programming languages that prevent them
- * add OS controls that prevent overwriting the stack
- * put things in different places on the stack
- * don't allow execution from places in memory where overflows occur
- protecting interprocess communications:
 - OS provides various kinds of IPC e.g. messages, semaphores, shared memory, sockets
 - *possible exploits*:
 - * convince system process is another process
 - an authentication problem
 - * can break into another process's memory
 - handled by page tables
 - * forge a message from someone else
 - OS tags IPC with identities
 - * eavesdrop on someone else who gets the secret
 - related to page reuse and internal OS buffers
 - mostly secure, but hard for certain scenarios:
 - * bug in the OS
 - * not a single machine
 - depends on strong authentication and authorization
 - * OS has to prevent cooperating processes from sharing information
 - * process wants to communicate with another process, but OS has been instructed to prevent that e.g. mandatory access control
 - in **covert channels**, we use something not ordinarily regarded as a communications mechanism to actively attempt to deceive the OS:
 - * e.g. disk activity, page swapping, time slice behavior, use of a peripheral device
 - * only need to send 0's and 1's
 - * very difficult to detect
- stored data protection:
 - files are a typically shared resource
 - data stored on disk is subject to many risks:
 - * if OS protections are bypassed, how can we protect data?
 - * store data in encrypted form
 - *issues*:
 - * when does cryptography occur?
 - which files, explicitly or implicitly, how long decrypted, where does it exist in decrypted form
 - * where does the key come from?
 - human user, file system, smart card, disk hardware, where and how long do we store

- * what is the granularity of cryptography?
 - disk, file system, block
- *practicality*:
 - * for improper users, why not just use access control
 - * no point in hiding from OS
 - * for data transfers, encrypt while in transit
 - * someone who physically accesses the device not using the OS
 - only relevant attack that encryption protects against
- in full disk encryption:
 - * all data on the disk is encrypted
 - data is encrypted and decrypted as it enters and leaves disk
 - * prevents improper access to stolen disks
 - * could be done in hardware or software

Network Security

- degree of locality:
 - some networks are very local e.g. Ethernet
 - benefits:
 - * physical locality
 - * small number of users and machines
 - * common goals and interests
 - other networks e.g. Internet are very non-local
 - * many users and sites share bandwidth
- network media e.g. wires, cables, telephone lines can be physically protected
 - satellite links and radio links have more limited *physical* protection possibilities
- implication of protocol type:
 - protocol defines a set of rules that will always be followed
 - specific attacks exist against specific protocols
- threats to networks include wiretapping, impersonation, confidentiality and integrity attacks, DoS attacks:
 - **passive wiretapping** is listening in illicitly on conversations
 - **active wiretapping** is injecting traffic illicitly
 - **packet sniffers** can listen to all traffic on a broadcast medium
 - wiretapping on wireless is often just putting up an antenna
 - message can be read or even altered at intermediary gateways and routers
 - * typically requires access to part of the path the message takes
 - in denial of service, legitimate users are prevented from doing their work by flooding the network or corrupting routing tables or flooding routers or destroying key packets
 - * all-inclusive nature of the Internet makes basic access trivial, universality of IP makes this easy
- SYN flood attack:
 - attacker uses initial request and response to start enough TCP sessions to fill a table that is used to keep track of connections at the server
 - * sends a bunch of SYN requests, without acknowledging the SYN/ACK
 - prevents new real TCP sessions
 - * server cannot delete half-open connections in case we have a slow, real client
 - can defend with SYN cookies and firewalls along with large tables
 - SYN cookie approach:
 - * when table is almost full, server sends back a SYN/ACK, *without*

- creating a new table entry, that contains a cookie
 - * cookie value is a secret function of various information e.g. client/server address and port, timer
 - * store the cookie as the sequence number itself!
 - no need to change the protocol to support cookies
 - * server doesn't need to save cookie values
 - * slows down attacker greatly since he would need to create full connections to actually take up space in the table
- **distributed denial of service (DDoS) attacks:**
 - send a large volume of packets from a large number of distributed machines
 - * no need to target a particular exploit like TCP tables
 - distribution harnesses multiple machines and makes defenses harder
 - if more packets sent than can be handled by target e.g. link or server, service is denied
 - * could be pure flooding, or overwhelming of CPU or memory resources, direct or reflected
 - *complications:*
 - * high availability of compromised machines
 - * Internet is designed to deliver traffic
 - * IP spoofing allows easy hiding
 - * distributed nature makes legal approaches hard
 - * attackers can mimic normal packets
 - *defense approaches:*
 - * overprovisioning
 - * dynamic increases in provisioning
 - * filtering
 - * traffic redirection e.g. content delivery networks
 - * reducing volume of attack
 - * none of these are totally effective

Traffic Control Mechanisms

- in **source address filtering**, we filter out some packets because of their source address value:
 - AKA ingress or egress filtering, address assurance
 - usually because we believe their address to be spoofed
 - router knows what network it sits in front of:
 - * filter outgoing packets with source addresses not in its range
 - * prevents users from spoofing other nodes' addresses, but not from spoofing each others
 - can also be done in the other direction, as packets leave the Internet and enter a border router

- * only prevents spoofed IPs that are in the local network (these packets should have never left the local network, so we can safely drop them)
- other forms of filtering e.g. worm signatures, unknown protocol identifiers, unallocated IP addresses, local use addresses only
 - can also redirect packets to a special filtering site on the edge of the network:
 - * expressively designed to deal with DDoS attacks with aggressive filtering criteria
 - * incurs serious delay penalties
- realistic limits on filtering:
 - little filtering possible in Internet core:
 - * packets handled too fast
 - * backbone providers typically don't want to filter
 - filtering near edges is also limited in terms of possibility, affordability, what router owners will do
- many routers can place **limits** on the traffic they send to a destination:
 - limits defined flexibility
 - often not good enough to differentiate good and bad traffic
- to better hide traffic characteristics, we can use **padding** to add extra traffic
 - fake traffic must look like real traffic
- similarly, use ability to control message routing to conceal the traffic in the network
 - i.e. using **onion routing** to hide who is sending traffic to whom for anonymization purposes
- a **firewall** is a machine to protect a network from malicious external attacks:
 - running special software to regulate network traffic and control entry and exit points
 - * examines each incoming packet and decide to let the packet through or not
 - a form of security called **perimeter of defense**
 - breaching the perimeter compromises all security
 - part of the solution, but not the entire solution
 - * i.e. defense in depth by combining different defenses
 - *types*:
 - * filtering gateways AKA screening routers
 - * application level gateways AKA proxy gateways
 - * reverse firewalls
- filtering gateways filter based on packet header information:
 - IP addresses can always be spoofed
 - * firewall should not always trust packet headers
 - can filter based on ports to drop packets sent to little-used ports
 - *pros*:

- * stateless
- * fast, cheap, flexible, transparent
- *cons*:
 - * limited capabilities
 - * dependent on header authentication
 - * generally poor logging
 - * may rely on router security