

CS174: Computer Graphics

Professor Law

Thilan Tran

Fall 2021

Contents

CS174: Introduction to Computer Graphics	3
Basic Graphics System	4
Output Devices	4
Modeling Objects	5
Transformations	7
Inverses	10
Combining Transformations	11
Rendering Pipeline	12
World and Camera	12
Projections	13
Window to Viewport Mapping	16
Hidden Surface Removal	17
Back Face Culling	17
Painter's Algorithm	18
Z-Buffer Algorithm	18
Lighting	21
Shading	24
Interpolations	25
Mappings	26
Appendix	27
Linear Algebra Review	27

Graphics Tips & Tricks	30
----------------------------------	----

CS174: Introduction to Computer Graphics

- computer graphics is the field of creating imagery by computer
 - used in the entertainment industry eg. movies and games, visualization applications, etc.
- basic elements of CG include:
 - **modeling** ie. mathematically representing objects:
 - * constructing models for specific objects
 - * makes use of 3D points, lines, curves, surfaces, polygons
 - volumetric vs. image-based representations
 - * primitives can have attributes like colors and texture maps
 - **animation** ie. representing the motions of objects:
 - * also give animators control of the motion
 - * eg. keyframe animation, motion capture, procedural animation
 - **rendering** ie. simulating the real-world behavior of light and formation of images:
 - * simulate light propagation
 - includes 3D scene, lighting, point-of-view, shading, projection
 - * properties include reflection, absorption, scattering, emission, interference
 - **interaction** ie. enabling humans and computers to interact

Basic Graphics System

- a basic graphics system includes:
 - input devices
 - CPU and GPU:
 - * CPU takes input from the user, and calculates how the world should update accordingly
 - * CPU passes to GPU the world changes, and GPU generates an image that can be stored as a **framebuffer**
 - computing and rendering system
 - output devices
 - * takes in framebuffer and displays it on the screen scanline by scanline

Output Devices

- in **cathode ray tube (CRT)** technology, electrons strike phosphorous coating and emit light:
 - direction of electron beam controlled by deflection plates
 - with random-scan or vector CRT, can only “burn” lines into the screen:
 - * deflection plates randomly target different parts of the screen to burn in points
 - * difficult to fill polygons, etc.
 - refresh rate between 60 to 85 Hz
 - * burning only lasts a few milliseconds
- in **raster CRT**, the screen is broken up into pixels eg. n by m phosphorous cells:
 - a pixel becomes the smallest element we can modify on the screen
 - intrinsically has a rasterizing or aliasing problem due to finiteness of pixels
 - framebuffer depth determines color complexity:
 - * 1 bit supports black and white
 - * 8 bits gray scale
 - * 8 bits per color (RGB) requires 24 bits and creates 16 M colors
 - * 12 bits per color is HD
 - 3 different colored electron guns
 - * each pixel has 3 different colors phosphors arranged in triads
 - a shadow mask helps to prevent electron beams from bleeding over into neighboring pixels
 - **interlaced** displays update odd to even scanlines:

- * rather than every scanlines, as in **non-interlaced** displays
- * human eye cannot catch the difference, while data bandwidth requirement is halved since only half of the framebuffer needs to be sent
- potential race condition with the framebuffer:
 - * as GPU is updating the framebuffer, the display may be reading it at the same time
 - * instead, **double buffering** uses two framebuffers to avoid this data conflict
- typical screen resolutions:
 - TV is 640x480
 - HD is 1920x1080
 - 4K LCD is 3840x2160
 - 35mm is 3000x2000
- memory speed and space requirements:
 - given $n \times m$ resolution, refresh rate r Hz, color depth b bits per pixel
 - memory space per second is $\frac{n*m*b*r}{8}$ bytes
 - if non-interlaced, memory read time is $\frac{1}{n*m*r}$ seconds per pixel
 - if interlaced, memory read time is $\frac{2}{n*m*r}$ seconds per pixel
- flat screen displays no longer have a CRT:
 - still raster based, with an active matrix of transistors at grid points
 - * vertical grid plus horizontal grid of wires allows voltages to be customized along the grid to change the lighting at each pixel
 - light sources for pixels may be:
 - * **light emitting diodes (LEDs)**
 - * **polarized liquid crystals (LCDs)**
 - * **plasma**, where gases are energized to glow

Modeling Objects

- a sphere may be most easily modeled by its origin and radius:
 - however, this is not easiest model to render in graphics
 - * complicated non-linear formulas
 - we only know how to render polygons
 - need to **discretize** the *surface* of the object into polygons:
 - * essentially linearizing shapes into line segments
 - * loses information of the inside of the object
 - eg. to tessellate a circle, can arrange many triangles around the origin
- polygons can be represented as collection of points connected with lines:
 - eg. for vertices v_1, v_2, v_3, v_4 , the connecting edges are assumed to be $v_1v_2, v_2v_3, v_3v_4, v_4v_1$

- * to model, represent with a list of vertices with their coordinates and then a list of faces with their vertices
- the **normal** of a polygon should face outward from the face
 - * thus vertices of faces are typically ordered counter-clockwise
- considerations include:
 - * closed / open
 - whether the last closing edge is inferred
 - * wireframe / filled
 - generally only consider wireframe polygons (discretized)
 - * planar / non-planar:
 - non-planar polygons span multiple planes
 - thus the normal points in different directions on the same polygon
 - * convex / concave:
 - concave polygons have internal angles that are greater than 180°
 - thus the normal points in the wrong direction at some concave vertices
 - * simple / non-simple
 - non-simple polygons intersect on themselves
- if we deal with only triangles, we satisfy many desired polygonal considerations:
 - simple, convex, and planar
 - modern GPUs can render 100 million triangles per second
- how can we check if a point is inside a polygon?
 - a point is *inside* a convex polygon if it lies to the *left* of all the directed edges
 - * in a convex polygon, the directed edges point counter-clockwise and the normals of two consecutive edges are always consistent
 - for concave polygons, these useful properties do not hold
 - * want to break up concave polygons into convex ones

Transformations

- **translations** can be written as:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} T_x \\ T_y \end{bmatrix} + \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} T_x + x \\ T_y + y \end{bmatrix}$$

- **scalings** can be written as:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} S_x & 0 \\ 0 & S_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} S_x x \\ S_y y \end{bmatrix}$$

- note that scaling happens with respect to the origin:
 - * scaling greater than one moves away from the origin, while scaling less than one moves towards the origin
 - * negative scales flip across the axis
- **rotations** can be written as the following, where θ is the angle of rotation and r is the distance of the initial point from the origin:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} x \cos \theta - y \sin \theta \\ x \sin \theta + y \cos \theta \end{bmatrix}$$

- rotations also happen with respect to the origin
- counter clockwise rotations are positive, while clockwise rotations are negative
- **shears** can be written as:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & a \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} x + ay \\ y \end{bmatrix}$$

- this is a horizontal shear
- **transforming lines and polygons:**
 - if we can *preserve* the linearity of the line, we only have to transform the endpoints of the line or polygon instead of all the contained or connected points
 - affine transformations on endpoints guarantee that straight lines will remain straight lines after the transformation!
 - * ie. affine transformations preserve affine combinations eg. line segment interpolations
 - affine transformations also preserve planarity, parallelism, and relative ratios of edge lengths

- these transformations can be *stacked* by appending matrices together:
 - eg. $M_R \times M_S \times M_S \times M_R \times \dots \times M_P = M_T \times M_P$
 - multiplying onto the left, AKA post-multiplication:
 - * note that order matters!
 - * except for special cases, such as pure transformations of a single type being commutative
 - eg. pure translations, pure scalings, pure rotation about an axis
 - can create a single transformation matrix out of the individual transformations
 - * to all points, apply the single transformation matrix
 - however, translation is a matrix addition rather than a matrix multiplication:
 - * cannot coalesce into a single transformation matrix!
 - * however, with homogeneous representation, we can bypass this issue and treat translation as matrix multiplication
- homogeneous transformations:

$$T = \begin{bmatrix} 1 & 0 & T_x \\ 0 & 1 & T_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x + T_x \\ y + T_y \\ 1 \end{bmatrix}$$

$$S = \begin{bmatrix} S_x & 0 & 0 \\ 0 & S_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} S_x x \\ S_y y \\ 1 \end{bmatrix}$$

$$R = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x \cos \theta - y \sin \theta \\ x \sin \theta + y \cos \theta \\ 1 \end{bmatrix}$$

$$Sh_x = \begin{bmatrix} 1 & a & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x + ay \\ y \\ 1 \end{bmatrix}$$

- 3D transformation matrices:

$$T = \begin{bmatrix} 1 & 0 & 0 & T_x \\ 0 & 1 & 0 & T_y \\ 0 & 0 & 1 & T_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$S = \begin{bmatrix} S_x & 0 & 0 & 0 \\ 0 & S_y & 0 & 0 \\ 0 & 0 & S_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$R_z = \begin{bmatrix} \cos \theta & -\sin \theta & 0 & 0 \\ \sin \theta & \cos \theta & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$Sh_x = \begin{bmatrix} 1 & a & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- note that this rotation is happening *around* the z-axis, since the z-coordinates are unchanged
- note that this shear is a purely horizontal shear along the x-axis
- the **general rotation matrix (GRM)** is a shortcut rotation matrix, given the orthonormal unit vectors of the rotated basis, say i, j, k , to rotate them *back* onto the normal axes:

$$GRM = \begin{bmatrix} i_x & i_y & i_z & 0 \\ j_x & j_y & j_z & 0 \\ k_x & k_y & k_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- looking at the 3×3 submatrix in the upper left, the first row is exactly i , the second row is exactly j , and the third row is exactly k
 - * these are the θ projections used in the rotation transformation
- to inverse, since orthonormal, we can just take the transpose, where the basis vectors are column vectors
 - * this inverse is exactly how to transform normal axes *into* the rotated basis ie. change of basis transformation
- consider the 3×3 submatrix in the upper left of each of these 3D transformations:
 - each of the rows in this submatrix can be taken as a vector, and the dot product of any of these two rows in the submatrix equals 0

- an affine transformation has this property, where the upper left 3×3 submatrix is orthogonal
- translations and rotations are **rigid body transformations** since the lines, angles, and distances between points do not change:
 - * the upper left 3×3 submatrix is orthonormal ie. vectors are orthogonal and unit vectors
 - * note that scaling and shears are *not* rigid body transformations, since the vectors are not normalized (orthogonal, but *not* orthonormal)
 - * for orthonormal matrices, $A^{-1} = A^T$
- rotations around the other axes:

$$R_x = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos \theta & -\sin \theta & 0 \\ 0 & \sin \theta & \cos \theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$R_y = \begin{bmatrix} \cos \theta & 0 & \sin \theta & 0 \\ 0 & 1 & 0 & 0 \\ -\sin \theta & 0 & \cos \theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- shear along the x and y-axes:

$$\begin{bmatrix} 1 & 0 & a & 0 \\ 0 & 1 & b & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- the z-axis is locked, and shearing occurs on the other two axes

Inverses

- inverse transformations:

$$T^{-1} = \begin{bmatrix} 1 & 0 & -T_x \\ 0 & 1 & -T_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x - T_x \\ y - T_y \\ 1 \end{bmatrix}$$

$$S^{-1} = \begin{bmatrix} \frac{1}{S_x} & 0 & 0 \\ 0 & \frac{1}{S_y} & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} \frac{x}{S_x} \\ \frac{y}{S_y} \\ 1 \end{bmatrix}$$

$$R^{-1}(\theta) = R(-\theta)$$

$$Sh_x^{-1}(a) = Sh_x(-a)$$

Combining Transformations

- how can we rotate by θ around an arbitrary point (x_R, y_R) , instead of just the origin?

$$M = T(x_R, y_R)R(\theta)T(-x_R, -y_R)$$

1. translate reference point to the origin
 2. rotate
 3. translate reference point back to its original location
- how can we scale an object by k in place?

$$M = T(x_R, y_R)S(k)T(-x_R, -y_R)$$

- normally, scaling only happens with respect to the origin:
 - * the object will move away from its original location if it is not centered
 - * we want to scale the object while retaining its position
- 1. translate reference point (eg. bottom left corner of object) to the origin
- 2. scale
- 3. translate reference point back to its original location
- how to find a point's new coordinates after a change of basis ie. frame of reference?
 - given the basis vectors defining the new basis, simply apply the inverse transformation in the change of basis to the point to get new coordinates
 - eg. if we have a new frame of reference with origin $(6, 2)$, we can translate all points by $(-6, -2)$ to get new coordinates
 - * similarly for stacked transforms eg. rotated and translated basis, have to inverse translate and then inverse rotate
- how to rotate a point by θ around a vector (with base point P_R)?

$$M = T(P_R)R_y(-\phi_y)R_z(-\phi_z)R_x(\theta)R_z(\phi_z)R_y(\phi_y)T(-P_R)$$

1. first, we need to align the vector to an axis
 - move vector base to origin, then rotate other axes so that the vector lies on some axis eg. x-axis
2. rotate around the aligned axis
3. move vector to original location
 - rotate other axes, and then translate back

Rendering Pipeline

World and Camera

- we need to define a **camera** or reference frame through which to view our world:
 - we can build a camera coordinate system that can be represented by a matrix
 - requires an **eye vector** to represent the eye direction, as well as a **top vector** to describe the tilt of the camera
 - three orthogonal vectors can be generated as follows:
 1. eye vector is the first vector
 - * can also be represented by an **eye point** and **reference point**
 2. then, we need a vector orthogonal to the plane of the eye and top vectors, so we can take the cross product
 3. take the cross product of the first two bases

- camera basis vectors:

$$k = \frac{P_{ref} - P_{eye}}{|P_{ref} - P_{eye}|}$$
$$i = \frac{v_{up} \times k}{|v_{up} \times k|}$$
$$j = k \times i$$

- general graphics pipeline:
 1. first, we want to place various models into our **worldspace**:
 - we perform various transformations such as translations, rotations, etc.
 - these are collapsed into a single transformation matrix that is applied to all vertices
 2. then, to view this world through a camera ie. eye view:
 - we need to adjust all the coordinates into the camera's frame of reference
 - * ie. moving eye view onto the origin
 - need to translate the eye point back to the origin, and rotate the camera bases
 - * essentially, eye is at origin, looking down the z-axis, and head is upright

- however, the x-axis from the eye's *point of view* has been flipped!
 - * points left instead of right
 - * we need to mirror across the yz-plane
- use the following transformation matrix:

$$M = Mirror_x \times GRM(i, j, k) \times T(-P_{eye})$$

- note that we can combine these matrices from different stages together!
 - * transformation matrix (TM), eye matrix (EM), and an upcoming projection matrix (PM) can be pre-multiplied to avoid unnecessary matrix multiplies
- the mirror matrix on the yz-plane is as follows:

$$Mirror_x = \begin{bmatrix} -1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Projections

- now that everything is in eye space, we still have to perform something called a **projection**:
 - need to take a 3D object in space, and capture it *onto* a plane to display to the user
 - draw projection lines from a center of projection to the vertices of the object in space
 - * the intersection of these projected vertices to the projection plane forms the projection image
 - two types of projections:
 - * **parallel projection**, where the eye is placed at infinity away
 - view volume is a parallelepiped
 - * **perspective projection**, where the eye's location gives a sense of depth to the projection:
 - view volume is more of a truncated pyramid, with **clipping planes** of different sizes that cut off the viewing volume
 - a viewing angle in x and y determine the slopes of this pyramid
 - * having a front clipping plane prevents division by zero errors
 - during projection, usually divide by distance from eye
- in a parallel projection, all we have to do is throw away the z-coordinate:

$$Parallel_M = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- the view volume is a parallelepiped of certain dimension
 - * any vertices lying outside of the box should be clipped and not rendered
- we can create a unit or normalized **parallel canonical view volume** that extends between -1 and 1 in the xy-plane and from 0 to 1 in the z-direction
- we can map an arbitrary view volume into a canonical one with the following normalized matrix:

$$Normal_M = \begin{bmatrix} \frac{2}{W} & 0 & 0 & 0 \\ 0 & \frac{2}{H} & 0 & 0 \\ 0 & 0 & \frac{1}{F-N} & \frac{-N}{F-N} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- * given width, height, far distance, and near distance of view volume
- in a perspective projection, we need to take into account where the projection plane lies between the eye and object position, say a distance d from the eye:

$$Perspective_M = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & \frac{1}{d} & 0 \end{bmatrix}$$

- simply use similar triangles in a ratio, we know that $\frac{x'}{d} = \frac{x}{z}$ ie. $x' = \frac{x}{z}d$
- when applied to an arbitrary point, after normalization, the projected point is as desired:

$$Perspective_M \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ z \\ \frac{z}{d} \end{bmatrix} = \begin{bmatrix} \frac{x}{z}d \\ \frac{y}{z}d \\ d \\ 1 \end{bmatrix}$$

- * z-coordinates are lost, collapsed into a constant d
- how can we handle non-square projections of other aspect ratios?
 - given parameters aspect ratio $A_R = \frac{W}{H}$ and half angle of view in the x-axis $\theta_x = \theta$
 - * note that $\tan \theta = \frac{W}{2d}$, where d is the distance of the projection screen from the eye
 - we can normalize our previous x' to a canonical view volume by dividing by $\frac{W}{2}$ to get $x' = \frac{x}{z} \frac{2d}{W}$, and replace d with θ
 - similarly, we can do the same for the y-axis and use the aspect ratio to remove H and get the following:

$$x' = \frac{x}{z \tan \theta}$$

$$y' = \frac{y}{z} \frac{2d}{H} = \frac{y A_R}{z \tan \theta}$$

- * this constrains our new axes between $[-1, 1]$, as in the canonical view volume
- * note that we are indeed dividing by the distance from the eye through z
- the perspective projection matrix for any aspect ratio is as follows:

$$Perspective_{M_{AR}} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & A_R & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & \tan \theta & 0 \end{bmatrix}$$

- when applied to an arbitrary point, after normalization, the projected point is as desired:

$$Perspective_{M_{AR}} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ A_R y \\ z \\ z \tan \theta \end{bmatrix} = \begin{bmatrix} \frac{x}{z \tan \theta} \\ \frac{A_R y}{z \tan \theta} \\ \frac{1}{\tan \theta} \\ 1 \end{bmatrix}$$

- however, our z values are now constant, and the depth data is lost for future calculations
- importantly, note that this perspective division step, where we normalize our points, cannot be encapsulated in a 4×4 matrix
 - transformation, eye, and projection matrices can be combined together
 - but perspective division has to be done separately, before the window-to-viewport mapping step
- to fix the loss of our depth during perspective division, we could recover our z values as follows:

$$\begin{aligned} z' &= A + \frac{B}{z} \\ A &= \frac{F}{F - N} \\ B &= -\frac{NF}{F - N} \end{aligned}$$

- given the distance of the near and far planes, N and F
- forms a system of equations where $z' = 0$ if $z = N$, and $z' = 1$ if $z = F$
- the full normalized perspective projection matrix is as follows, which maps directly into a parallel canonical view volume that expresses height in z from 0 to 1 (rather than a flat plane with no depth):

$$Perspective_{M_{AR}} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & A_R & 0 & 0 \\ 0 & 0 & A \tan \theta & B \tan \theta \\ 0 & 0 & \tan \theta & 0 \end{bmatrix}$$

- * note that after multiplying with the projection matrix, perspective division ie. normalization must be done

Window to Viewport Mapping

- the final rendering stage is to map our viewing window to a viewport of arbitrary size:

$$M = T(V_L, V_B)S\left(\frac{V_R - V_L}{W_R - W_L}, \frac{V_T - V_B}{W_T - W_B}\right)T(-W_L, -W_B)$$

- requires viewport left, right, bottom, and top or V_L, V_R, V_B, V_T , respectively
 - * as well as the same values for the window, if not normalized to the canonical window

Hidden Surface Removal

- through **hidden surface removal**, we remove parts of certain objects that are obscured by other objects:
 - saves on computations for graphics that would never be seen on the screen
 - basic operations:
 1. establish priorities among objects
 2. collect overlapping elements and resolve visibility using the priorities
 - different algorithm types:
 - * **image space** algorithms perform both operations at pixel resolution
 - * **list-priority** algorithms perform (1) at object resolution and (2) at pixel resolution
 - * **object / world / eye space** algorithms perform both operations at object resolution
 - different surface removal algorithms may not support all object types
 - * polymesh, free form, volume, CSG, implicit
 - algorithm considerations include flexibility, memory, speed

Back Face Culling

- in **back face culling**, we consider the normals of faces:
 - only works for polymesh objects
 - as soon as possible, we should throw away **back faces** in which the normal faces away from the eye
 - * for any polygon, on average, half of the faces are not visible
 - to determine back faces in world space:
 - * take the dot product of the normal and the vector created from any point on the face towards the eye
 - * if dot product is positive, the face is a front face, otherwise it is a back face
 - when zero, the face exactly “grazes” the eyesight, and we can still ignore it
 - in eye space, we can perform the same calculation with the eye placed at the origin
 - * note that the calculation is simplified since we can just take a dot product of the point and normal, and if negative, the face is a front face
 - in projection space (after perspective division), the eye is now placed at

negative infinity after projecting objects onto a flat screen:

- * calculation is even simpler, front faces have a negative z-component
- * all other faces are back faces

Painter's Algorithm

- the **painter's algorithm** is an object space algorithm that runs as follows:
 1. sort all polygons by z-depth
 2. scan-convert polygons in back to front order
 - ie. take polygons through the pipeline one at a time, and override the pixels on the screen accordingly
 - cannot handle cyclic or intersecting polygons!
 - * need to break up such polygons into the simpler case

Z-Buffer Algorithm

- the **z-buffer algorithm** is an image space algorithm that runs as follows:
 - maintain a depth and color buffer containing information on each pixel
 - * initialized to infinite distance and desired background color
 - 1. for each pixel covered by a polygon:
 - calculate depth z for the polygon at the pixel (x, y) using plane equations
 - if z is less than the entry in the depth buffer, update the depth and color buffers accordingly with the new z and color of the polygon
 - note that order of polygon processing doesn't matter, since we are proceeding pixel by pixel
- properties:
 - image precision algorithm
 - easy to implement in software and hardware
 - polygons scan-converted into framebuffer in random order
 - * no pre-sorting of polygons is necessary
 - *pros*:
 - * handles cyclic and intersecting polygons
 - * extends to various faces other than polygons, all we need is a depth calculation
 - * simplicity and ease of implementation
 - * can be modified to reduce memory requirements
 - * can be extended to A-buffer to reduce aliasing
 - * theoretically can handle any number of polygons

- *cons*:
 - * buffer memory requirements, not much of an issue in modern PCs
 - * aliasing issues ie. smooth edges or lines become jagged due to the pixel granularity
 - depends on screen resolution
 - * complexity depends on polygons' projection area on the screen
 - * difficult to handle transparency, since blending of colors must be done in a specific order
- **area weighted averaging** is one technique to perform anti-aliasing:
 - the overall color of a pixel is a weighted average combination of the polygon and background colors
 - * weighted by area of the pixel covered by the polygon vs. background
 - resultant color is a blend of colors, creates a softer transition on the polygon edges
 - pixel area calculations are expensive operations
- the **scanline z-buffer algorithm** is an adaptation on the algorithm that reduces the memory usage:
 - proceed one scanline at a time
 - * requires all polygons overlapping the current scanline, rather than any single polygon at a time
 - 1. for each scanline at height y :
 - 2. for each polygon which intersects scanline:
 - scan convert for specific scanline and determine span segments
 - 3. for each pixel in span segment at location x :
 - calculate depth z for polygon at pixel (x, y)
 - if z is less than the entry in the depth buffer, update the depth and color buffers
 - *pros*:
 - * less memory requirement than z-buffer, proportional to window width instead of area
 - *cons*:
 - * multiple passes through polygon database
- z-buffer speedup considerations:
 - bounding box testing:
 - * associate y_{min}, y_{max} for each face (step 2)
 - * calculate left and right ends ie. x-intercepts with scanlines (step 3)
 - can subdivide the entire screen into spaces to parallelize
 - can subdivide hierarchically (some screen regions more busy than others) to parallelize

- incremental depth calculation along a scanline:

$$\begin{aligned} Ax + By + Cz + D &= 0 \\ z_{x,y} &= -\frac{Ax + By + D}{C} \\ z_{x+1,y} &= -\frac{A(x+1) + By + D}{C} \\ &= z_{x,y} - \frac{A}{C} \end{aligned}$$

Lighting

- in global illumination, reflected light reaches the eye to create an image
 - light energy is bounced around or scattered between various surfaces of objects
 - * this scattering is called **radiosity**
 - **reflection** of objects on other surfaces is a different phenomenon
 - in graphics, we instead trace rays of light from the eye onto objects in order to color pixels:
 - * rather than illuminating unnecessary pixels from the light source that may not eventually reach the eye
 - * in addition, fully simulating radiosity is an enormously expensive calculation
 - instead, we approximate illumination through different types of lighting
 - types of lights:
 1. ambient light
 2. spot light
 3. directional light with no location eg. sun
 4. spotlight with location and direction:
 - * has some spotlight angle β
 - * only points hit by the spotlight within β are illuminated
 - * shoot a ray from the spotlight location to the point, and illuminate if $\alpha < \beta$ ie. $L \cdot D > \cos \beta$
- lighting depends on various properties:
 - geometric:
 - * position and orientation of object or eye
 - * position, direction, point vs. spot vs. area of light
 - material:
 - * color, reflectivity, shininess, translucency, bumpiness of object
 - * color of light
 - * filter or color blindness of eye
- lighting types include ambient, diffuse, and specular lighting
- **ambient** lighting is produced from a background light:
 - unrealistic, only outlines the silhouette of the object
 - good approximation of scattered light
 - does not depend on position or orientation of light, object, or eye
 - * only object and light's material property

- $L_A = k_a \cdot I_a$, ambient reflection coefficient times intensity of ambient light source
- **diffuse** lighting is produced from a point light source:
 - AKA Lambertian reflection for dull, matte surfaces
 - surfaces look equally bright from all directions, regardless of where the eye is placed
 - Lambert’s Law states that the amount of light is reflected is proportional to the angle between the incident light and the normal:
 - * $L_D = k_d I_p \cos \theta = k_d I_p (N \cdot L)$
 - * diffuse reflection coefficient times intensity of point light source times dot product of normal and incident light
 - ignore negatives by “clamping” at zero
 - * negatives indicate **self-occlusion**
 - for a light at sufficient distance, L is constant for the entire scene, and $N \cdot L$ is constant on a polygon
 - can attenuate the diffuse lighting depending on distance:

$$L_D = f_{att} k_d I_p (N \cdot L)$$

$$f_{att} = \frac{1}{d^2}, f_{att} = \frac{1}{c_1 + c_2 d + c_3 d^2}$$

$$f_{att} = \min(f_{att}, 1)$$

- * in atmospheric attenuation, immerse object in “fog” past a certain distance by blending colors
- **specular** lighting is created off of “shiny” surfaces:
 - does depend on location of the eye
 - color is of the light source, not object
 - light not reflects *unequally* in different directions
 - $L_S = f_{att} k_s I_p \cos^n \theta = f_{att} k_s I_p (R \cdot V)^n$
 - * attenuation factor times specular reflection coefficient times intensity of point light source times dot product raised to specular reflection exponent
 - n is the material’s specular reflection exponent, higher indicates a more shiny surface
 - * ie. metallic surfaces are closest to a perfect reflector which have a very bright, concentrated reflection
 - R is the reflection vector, V is the view vector from the point to the eye:
 - * R is equally reflected off of N by the angle between N and L
 - * $R = 2(L \cdot N)N - L$
 - * is there a way we can approximate $R \cdot V$ with something else?

- use a halfway vector H that is halfway between L and V , and use $N \cdot H$ instead
 - this is Blinn-Phong illumination, faster but less accurate
 - a faster alternative also exists for calculating t^n
- exponentiating a term that is less than 1 draws it closer to 0:
 - * higher exponent indicates a smaller region where the point light reflection is aligned with the viewer, leading to a smaller shiny spot
 - * should clamp the cos to zero
- overall equation for illuminated color at a point:

$$I_\lambda = k_a I_a O_{d\lambda} + f_{att} k_d I_p (N \cdot L) O_{d\lambda} + f_{att} k_s I_p (R \cdot V)^n$$

- I_λ is overall illuminated color (with three parts RGB), $O_{d\lambda}$ is the object's diffuse color
- there may be an additional emissive term k_e if the object itself creates light
- to handle multiple point light sources, just add up the diffuse and specular terms for each point light
- k, O, I, n are all material properties, while remaining variables are geometric properties
- note that this illumination equation does not take into account shadows from other objects
 - * requires additional energy conservation equations involving ray tracing and radiosity

Shading

- now that we know how to light certain points, how do we determine which points on our world to illuminate?
 - infinite possible points to light up
 - polygons have a color and normal
 - * while vertices have position, color, normal, as well as texture coordinates
 - typically, illumination is done in world space, and interpolation is done in screen space
- in **flat / faceted shading**, each polygon is shaded and illuminated uniformly:
 1. use N of polygon
 2. find I_λ at center or any vertex of polygon
 3. apply that same color to all points inside polygon
 - means that:
 - * normal is constant across polygon
 - * light is at infinity so $N \cdot L$ is constant across polygon
 - * viewer is at infinity so $N \cdot V$ is constant across polygon
 - no interpolation is done
 - *pros*:
 - * simple
 - *cons*:
 - * drastic jump in illumination at the edges
 - to the human eye, the jump is perceived as even stronger
- in **Gouraud smooth shading**, we illuminate each vertex:
 - AKA intensity interpolation or color interpolation shading
 - how to find the normal for a vertex?
 - * average the normals of the neighboring polygons
 - * to emphasize an edge, additionally bias the average with that polygon normal
 - 1. find color at each vertex and illuminate in world space
 - 2. in screen space, interpolate ie. smoothen vertex colors across the polygon face
 - *pros*:
 - * no more jump in illumination on polygon edges
 - *cons*:
 - * for large polygons, we only considered the vertices, when the faces can have important illumination consequences
 - eg. accidentally skipping specular highlights in the middle of a polygon or unnecessarily interpolating specular highlights from a vertex

- note that the polygonal geometry is not being smoothed out, only the coloring
 - * though smooth shading helps to create an illusion of a smoother object
- how to solve some of the Gouraud shading drawbacks?
 - illuminate *every* pixel inside a polygon
 - * this is **Phong shading**
 - but we have to illuminate in world space (but no notion of pixels in world space), instead of screen space (no notion of lights, object colors, eye):
 - * need to perform an inverse transformation while in screen space back to world space to color each pixel
 - * interpolate the normals at vertices in screen space to get a normal to use for coloring

Interpolations

- given a polygon with colored vertices, how can we interpolate and smoothen out the colors across the polygon?
 - in 1D, we can perform linear interpolation using an affine combination of colors ie. $C = (1 - t)C_1 + tC_2$
 - in 2D, we want to perform **bilinear interpolation**
 - * an alternative interpolation technique is **barycentric interpolation**, which uses percentage areas from vertices
 - can be used to generally interpolate colors, depth, normals, texture coordinates, etc.
 - * any values given at vertices can be interpolated
- bilinear interpolation algorithm for a certain point:
 1. draw a horizontal scanline through the point
 2. perform two linear interpolations for color along the left and right intersecting polygon edges of the scanline
 - gives us x_a, y_a, λ_a and x_b, y_b, λ_b
 3. perform a final linear interpolation for the point on the left and right colored intersecting points
 - but this algorithm will use different vertices to interpolate depending on how the object rotates
 - * however, when using triangles, the same vertices are always utilized
 - in Gouraud shading, when trying to interpolate colors, we will also run into issues where specular lighting details are lost

Mappings

- through mappings, we adjust the shading or sometimes positions of parts of an object based on some kind of image map
- **texture / pattern mapping** is the process of mapping a digitized image onto a poly face:
 - individual elements of a texture are **texels** rather than pixels
 - * usually mapped onto an s, t plane with $0 \leq s, t \leq 1$
 - given u, v coordinates for a texture, we want to use the texel color as the diffuse color to color the pixel
 - if we sample the texture colors as a low rate, we may miss out on high-frequency patterns in the texture
 - * to fix this aliasing issue, we can sample at a higher resolution, or perform area averaging
 - can also apply multiple textures on the same object and blend them
- displacement and bump mappings distort the 3D appearance of the object given a map:
 - these maps are greyscale images that tell points how far to move
 - in **displacement mapping**, we displace the actual point on the surface of the object, and change the object's geometry:
 - * $P' = P + BN$ ie. bump along the normal
 - * points actually move, requires additional hidden surface removal
 - in **bump mapping**, we distort the normal to simulate a bump without altering the object's geometry:
 - * $N' = \frac{B_u(N \times P_t) - B_v(N \times P_s)}{|N|}$
 - * B_u, B_v, P_s, P_t are all partial derivatives
- with **environmental / reflection mapping**, we use polar coordinates of the reflected ray to map:
 - quick way to graphically simulate reflections
 - * ray tracing allows for true reflections
 - map usually defined on faces of a box or sphere
 - add a reflection term $k_r I_r$ to the illumination equation
- with **procedural mapping**, we use some formula to procedurally adjust point displacement, color, etc.
 - eg. create impression of water waves without full liquid simulations

Appendix

Linear Algebra Review

- **points** have a location, but no size, shape, or direction
 - lie on a coordinate plane
- **vectors** have a direction and length, but no location:
 - can define a vector along two basis vectors (in 2D)
 - vectors v_i, \dots, v_m are **linearly independent** if $a_1 v_1 + \dots + a_m v_m = 0$ iff. $a_i = 0$
 - * ie. no projection of one vector on any of the others
 - **linear dependent** vectors are scalar multiples of each other
- a difference between two points is a vector $v = Q - P$
 - similarly, a base point plus vector offset is another point $Q = P + v$
- the **homogeneous representation** for points and vectors allows us to distinguish between the two:

– a point is represented as $\begin{bmatrix} P_x \\ P_y \\ 1 \end{bmatrix}$

– a vector is represented as $\begin{bmatrix} V_x \\ V_y \\ 0 \end{bmatrix}$

- similarly for 3D, we have a 4th element to distinguish the two
- we can now define vectors and points in matrix multiplication:

$$v = \beta_1 v_1 + \beta_2 v_2 + \beta_3 v_3 = [\beta_1 \ \beta_2 \ \beta_3 \ 0] \begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ P_0 \end{bmatrix}$$

$$P = P_0 + v = P_0 + \alpha_1 v_1 + \alpha_2 v_2 + \alpha_3 v_3$$

$$= [\alpha_1 \ \alpha_2 \ \alpha_3 \ 1] \begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ P_0 \end{bmatrix}$$

- more operations in homogeneous representation:

$$v + w = \begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ 0 \end{bmatrix} + \begin{bmatrix} w_1 \\ w_2 \\ w_3 \\ 0 \end{bmatrix} = \begin{bmatrix} v_1 + w_1 \\ v_2 + w_2 \\ v_3 + w_3 \\ 0 \end{bmatrix}$$

$$av + bw = a \begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ 0 \end{bmatrix} + b \begin{bmatrix} w_1 \\ w_2 \\ w_3 \\ 0 \end{bmatrix} = \begin{bmatrix} av_1 + bw_1 \\ av_2 + bw_2 \\ av_3 + bw_3 \\ 0 \end{bmatrix}$$

$$P + v = \begin{bmatrix} p_1 \\ p_2 \\ p_3 \\ 1 \end{bmatrix} + \begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ 0 \end{bmatrix} = \begin{bmatrix} p_1 + v_1 \\ p_2 + v_2 \\ p_3 + v_3 \\ 1 \end{bmatrix}$$

$$P - Q = \begin{bmatrix} p_1 \\ p_2 \\ p_3 \\ 1 \end{bmatrix} - \begin{bmatrix} q_1 \\ q_2 \\ q_3 \\ 1 \end{bmatrix} = \begin{bmatrix} p_1 - q_1 \\ p_2 - q_2 \\ p_3 - q_3 \\ 0 \end{bmatrix}$$

- linear combination in homogeneous representation:

$$aP + bQ = a \begin{bmatrix} p_1 \\ p_2 \\ p_3 \\ 1 \end{bmatrix} + b \begin{bmatrix} q_1 \\ q_2 \\ q_3 \\ 1 \end{bmatrix} = \begin{bmatrix} ap_1 + bq_1 \\ ap_2 + bq_2 \\ ap_3 + bq_3 \\ a + b \end{bmatrix}$$

- if affine, $a + b = 1$, and the combination creates a point
- if $a + b = 0$, the combination creates a vector
- otherwise, we can normalize the result so that the last element is 1
 - * combination still creates a point
- a **vector space** is a space defined with respect to certain **basis vectors**:
 - eg. in 2D, we need two bases in order to define any unique vector
 - the magnitudes in the direction of the basis vectors, added together, defines any unique vector
 - * eg. $v = v_x \vec{v}_1 + v_y \vec{v}_2 = v_x i + v_y j$
 - basis vectors do not necessarily have to be orthogonal, or even unit magnitude:
 - * however, cannot be on the same line ie. linear dependent
 - * good practice to have unit basis vectors to only specify direction
- a **generator set** is a set of vectors that generate a vector space:
 - for a vector space \mathbb{R}^n we need minimum n vectors to generate all vectors

- a generator set with minimum size is called a **basis** for the specified vector space
 - * basis is purely defined by vectors, and creates a vector space that only supports vectors
- a **frame** has a point of origin along with a basis, and creates an **affine space** that supports vectors and points
- the **dot product** of two vectors is defined as $v_1 \cdot v_2 = |v_1||v_2| \cos \theta$:
 - alternatively $v_{1_x}v_{2_x} + v_{1_y}v_{2_y}$
 - a scalar value
 - when the dot product is 0, the vectors are orthogonal
 - when the dot product is negative, the angle is greater than 90 degrees
 - when the dot product is positive, the angle is less than 90 degrees
 - $|u|\cos\theta$ gives the projection of vector u on in the direction of v
- the **cross product** of two vectors gives a vector:

$$v = v_1 \times v_2 = \det \begin{bmatrix} i & j & k \\ v_{1_x} & v_{1_y} & v_{1_z} \\ v_{2_x} & v_{2_y} & v_{2_z} \end{bmatrix}$$

- resultant vector is perpendicular to the plane of the two vectors, pointing as defined by the right-hand rule
- $|a \times b| = |a||b| \sin \theta$
- not commutative
- **polygons** can be defined as a set of directed edges or connected vectors
 - the vectors can be calculated as the differences between the connected points
- **lines** can be written in parametric form as $P = P_1 + \alpha \vec{d}$:
 - where $d = P_2 - P_1$ and P_1, P_2 are the endpoints of the directed line pointing towards P_2
 - equivalently, $P = (1 - \alpha)P_1 + \alpha P_2$
 - traces ie. interpolates a line between the two endpoints that is infinite in both directions
- adding points has no meaning, but by linearly combining them with additional constraints, we can interpolate useful constructs:
 - $P = \alpha_1 P_1 + \alpha_2 P_2$ is a **linear combination**
 - with the condition $\alpha_1 + \alpha_2 = 1$, the parametric equation becomes an **affine linear combination**

- * represents a point lying on the line passing through P_1, P_2 that is infinite in both directions
 - with the *additional* condition that $\alpha_i \geq 0$, the parametric equation becomes a **convex linear combination**:
 - * represents a point on the line segment between P_1, P_2
 - * note convex necessitates affine
 - if we only have the constraint $\alpha_i \geq 0$ (not affine), we have a **ray** that is infinite in one direction only
- consider defining a polygon in terms of parameteric form as $P = \alpha_1 P_1 + \dots + \alpha_n P_n$:
 - without any constraints, we cannot guarantee that the interpolated points are on the same plane as the polygon i.e. only a linear combination
 - with constraint $\sum \alpha_i = 1$, we have an affine combination, and the interpolated points will lie on the same plane as the polygon
 - with additional constraint $\alpha_i \geq 0$, we have a convex combination, and the interpolated points will lie within the convex hull of the polygon
- the **convex hull** can be imagined as taking a string around pegs at each corner of the polygon:
 - technically, the smallest convex polygon that contains all the points of the actual polygon
 - can be larger than the actual drawn polygon

Table 1: Summary of Scalar, Point, and Vector Operations (* affine only)

Operands	Add	Subtract	Multiply
point-point	$P = a_1 * P_1 + a_2 * P_2$ (*)	$v = P_2 - P_1$	
vector-vector	$v = v_1 + v_2$	$v = v_1 - v_2$	
scalar-point			$P = s * P_1$ (*)
scalar-vector			$v = s * v_1$
point-vector	$P_2 = P_1 + v_1$	$P_2 = P_1 - v_1$	

Graphics Tips & Tricks

- to transform lines:
 1. described by 2 end points
 - if we are performing an affine transformation, we can simply transform the end points and connect the line since the points will re-

- main collinear
- 2. described by equation $y = mx + b$:
 - find two points, transform them, and connect the line
 - for translations, we can simply adjust b
 - for rotations, we can simply adjust m
- to transform planes:
 1. described by 3 non-collinear points
 - if we are performing an affine transformation, transform the points and draw the new plane
 2. described by plane equation $Ax + By + Cz + D = 0$:
 - or $A(x - x_i) + B(y - y_i) + C(z - z_i) = 0$
 - the normal is (A, B, C)
 - if M_{point} is the matrix to transform a point, $M_{normal} = (M_{point}^T)^{-1}$
 - * for rigid body transformations, $M_{normal} = M_{point}$
- point in polygon test:
 1. if convex, check if point lies to the left of every edge
 2. extend a semi-infinite ray from the point:
 - if there are an odd number of intersections with the polygon, the point is inside, else outside
 - also works for concave polygons
 3. perform angle summation from the point to each pair of vertices:
 - if the sum of subtended angles is 360° , the point is inside, else outside
 - also works for concave polygons
- on-line test:

$$\frac{x - x_1}{y - y_1} = \frac{x_2 - x_1}{y_2 - y_1}$$

$$T(P) = (x - x_1)(y_2 - y_1) - (x_2 - x_1)(y - y_1)$$

- if the test value is 0, the point P lies on the line connected through P_1, P_2
- if positive, P is on the right
- if negative, P is on the left
- edge intersection test:
 - given two edges represented by 4 points, can we determine if the edges intersect?
 - check if one endpoint of an edge is to the left of the other edge, and the other endpoint is to the right of that edge
 - * and vice versa with the other edge's endpoints
 - ie. the products of the on-line tests for the endpoints against the other edge should be negative

- collinearity test:

$$\begin{aligned}
 t &= |P_1 P| \sin \theta \\
 &= \frac{|P_1 P| |P_1 P_2| \sin \theta}{|P_1 P_2|} \\
 &= \frac{|P_1 P \times P_1 P_2|}{|P_1 P_2|}
 \end{aligned}$$

- t represents the shortest distance from a point P to an edge P_1, P_2
- θ is the angle between P_1, P_2 and P_1, P
- if t is sufficiently small, we can say the point is collinear
- calculating a normal vector:
 1. give 3 consecutive convex vertices, simply find the cross product
 2. use summation method, which works for convex and concave polygons:

$$(\sum (y_i - y_j)(z_i + z_j), \sum (z_i - z_j)(x_i + x_j), \sum (x_i - x_j)(y_i + y_j))$$

- transformation matrices:

$$M = \begin{bmatrix} m_{11} & m_{12} & m_{13} & m_{14} \\ m_{21} & m_{22} & m_{23} & m_{24} \\ m_{31} & m_{32} & m_{33} & m_{34} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- the upper-left 3×3 matrix defines rotations, shears, and scalings
- (m_{14}, m_{24}, m_{34}) defines translations
- orthogonal transformation matrices:
 - eg. translations and rotations
 - for the upper-left 3×3 matrix in a 4×4 transformation matrix:
 - * each row is a unit vector, and each row is orthogonal to the others
 - * can be thought of rotating these vectors to align with the xyz-axis
 - determinant is 1
 - $M^{-1} = M^T$
 - if orthonormal as well, preserves angles and lengths ie. is a rigid body transformation