

CS213A: Embedded Systems

Professor Srivastava

Thilan Tran

Fall 2021

Contents

CS213A: Embedded Systems	2
Embedded Hardware Platforms	4

CS213A: Embedded Systems

- an **embedded system** is a computer system with a dedicated function within a larger mechanical or electrical system, often with real-time computing constraints:
 - vs. general-purpose computers that support a wider range of needs and typically do not have physical world interaction
 - eg. home thermometers, smart speakers, smart watches, Roombas, self driving technologies:
 - * embedded systems make rich inferences about contexts, life patterns, behaviors, intents
 - * diverse domains include health, transportation, security, entertainment, etc.
 - embedded trends include networked, miniaturization, human-coupling, AI-enabled, security and hacking concerns
- these embedded systems are typically connected to the Internet, ie. **Internet of Things (IoT)** devices:
 - IoT devices are these embedded computing devices that are interconnected in the Internet
 - * allows for advanced connectivity of devices, systems, services
 - these “things” now dominate the edge of the internet rather than traditional computers
 - unique challenges include security, privacy, and configuration
 - *passive* IoT devices unilaterally upload sensor data
 - * while *active* IoT devices upload data while also reading and analyzing results
 - common three-tier IoT systems:
 - * the actual **embedded system** or device that senses, actuates, buffers, processes, provides UI
 - * a **gateway** that relays, senses, buffers, processes, provides UI
 - * the **cloud services** that archive, process, disseminate
- other related concepts:
 - **ubicomp** or ubiquitous computing, where computing can be made to appear anytime and everywhere in everyday life
 - **Cyber-Physical Systems (CPS)** have perception-cognition-action loops:
 - * integrates computation, networking, sensing, and physical processes
 - * build upon embedded systems with modeling of physical processes and their dynamics

- **edge computing** and **fog computing** aim to extend the cloud closer to the devices producing and consuming the data:
 - * edge computing is a paradigm where much processing occurs on the end device, rather than in the centralized cloud
 - * fog computing is a similar paradigm where processing occurs on network nodes near the end device

Embedded Hardware Platforms

- a generic embedded system platform has the following components:
 - human interface to interact with users
 - sensors and actuators to interact with the physical world
 - network and I/O interface to interact with servers and other devices
 - computation and storage system
 - power subsystem
- computation options:
 - custom SoC with ASIC or FPGA
 - custom board with COTS components
 - custom system using standard chassis like PCI, peripheral cards, etc.
 - use standard platform eg. ruggedized PC
 - specific ICs typically take least power and are fastest, while microprocessors and more flexible computation options take more power and are less fast
 - * in terms of cost, mass volume production of specific ICs or SoCs will bring down their prices significantly
 - microcontroller units (MCUs) are the most common option:
 - * microcontrollers package a microprocessor (CPU) with RAM, ROM, I/O, etc.
- power management options:
 - lower-power reduced-functionality modes (dynamic power management):
 - * eg. idle, stop clock, power off, etc.
 - * save power by changing modes
 - frequency scaling (dynamic frequency scaling) requires an appropriate clocking subsystem
 - * power is proportional to frequency
 - voltage and frequency scaling (dynamic voltage-frequency scaling)
 - * power is proportional to frequency and voltage squared
 - ideally, we want high performance for intensive workloads:
 - * while conserving energy as workloads mitigate
 - * can switch scaling type depending on workload type
 - * one approach is big-little, where we have two cores that can be used depending on workload eg. by tracking a weighted average of CPU load
- I/O options:
 - inputs can include keyboard, mouse, pointing devices, speech

- * outputs can include monitor, projector, displays, audio
 - challenges include performance requirements eg. data rate, latency, synchronization
 - * pin constraints, power supply requirements
 - typical supported API are get & set parameters, transmit & receive data, enabled & disable functions
 - I/O compared to data and program memory:
 - * I/O device registers can't cache
 - * timing matters
 - * reading and writing can trigger hardware to perform some function
 - * sharing across applications can be challenging
 - one technique is to have specialized I/O instructions just for I.O
 - * restrict to kernel mode
 - another technique is to have memory-mapped I/O:
 - * access devices just like they are memory
 - * while using memory protection mechanisms
- pin types:
 - **general purpose I/O (GPIO)** pins are digital pins that can be read and written by software
 - * eg. control external peripherals, LEDs, software serial protocols, etc.
 - analog I/O pins can be read and written by software:
 - * AI pins perform analog to digital conversion at a certain bit resolution
 - * AO pins perform digital to analog to digital conversion
 - * **pulse width modulation (PWM)** is another way to convey analog values on a single digital pin using varying duty cycles
- serial interface protocols:
 - *pros*: serial interfaces reduce the pin counts, avoiding timing issues
 - *cons*: may need clock recovery and are difficult to memory map
 - eg. UART / RS232, I2C, SPI
 - synchronous vs. asynchronous is whether data is sent with or without a clock
 - full duplex vs. half duplex is whether data can be sent and received simultaneously or not
 - in a master / slave approach, one device is the master and others are slaves
 - * usually synchronous, as the master supplies the timing clock for data in both directions
 - in a multi-master bus, the master / slave bus may have more than one master
 - * needs an arbitration scheme
 - in a point-to-point or peer interface approach, there are no masters or

slaves

- * usually asynchronous
- in a multi-drop approach, there is one transmitter and several receivers
- in a multi-point approach, there are more than two peer transceivers allowing for bidirectional data exchange
- UART and RS-232 are a common serial interface found on MCUs and peripherals:
 - products that support RS-232 are called **universal asynchronous receiver transmitters (UARTs)**
 - unbalanced bus capable of full-duplex communication between two receiver / transmitter pairs
 - speeds up to 115.2 Kbps, with distance up to over 200 feet
 - typically consists of a start bit, data bits, and parity or stop bits (optional)
 - logic high typically is a negative voltage, while logic low is a positive voltage
 - communication protocol:
 1. start bit indicates a character is coming, and receiver resets its timers
 2. then, the receiver samples in the middle of the bits using a timer
 - * transmission rate ie. bit width must be known!
 3. stop bit should be 1
- the **inter-integrated circuit (I2C)** bus is another popular, patented interface:
 - half-duplex, synchronous, multi-master bus requiring only two signal wires, data (SDA) and clock (SCL)
 - these lines are pulled high via pull-up resistors and controlled by the hardware via open-drain drivers
 - I2C is an addressable communications protocol:
 - * allows master to communicate with individual slaves
 - * 7 or 10-bit address
 - master initiates communication and generates all clock signals for data transfer in either direction
 - * through clock stretching, the slave can pull SCL low if it needs more time to process data bits
 - three speeds include slow or 100 Kbps, fast or 400 Kbps, and high-speed or 3.4 Mbps
 - * each speed is downward compatible
 - communication protocol:
 1. master generates a start condition
 - * start condition is a high-to-low transition of SDA while SCL is high
 2. master sends an 8-bit data word
 - * data bit transitions take place while SCL is low

3. the receiver generates the acknowledge bit:
 - * pulls SDA low while master releases the line and allows it to float high
 - * if the acknowledge bit is read by the master as high, it should consider the last communication word as not received
 4. master generates a stop condition or a repeated start
 - * stop condition is a low-to-high transition of SDA while SCL is high
- the **serial peripheral interface (SPI)** is a synchronous serial bus with multi-master / slave protocol:
 - present on many microcontrollers
 - four signals are master out slave in (MOSI), master in slave out (MISO), serial clock (SCK), and active-low slave select (/SS)
 - allows for data rates over 1 Mbps in full duplex mode
 - generally limited to on-board communications and short traces
 - data is clocked simultaneously into slave and master based on SCK pulses:
 - * master supplies the SCK
 - * allows for four different clocking types based on polarity and phase of SCK
 - typically requires many pins when using multiple slaves
 - other I/O interfaces:
 - Microwire similar to SPI
 - 1-wire
 - USB, which allows devices to switch back and forth between host and device